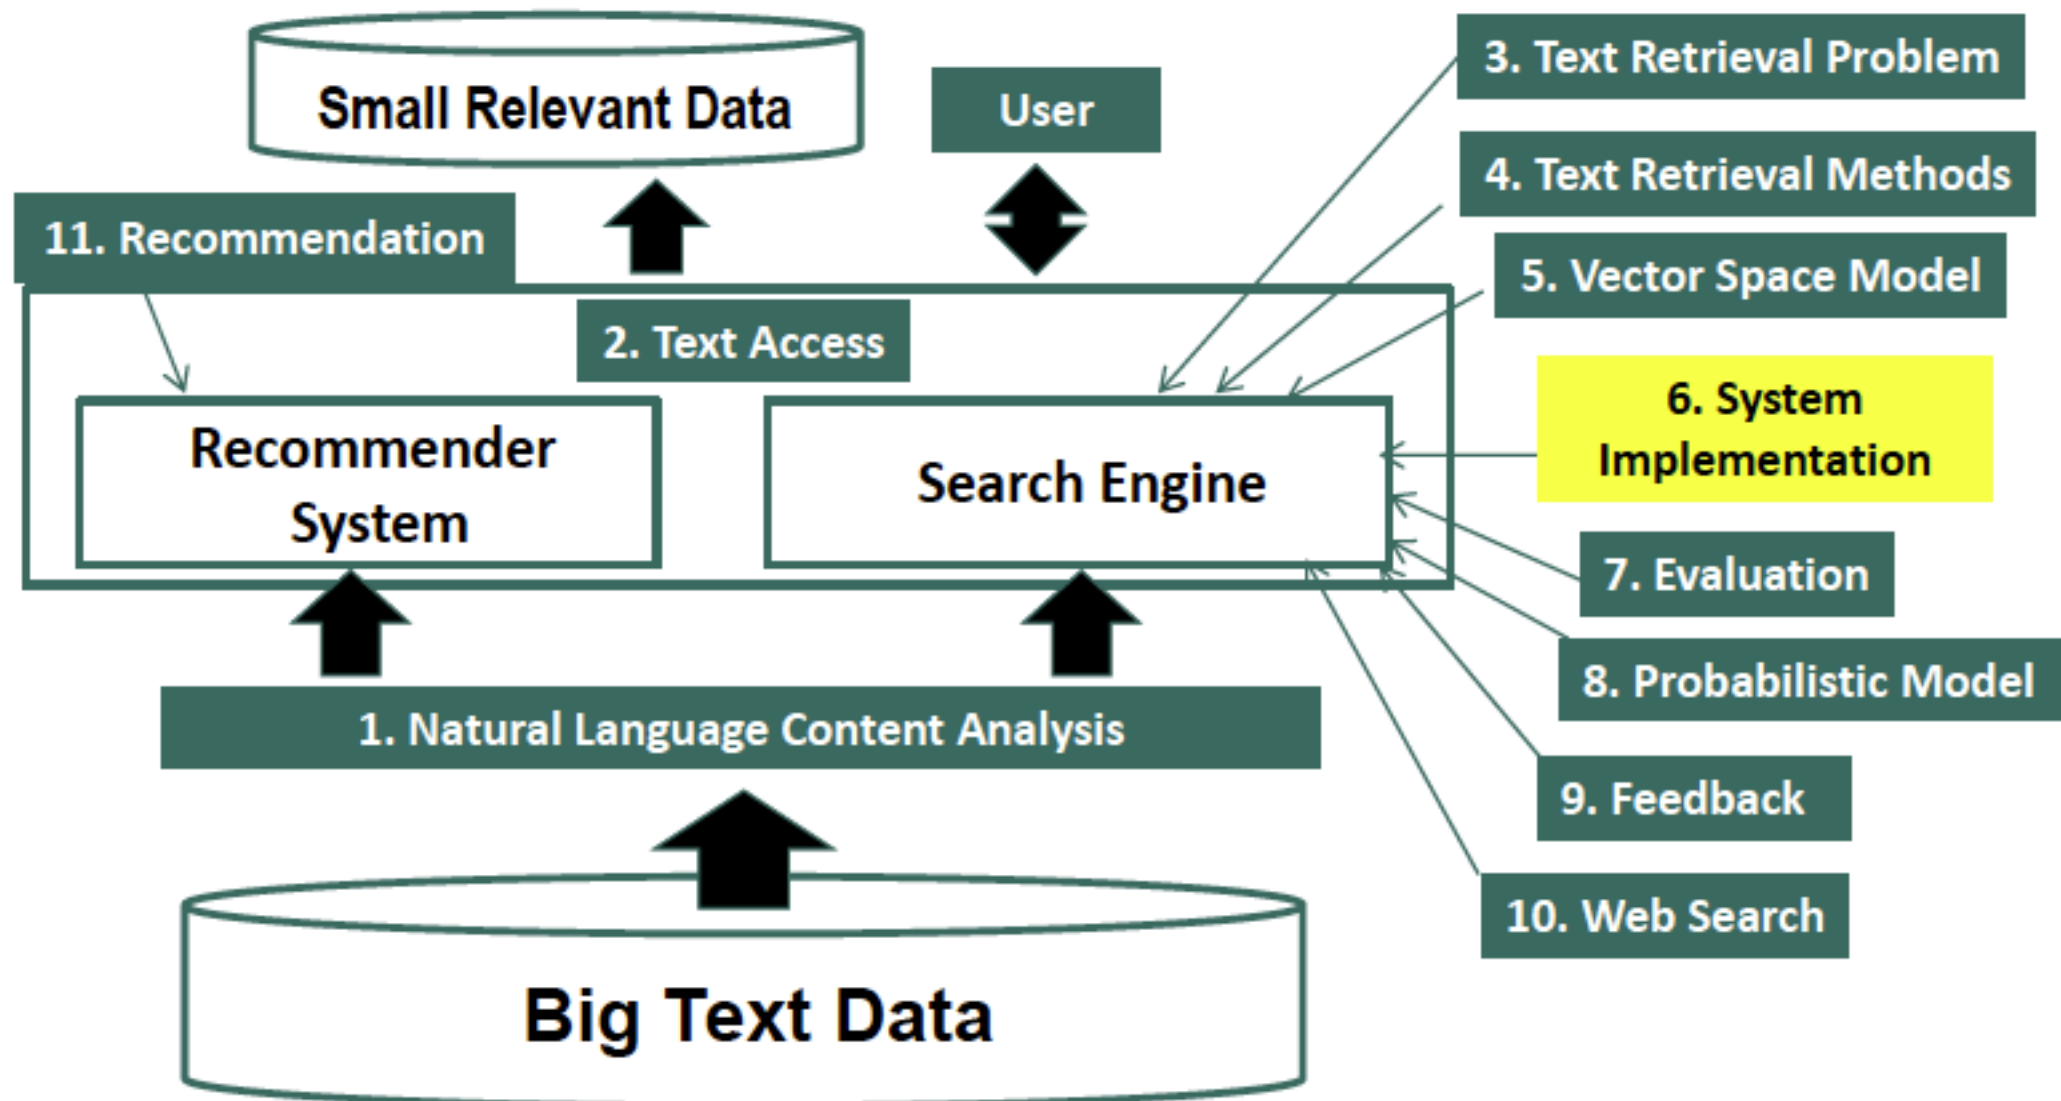


Information Retrieval & Text Mining

System Implementation
Inverted Index Construction

Dr. Saeed Ul Hassan
Information Technology University

Implementation of Text Retrieval Systems



Constructing Inverted Index

- The main difficulty is to build a huge index with limited memory
- Memory-based methods: not usable for large collections
- Sort-based methods:
 - Step 1: Collect local (termID, docID, freq) tuples
 - Step 2: Sort local tuples (to make “runs”)
 - Step 3: Pair-wise merge runs
 - Step 4: Output inverted file

Sort-based Inversion

doc1



doc2



...

doc300



Mapping Strings to integer>>

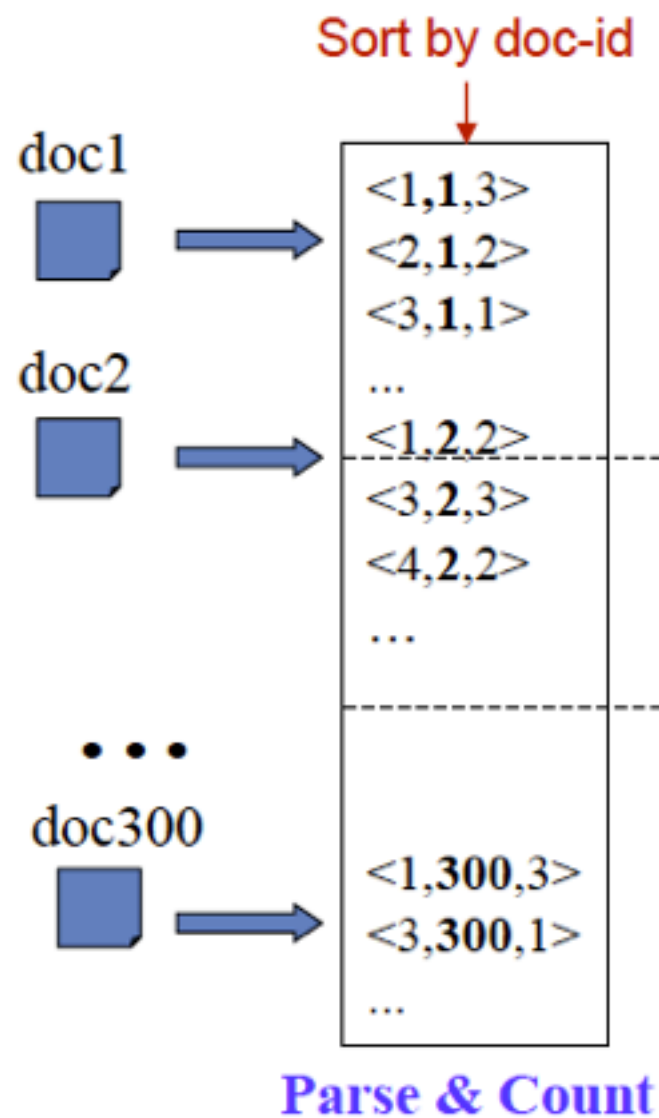
**Term
Lexicon:**

the 1
campaign 2
news 3
a 4
...

**DocID
Lexicon:**

doc1 1
doc2 2
doc3 3
...

Sort-based Inversion



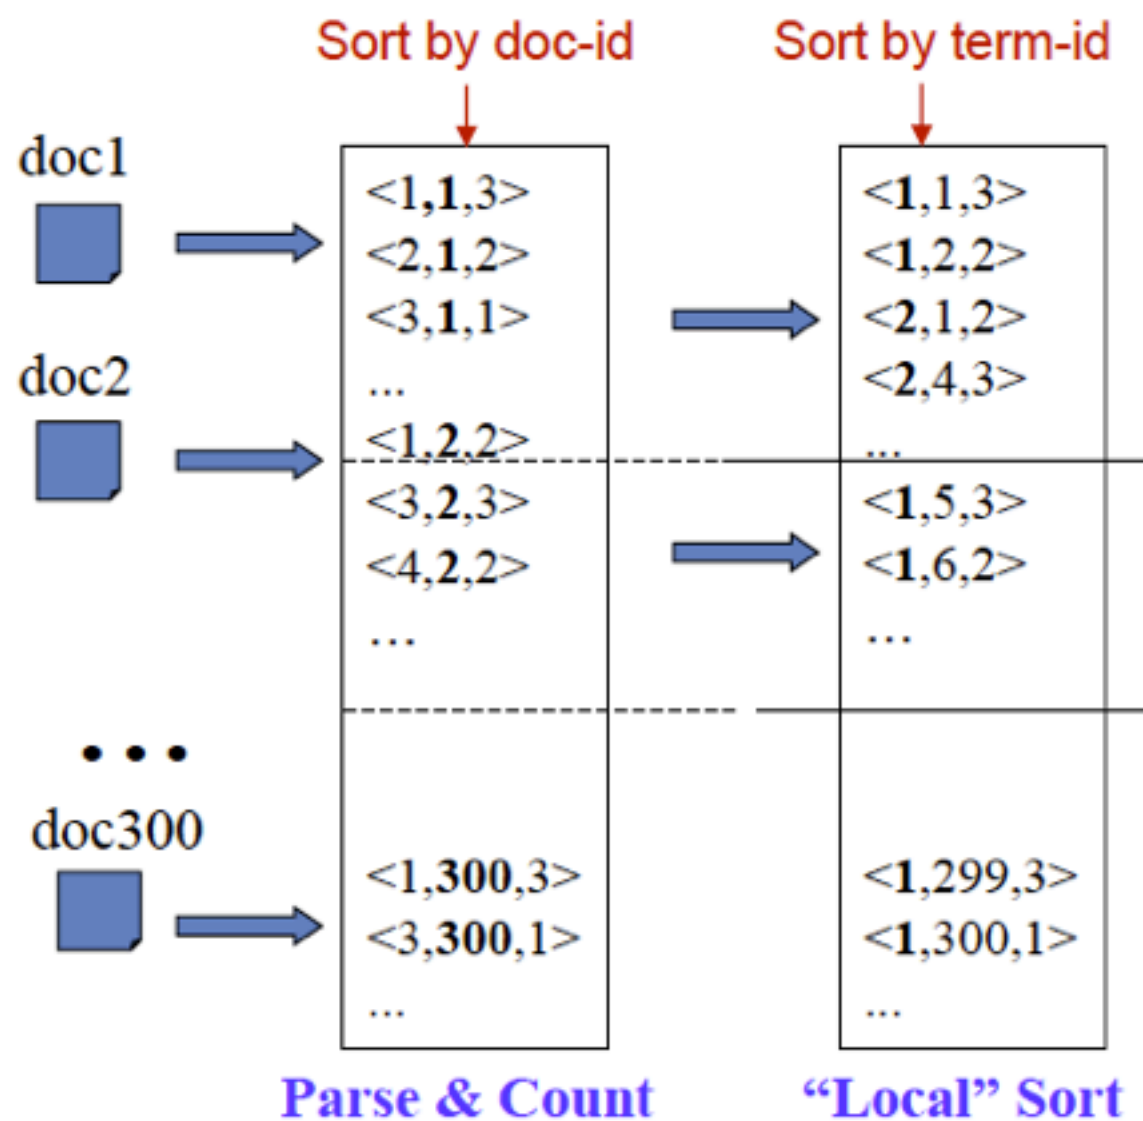
**Term
Lexicon:**

the 1
campaign 2
news 3
a 4
...

**DocID
Lexicon:**

doc1 1
doc2 2
doc3 3
...

Sort-based Inversion




Term Lexicon:

the 1
campaign 2
news 3
a 4
...

DocID Lexicon:

doc1 1
doc2 2
doc3 3
...

Sort-based Inversion



Inverted Index Compression

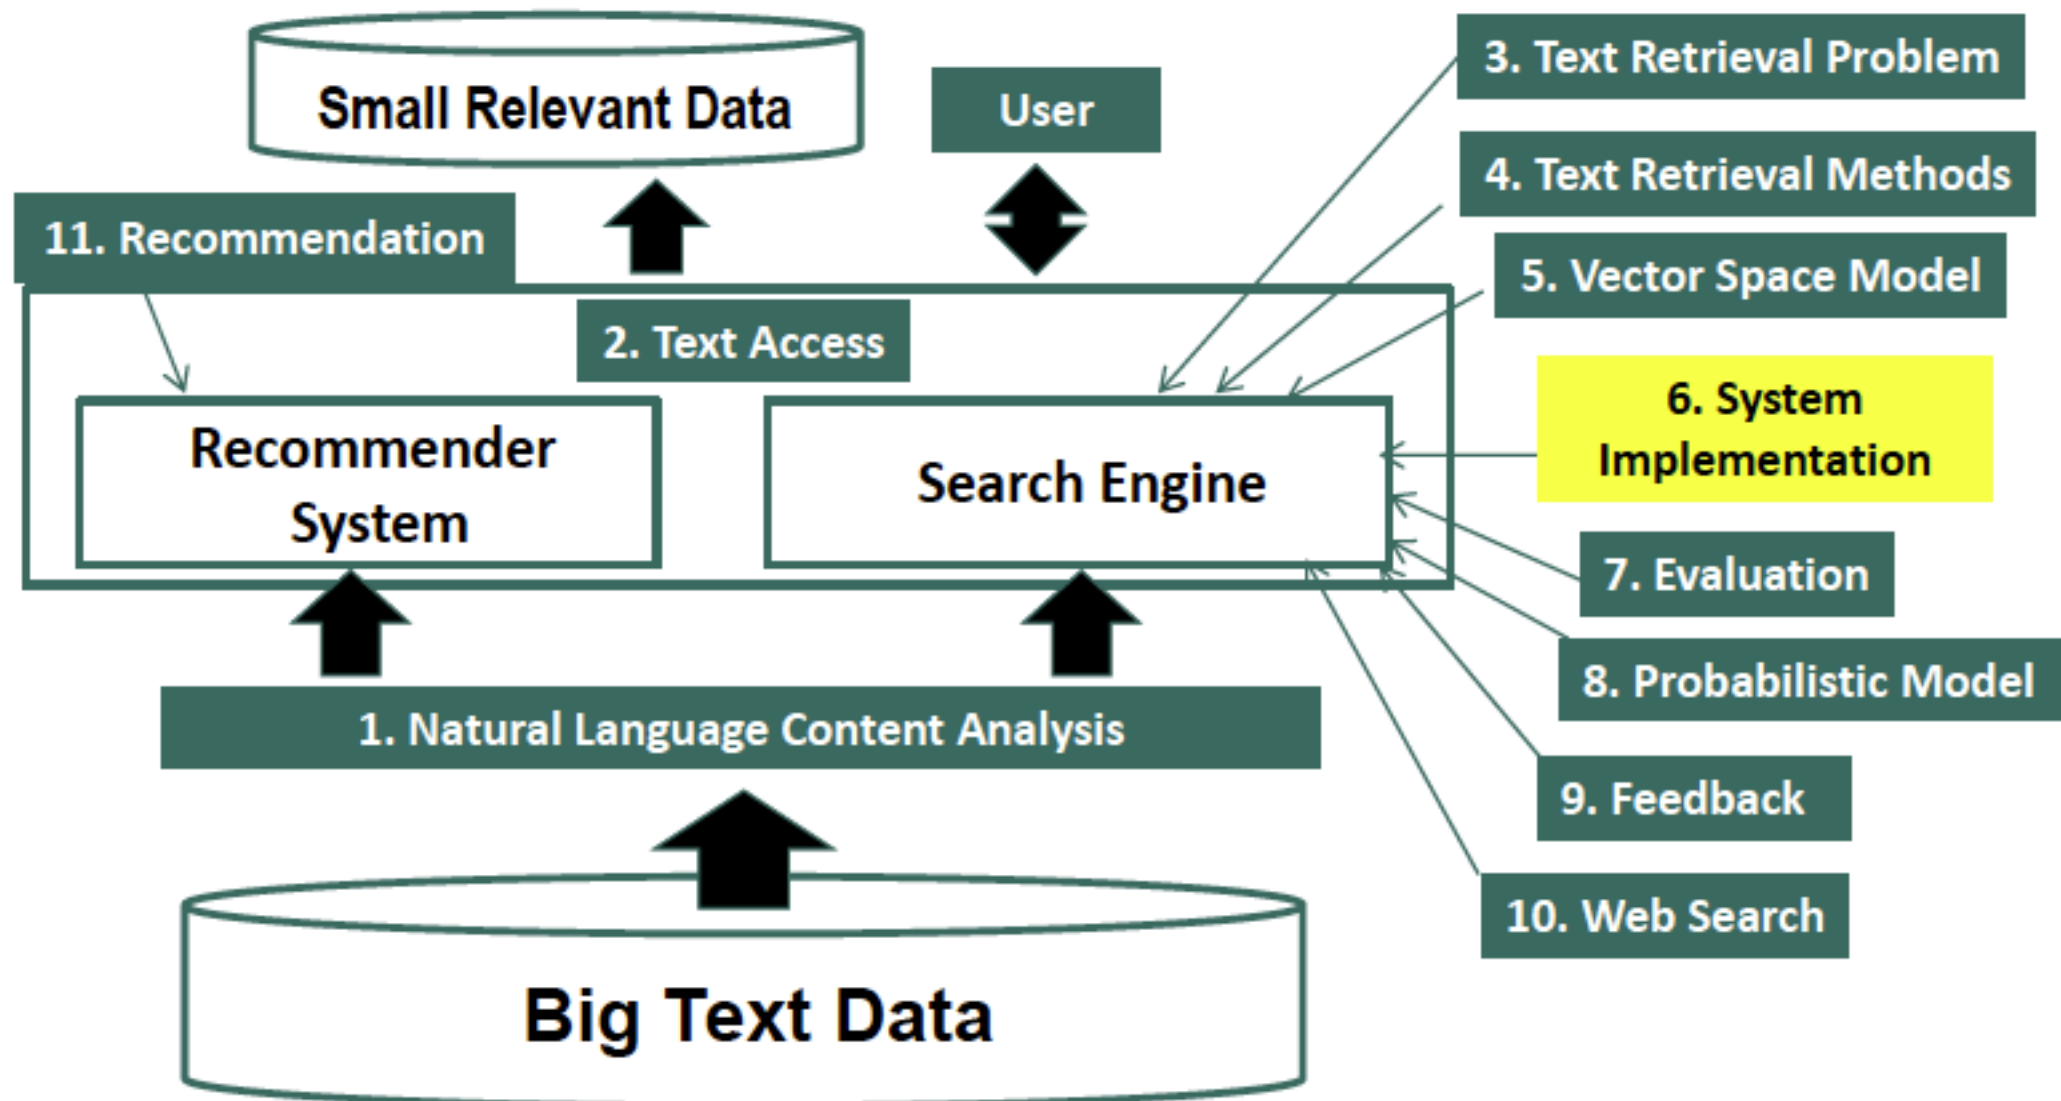
- In general, leverage skewed distribution of values and use variable-length encoding
- TF compression
 - Small numbers tend to occur far more frequently than large numbers (why?)
 - Fewer bits for small (high frequency) integers at the cost of more bits for large integers
- Doc ID compression
 - “d-gap” (store difference): $d_1, d_2-d_1, d_3-d_2, \dots$
 - Feasible due to sequential access

Information Retrieval & Text Mining

System Implementation
Fast Search

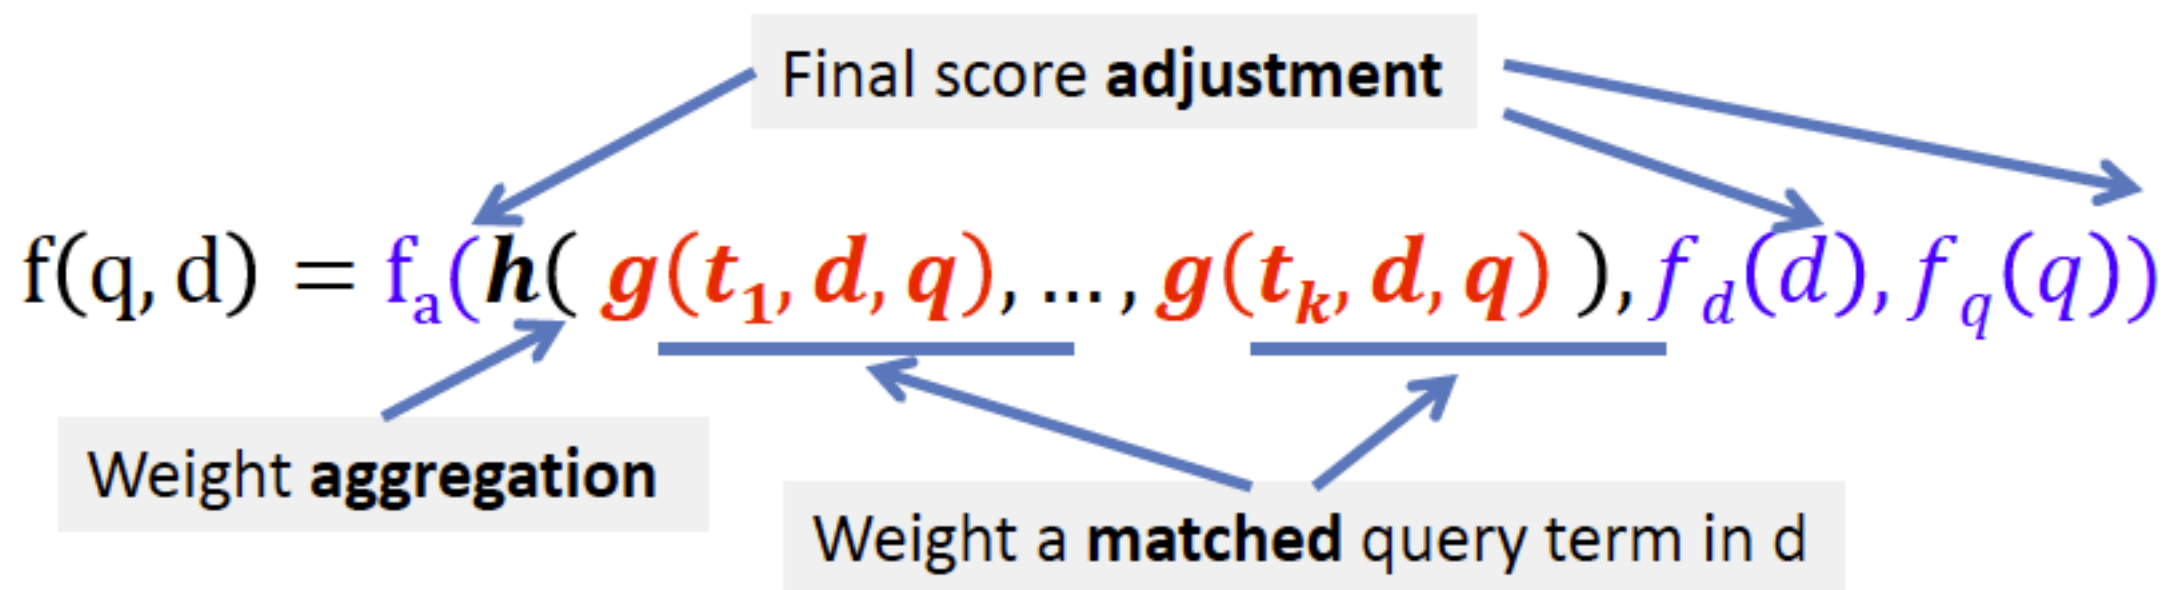
Dr. Saeed Ul Hassan
Information Technology University

Implementation of Text Retrieval Systems



How to Score Documents Quickly

General Form of Scoring Function



An Example: Ranking Based on TF Sum

$$f(d,q)=g(t_1,d,q)+\dots+g(t_k,d,q)$$

$$\text{where } g(t_i,d,q) = c(t_i,d)$$

Query = “info security”

Info: (d1, 3), (d2, 4), (d3, 1), (d4, 5)

Security: (d2, 3), (d4, 1), (d5, 3)

Accumulators:	d1	d2	d3	d4	d5
	0	0	0	0	0

An Example: Ranking Based on TF Sum

$$f(d,q)=g(t_1,d,q)+\dots+g(t_k,d,q)$$

$$\text{where } g(t_i,d,q) = c(t_i,d)$$

Query = “info security”

Info: (d1, 3), (d2, 4), (d3, 1), (d4, 5)

Security: (d2, 3), (d4, 1), (d5, 3)

Accumulators:	d1	d2	d3	d4	d5
	0	0	0	0	0
(d1,3) =>	3	0	0	0	0

An Example: Ranking Based on TF Sum

$$f(d,q)=g(t_1,d,q)+\dots+g(t_k,d,q)$$

$$\text{where } g(t_i,d,q) = c(t_i,d)$$

Query = “info security”

Info: (d1, 3), (d2, 4), (d3, 1), (d4, 5)

Security: (d2, 3), (d4, 1), (d5, 3)

Accumulators:	d1	d2	d3	d4	d5
	0	0	0	0	0
(d1,3) =>	3	0	0	0	0
(d2,4) =>	3	4	0	0	0

An Example: Ranking Based on TF Sum

$$f(d,q)=g(t_1,d,q)+\dots+g(t_k,d,q)$$

$$\text{where } g(t_i,d,q) = c(t_i,d)$$

Query = “info security”

Info: (d1, 3), (d2, 4), (d3, 1), (d4, 5)

Security: (d2, 3), (d4, 1), (d5, 3)

Accumulators:	d1	d2	d3	d4	d5
	0	0	0	0	0
(d1,3) =>	3	0	0	0	0
(d2,4) =>	3	4	0	0	0
(d3,1) =>	3	4	1	0	0

An Example: Ranking Based on TF Sum

$$f(d,q)=g(t_1,d,q)+\dots+g(t_k,d,q)$$

$$\text{where } g(t_i,d,q) = c(t_i,d)$$

Query = “info security”

Info: (d1, 3), (d2, 4), (d3, 1), (d4, 5)

Security: (d2, 3), (d4, 1), (d5, 3)

Accumulators:	d1	d2	d3	d4	d5
	0	0	0	0	0
(d1,3) =>	3	0	0	0	0
(d2,4) =>	3	4	0	0	0
(d3,1) =>	3	4	1	0	0
(d4,5) =>	3	4	1	5	0

An Example: Ranking Based on TF Sum

$$f(d,q)=g(t_1,d,q)+\dots+g(t_k,d,q)$$

$$\text{where } g(t_i,d,q) = c(t_i,d)$$

Query = “info security”

Info: (d1, 3), (d2, 4), (d3, 1), (d4, 5)

Security: (d2, 3), (d4, 1), (d5, 3)

Accumulators:	d1	d2	d3	d4	d5
	0	0	0	0	0
(d1,3) =>	3	0	0	0	0
(d2,4) =>	3	4	0	0	0
(d3,1) =>	3	4	1	0	0
(d4,5) =>	3	4	1	5	0
(d2,3) =>	3	7	1	5	0

An Example: Ranking Based on TF Sum

$$f(d,q)=g(t_1,d,q)+\dots+g(t_k,d,q)$$

$$\text{where } g(t_i,d,q) = c(t_i,d)$$

Query = “info security”

Info: (d1, 3), (d2, 4), (d3, 1), (d4, 5)

Security: (d2, 3), (d4, 1), (d5, 3)

Accumulators:	d1	d2	d3	d4	d5
	0	0	0	0	0
(d1,3) =>	3	0	0	0	0
(d2,4) =>	3	4	0	0	0
(d3,1) =>	3	4	1	0	0
(d4,5) =>	3	4	1	5	0
(d2,3) =>	3	7	1	5	0
(d4,1) =>	3	7	1	6	0

An Example: Ranking Based on TF Sum

$$f(d,q)=g(t_1,d,q)+\dots+g(t_k,d,q)$$

$$\text{where } g(t_i,d,q) = c(t_i,d)$$

Query = “info security”

Info: (d1, 3), (d2, 4), (d3, 1), (d4, 5)

Security: (d2, 3), (d4, 1), (d5, 3)

Accumulators:	d1	d2	d3	d4	d5
	0	0	0	0	0
(d1,3) =>	3	0	0	0	0
(d2,4) =>	3	4	0	0	0
(d3,1) =>	3	4	1	0	0
(d4,5) =>	3	4	1	5	0
(d2,3) =>	3	7	1	5	0
(d4,1) =>	3	7	1	6	0
(d5,3) =>	3	7	1	6	3

An Example: Ranking Based on TF Sum

$$f(d,q)=g(t_1,d,q)+\dots+ g(t_k,d,q)$$

$$\text{where } g(t_i,d,q) = c(t_i,d)$$

Query = “info security”

Info: (d1, 3), (d2, 4), (d3, 1), (d4, 5)

Security: (d2, 3), (d4, 1), (d5, 3)

Accumulators:		d1	d2	d3	d4	d5
		0	0	0	0	0
info	(d1,3) =>	3	0	0	0	0
	(d2,4) =>	3	4	0	0	0
	(d3,1) =>	3	4	1	0	0
	(d4,5) =>	3	4	1	5	0
security	(d2,3) =>	3	7	1	5	0
	(d4,1) =>	3	7	1	6	0
	(d5,3) =>	3	7	1	6	3

Further Improving Efficiency

- Caching (e.g., query results, list of inverted index)
- Keep only the most promising accumulators
- Scaling up to the Web-scale? (need parallel processing)

Some Text Retrieval Toolkits

- Lucene: <http://lucene.apache.org/>
- Lemur/Indri: <http://www.lemurproject.org/>
- Terrier: <http://terrier.org/>
- MeTA: <http://meta-toolkit.github.io/meta/>
- More can be found at <http://timan.cs.uiuc.edu/resources>

Summary of System Implementation

- Inverted index and its construction
 - Preprocess data as much as we can
 - Compression when appropriate
- Fast search using inverted index
 - Exploit inverted index to accumulate scores for documents matching a query term
 - Exploit Zipf's law to avoid touching many documents not matching any query term
 - Can support a wide range of ranking algorithms
- Great potential for further scaling up using distributed file system, parallel processing, and caching

Additional Readings

- Ian H. Witten, Alistair Moffat, Timothy C. Bell: Managing Gigabytes: Compressing and Indexing Documents and Images, Second Edition. Morgan Kaufmann, 1999.
- Stefan Büttcher, Charles L. A. Clarke, Gordon V. Cormack: Information Retrieval - Implementing and Evaluating Search Engines. MIT Press, 2010.