

Lucene Tutorial

Open Source IR Library

Overview

- Open source IR systems
- Lucene Intro
- Installation
- Lucene core classes
- Scoring
- Index format
- Useful resources.

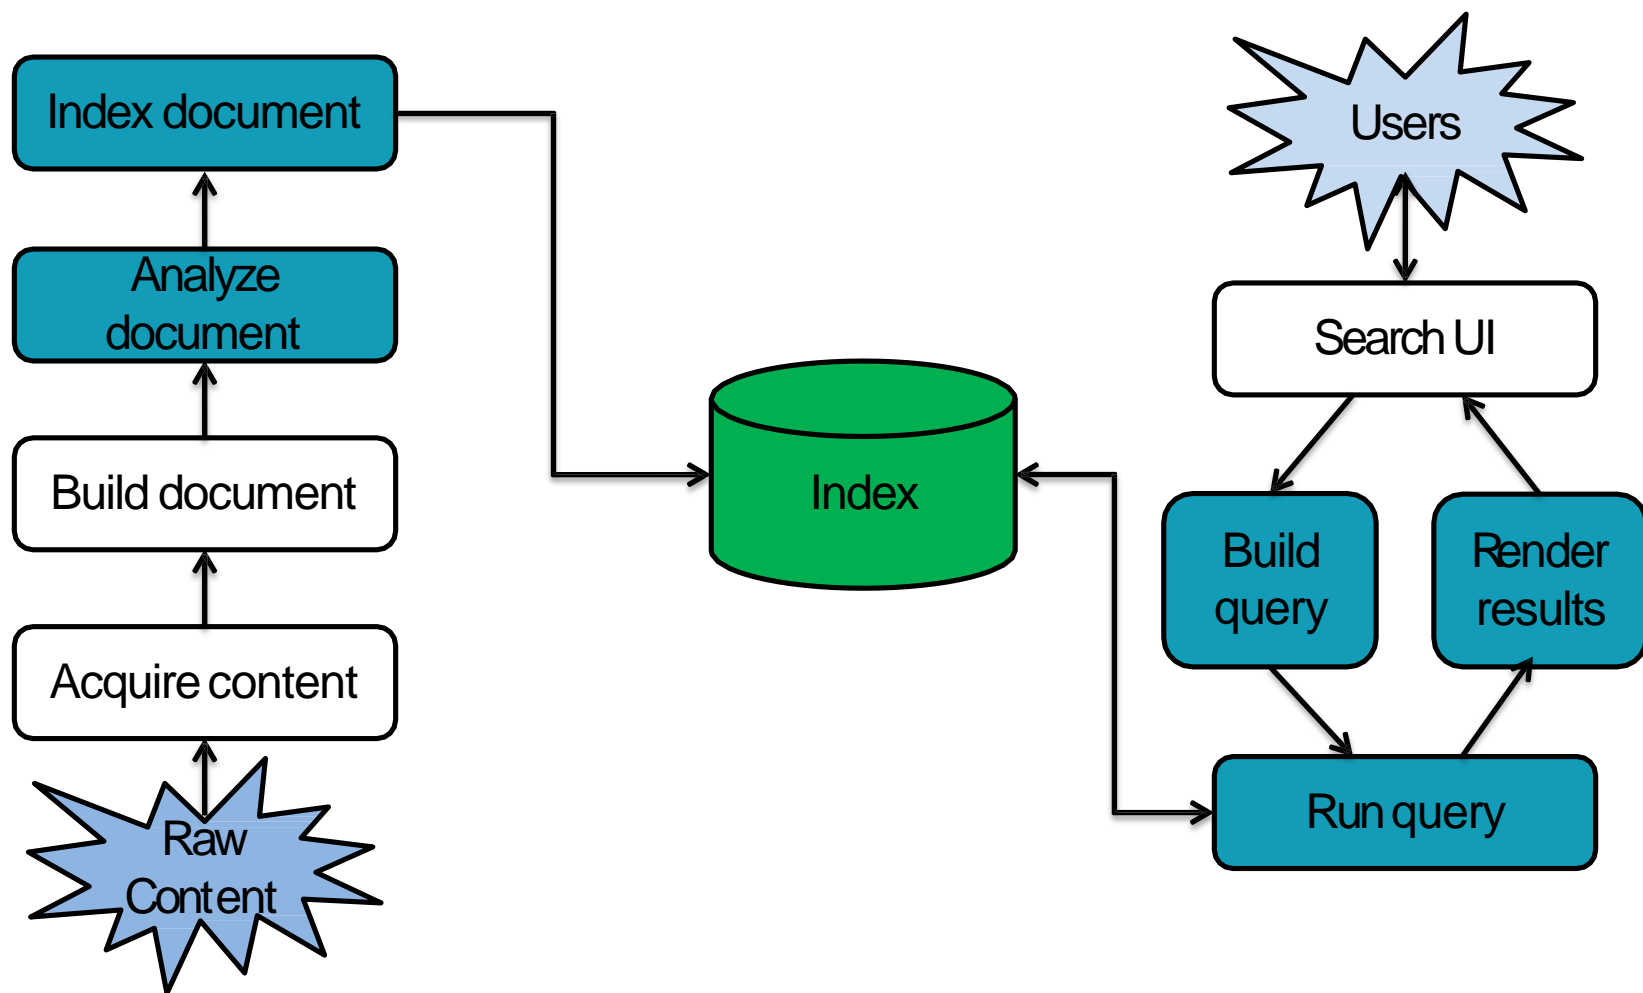
Open source IR systems

- IR systems to reduce information overload.
- Widely used academic systems
 - Terrier (Java, U. Glasgow) <http://terrier.org>
 - Indri/Galago/Lemur (C++ (& Java), U. Mass & CMU)
 - Zettair (RMIT U.)..
- Widely used non-academic open source systems
 - **Lucene**
 - Things built on it: Solr, ElasticSearch
 - A few others (Xapian, ...)

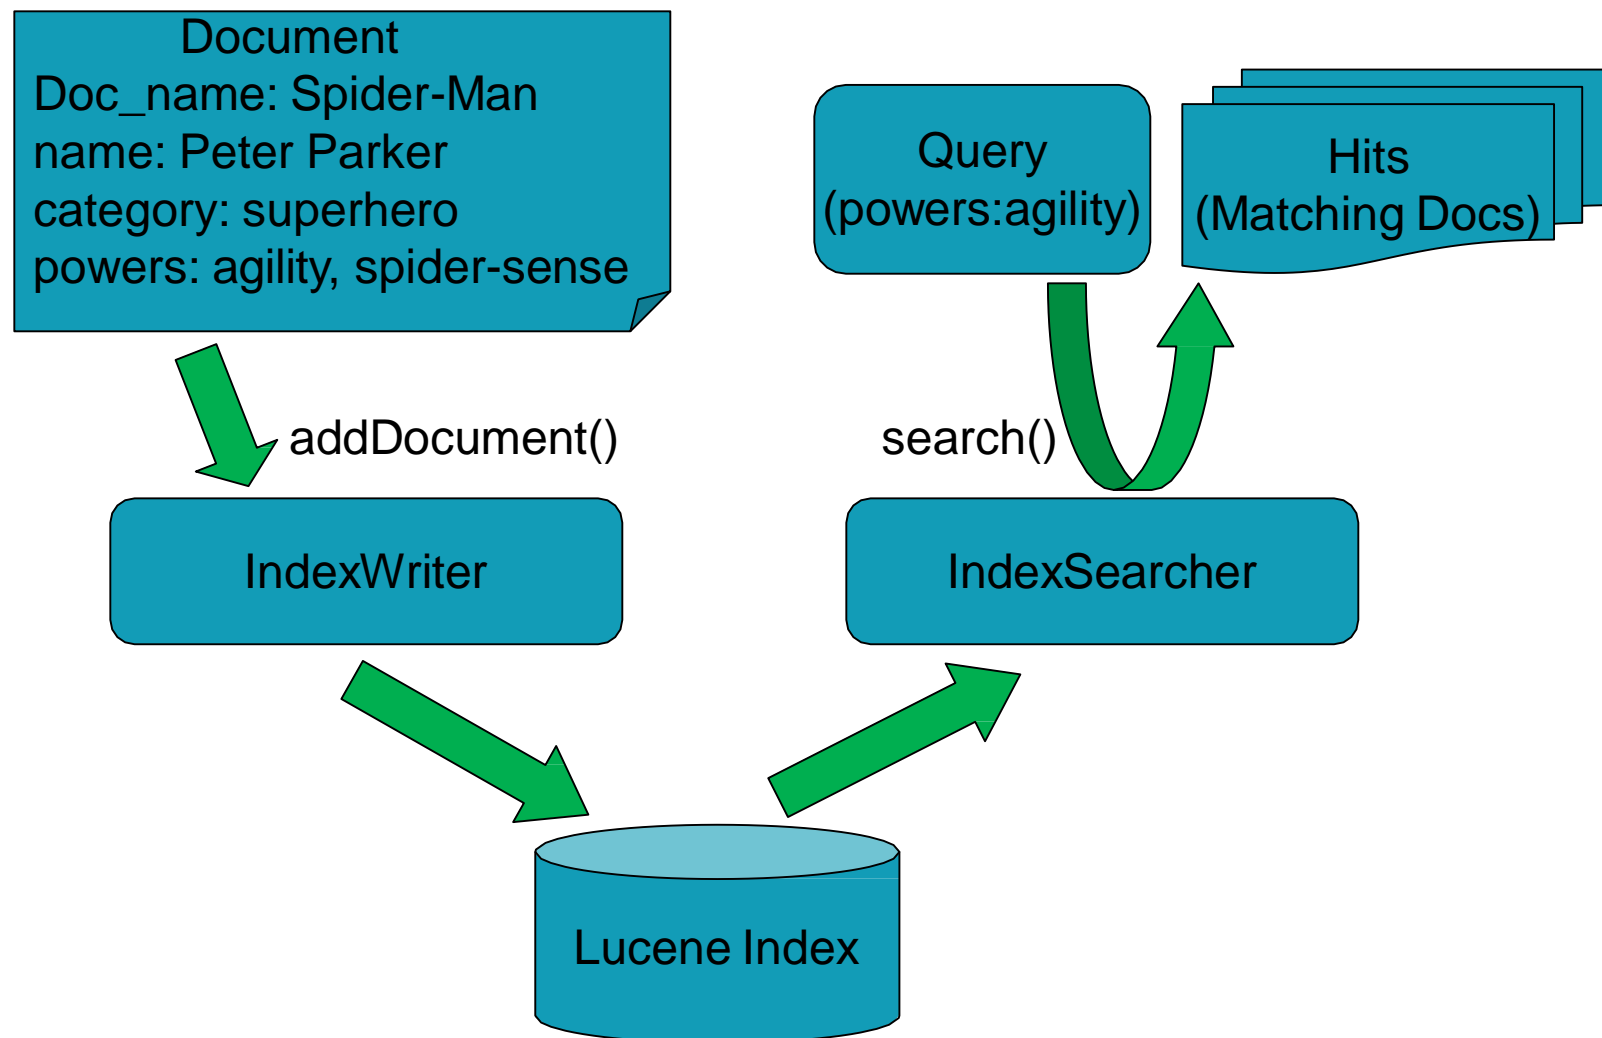
Lucene

- Open source Java library for indexing and searching
 - Lets you add search to your application
 - Not a complete search system by itself
 - Written by Doug Cutting in 1990
- Used by: Twitter, LinkedIn, Zappos, CiteSeer, Eclipse, ...
 - ...and many more (see <http://wiki.apache.org/lucene-java/PoweredBy>)
- Ports/integrations to other languages
 - C/C++, C#, Ruby, Perl, Python (Pylucene), PHP, ...

Basic Application



Lucene in search systems



Installation

- Apache Lucene
<https://archive.apache.org/dist/lucene/java/3.6.2/>
after downloading extract the files to the desktop.
- JDK/JRE
- Eclipse

Add External Jar Files

- lucene-core-3.6.2

Demo Files

- **LuceneConstants**- provide various constants to be used across the sample application.
- **TextFileFilter**- Code is used as a **.txt file** filter.
- **Indexer**- Code to create a Lucene index.
- **Searcher**- Code to search a Lucene Index.
- **LuceneTester**- Code used to test the indexing and search capability of lucene library.

Set directory Paths

- IndexFiles

 - set docsPath= “Documents folder path”

 - set indexPath = “Index folder path”

- SearchFiles

 - set index= “Index folder path”

- Write a query

- Top results

Core indexing classes

- **IndexWriter**
 - Central component that allows you to create a new index, open an existing one, and add, remove, or update documents in an index
 - Built on an IndexWriterConfig and a Directory
- **Directory**
 - Abstract class that represents the location of an index
- **Analyzer**
 - Extracts tokens from a text stream

Creating an IndexWriter

```
Import org.apache.lucene.analysis.Analyzer;  
import org.apache.lucene.index.IndexWriter;  
import org.apache.lucene.index.IndexWriterConfig; import  
org.apache.lucene.store.Directory;  
  
...  
private IndexWriter writer;  
public Indexer(String dir) throws IOException {  
    //this //directory will contain  
    the indexes  
    Directory indexDirectory =  
        FSDirectory.open(new File(indexDirectoryPath));  
  
    //create the indexer  
    writer = new IndexWriter(indexDirectory,  
        new StandardAnalyzer(Version.LUCENE_36), true,  
        IndexWriter.MaxFieldLength.UNLIMITED);  
}
```

Index a Document with IndexWriter

```
private IndexWriter writer;
```

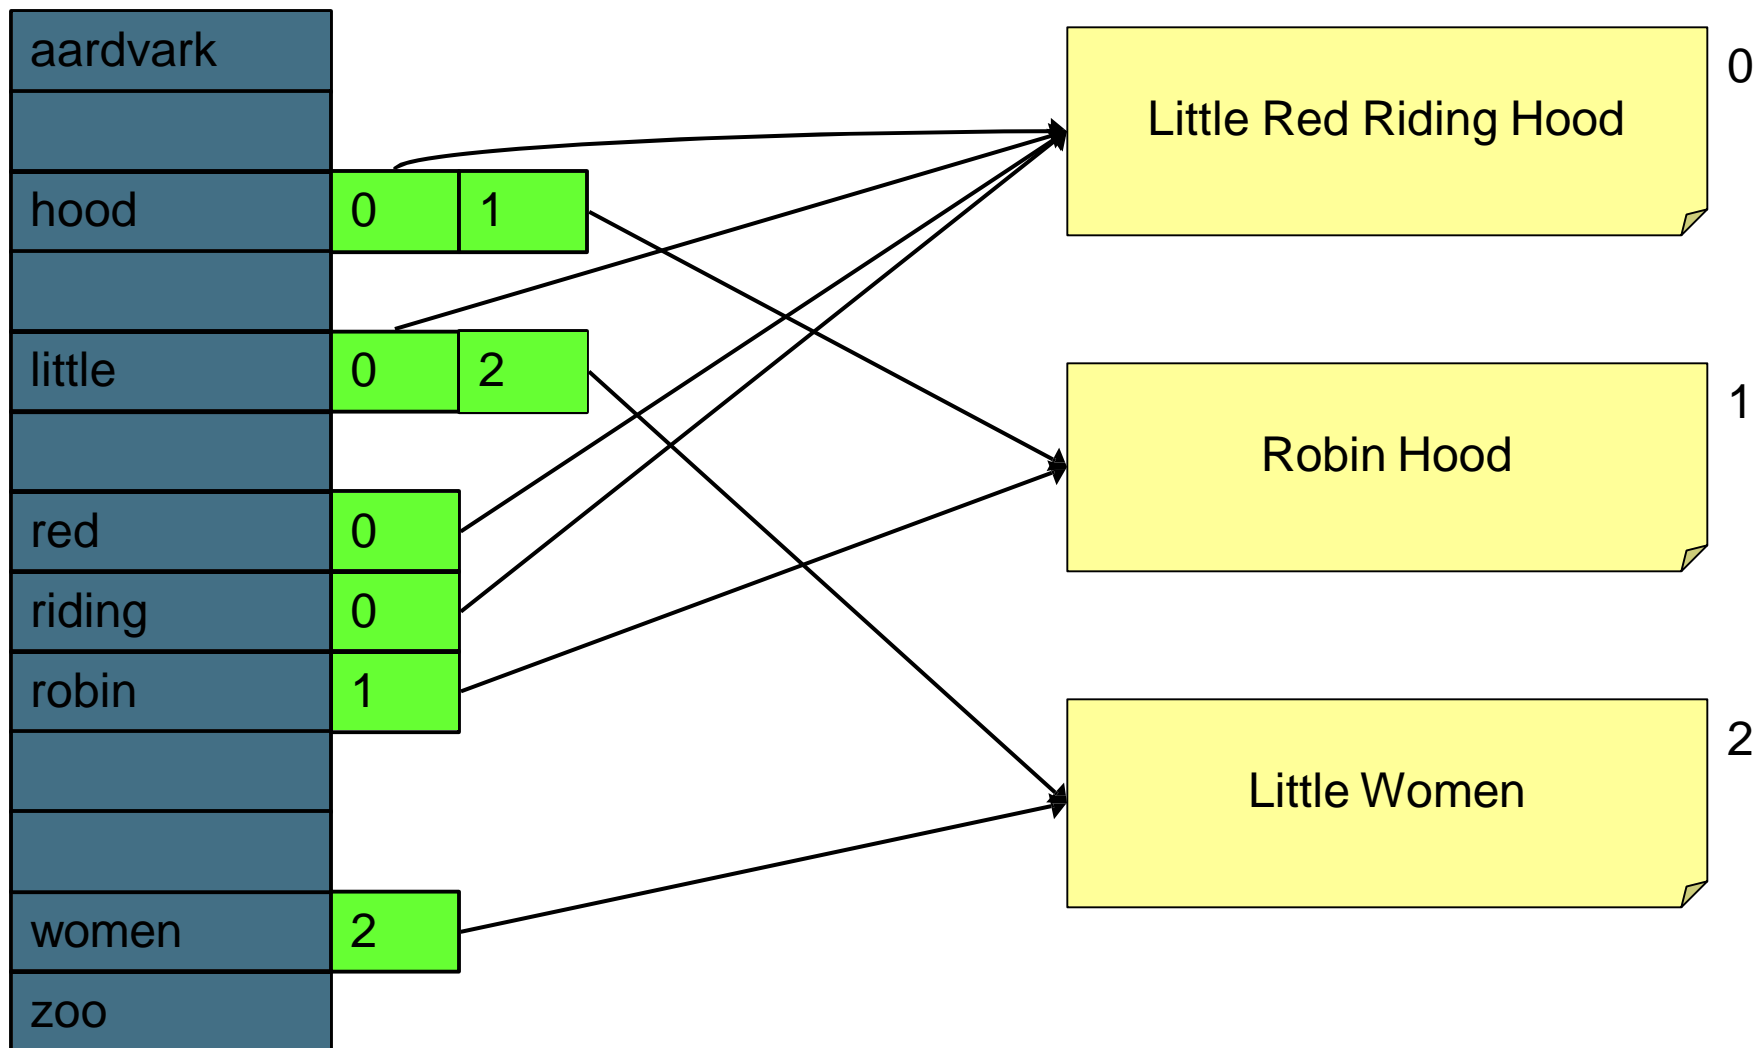
```
...
```

```
private void indexFile(File file) throws IOException {  
    System.out.println("Indexing  
        "+file.getCanonicalPath());  
    Document document = getDocument(file);  
    writer.addDocument(document);  
}
```

The Index

- The Index is the kind of inverted index we know and love
 - natural ordering of docIDs
 - encodes both term frequencies and positional information
- APIs to customize the codec

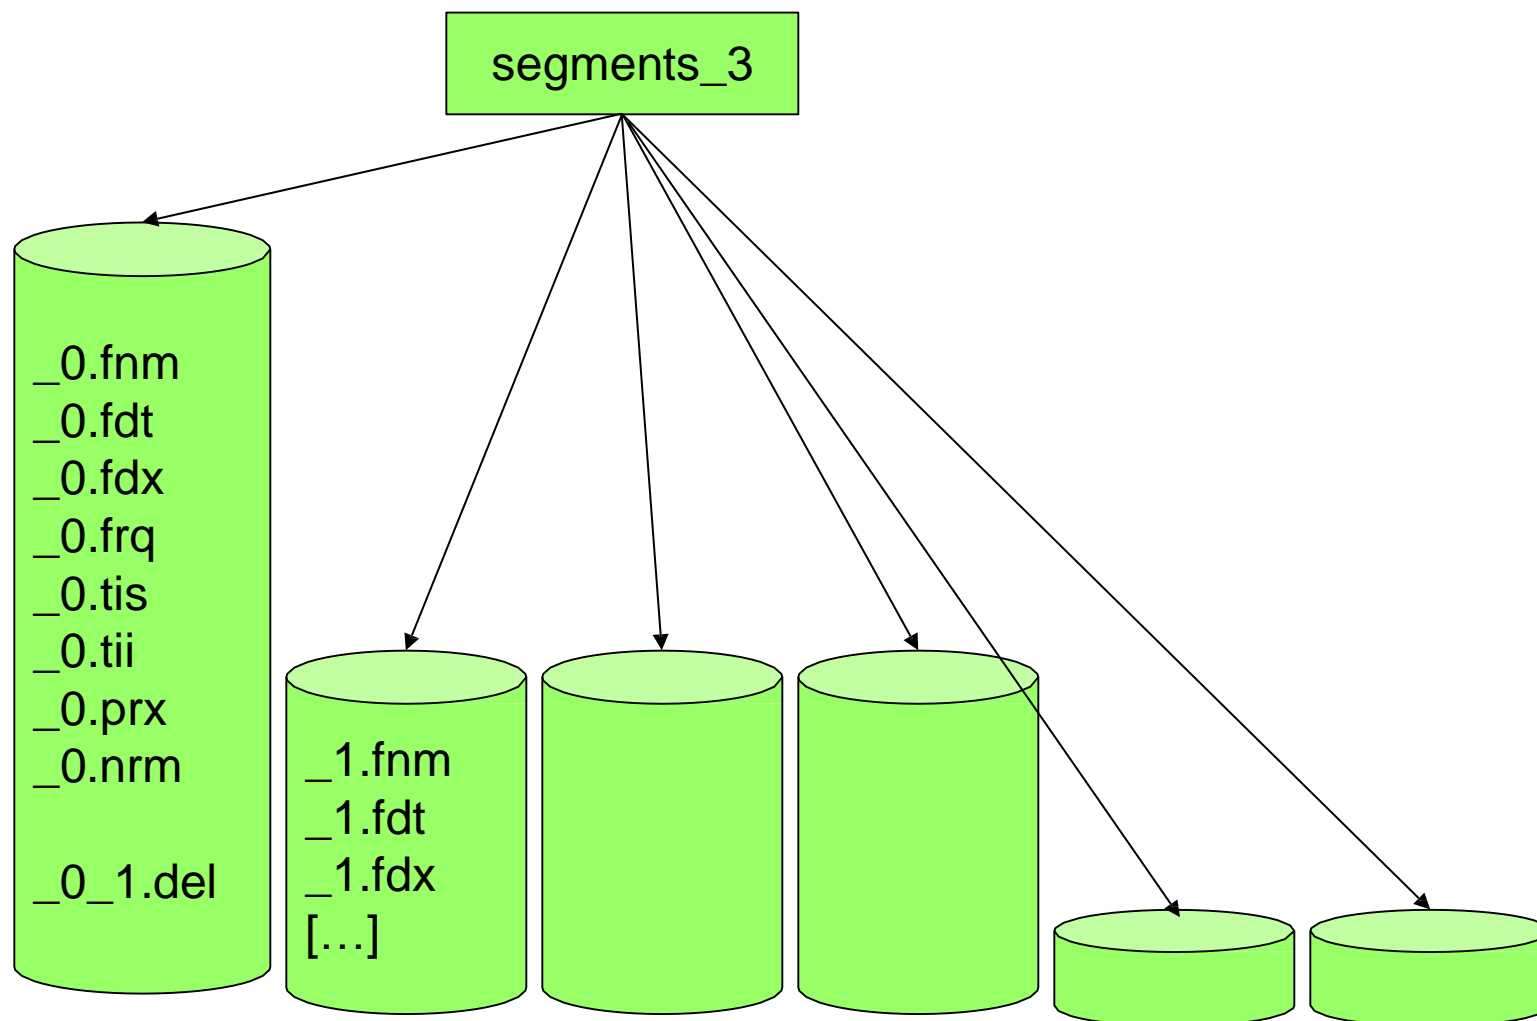
Inverted Index



Index format

- Each Lucene index consists of one or more segments
 - A segment is a standalone index for a subset of documents
 - All segments are searched
 - A segment is created whenever IndexWriter flushes adds/deletes
- Periodically, IndexWriter will merge a set of segments into a single segment
 - Policy specified by a MergePolicy


Index Structure



Core searching classes

- **IndexSearcher**
 - Central class that exposes several search methods on an index
 - Accessed via an IndexReader
- **Query**
 - Abstract query class. Concrete subclasses represent specific types of queries, e.g., matching terms in fields, boolean queries, phrase queries, ...
- **QueryParser**
 - Parses a textual representation of a query into a Query instance

IndexSearcher



Creating an IndexSearcher

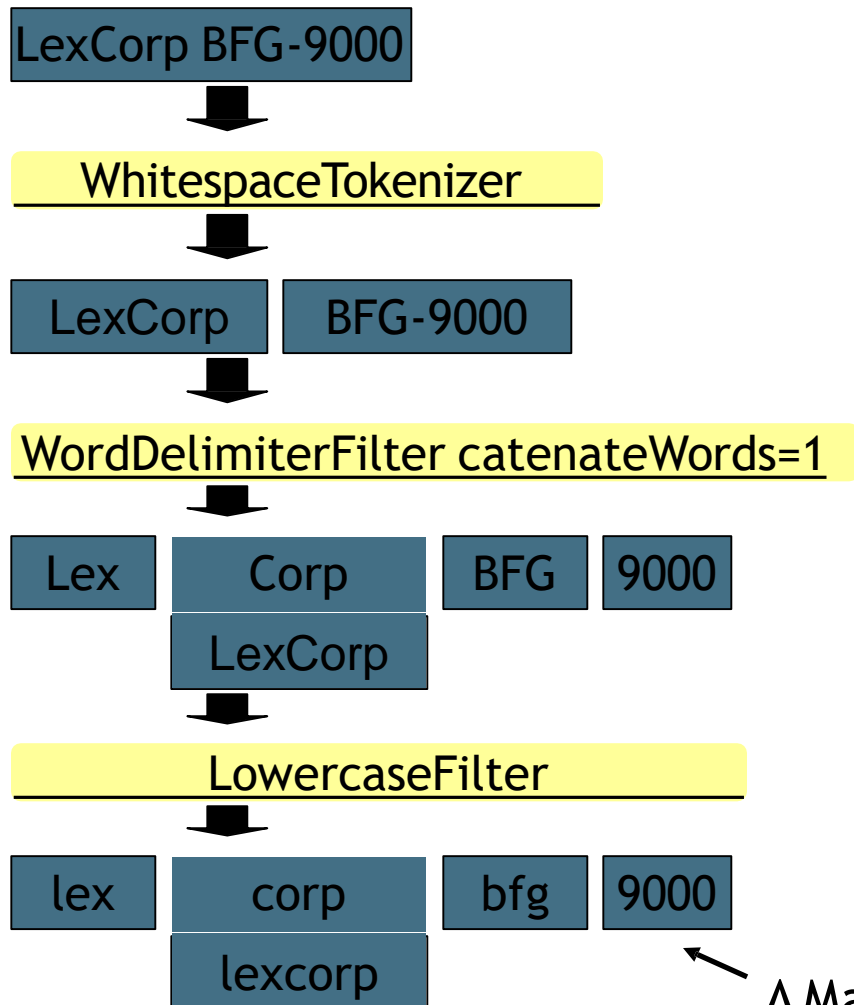
```
import org.apache.lucene.search.IndexSearcher;  
  
...  
public Searcher(String indexDirectoryPath) throws IOException {  
    Directory indexDirectory = FSDirectory.open(new  
    File(indexDirectoryPath));  
    indexSearcher = new IndexSearcher(indexDirectory);  
}
```

Query and QueryParser and search() returns TopDocs

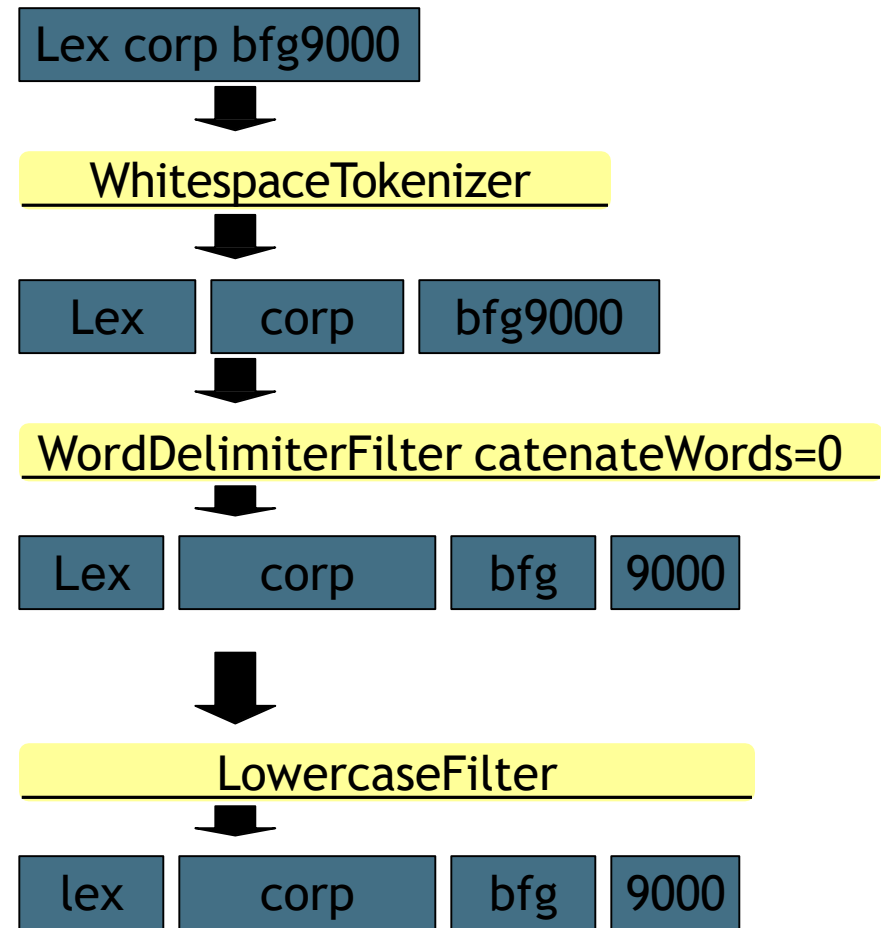
```
import org.apache.lucene.queryParser.QueryParser;  
import org.apache.lucene.search.Query;  
  
public static void search(String indexDir, String q) throws IOException,  
    ParseException {  
    ...  
    queryParser = new  
    QueryParser(Version.LUCENE_36, LuceneConstants.CONTENTS,  
        new StandardAnalyzer(Version.LUCENE_36));  
}  
  
public TopDocs search( String searchQuery) throws IOException,  
    ParseException {  
    query = queryParser.parse(searchQuery);  
    return indexSearcher.search(query, LuceneConstants.MAX_SEARCH);  
}
```

Analysis & Search Relevancy

Document Indexing Analysis



Query Analysis



A Match!

Scoring

- VSM – Vector Space Model
- tf – term frequency: number of matching terms in field
- lengthNorm – number of tokens in field
- idf – inverse document frequency
- coord – coordination factor, number of matching terms
- Boosting techniques

<http://lucene.apache.org/java/docs/scoring.html>

Useful Resources

- https://lucene.apache.org/core/7_0_1/index.html
- https://lucene.apache.org/core/3_5_0/scoring.html
- https://lucene.apache.org/core/2_9_4/fileformats.pdf
- https://lucene.apache.org/core/7_2_1/demo/overview-summary.html
- https://www.youtube.com/watch?v=pVDVURw_AJQ