# Information Retrieval & Text Mining

## Document Clustering

**Dr. Saeed Ul Hassan**
**Information Technology University**

# Today's Topic

- **Document clustering**
  - Why we need it?
  - Meta-data representation
    - Textual, non textual (images, tables, algorithms, etc)
  - Evaluation

- **Clustering algorithms**
  - Partitioning algorithms
  - Hierarchical algorithms

# What is clustering?

- **Clustering**: *The process of grouping a set of objects into classes of similar objects.*
  - Documents within a cluster should be similar.
  - Documents from different clusters should be dissimilar.

- **The commonest form of *unsupervised learning***
    - **Unsupervised learning** = learning from raw data, as opposed to **Supervised learning** = where a classification of examples are available.
  - A common and important task that finds many applications in IR and other domains.

# A data set with clear cluster structure

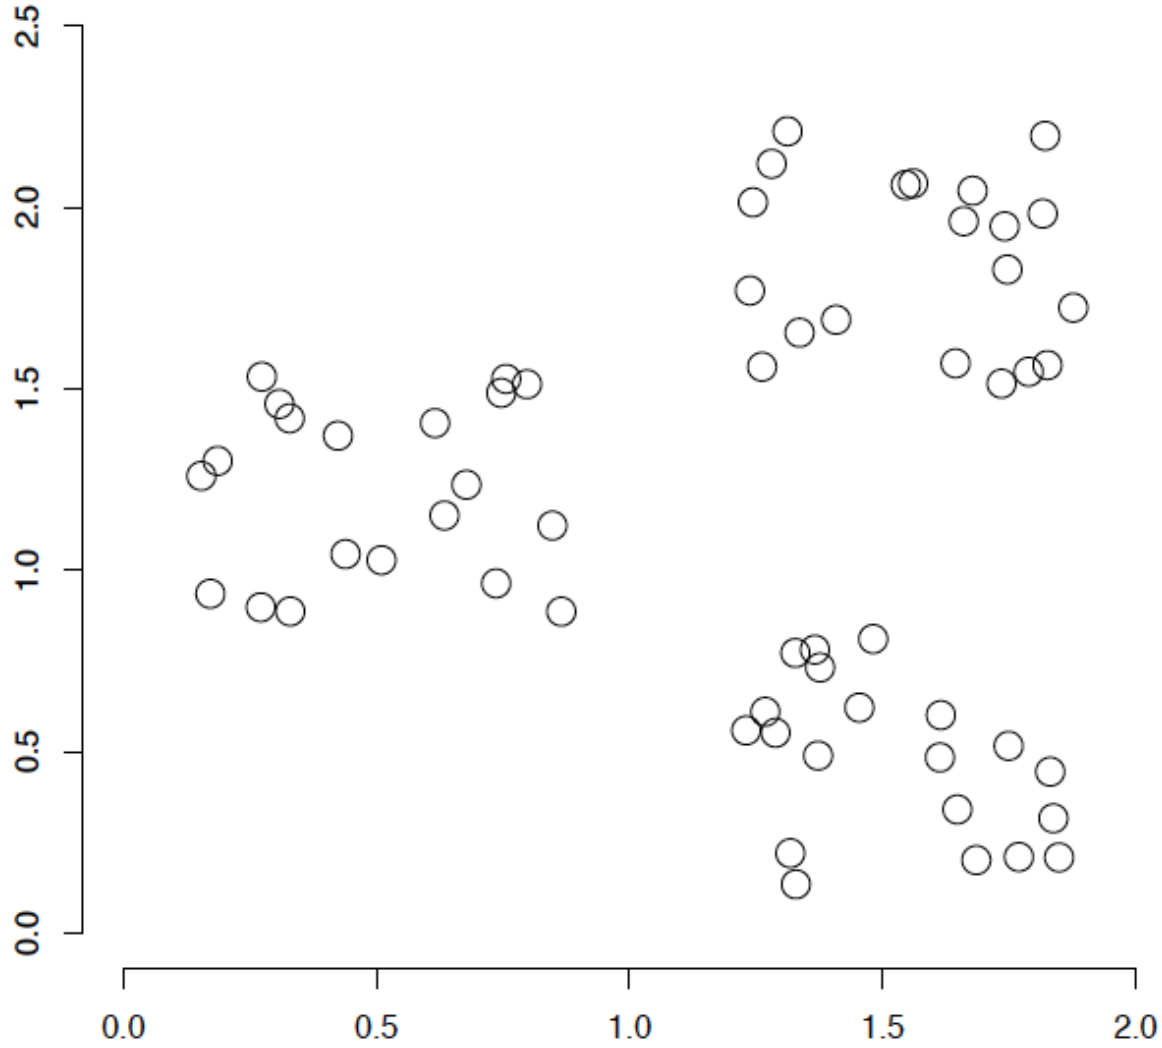


- How would you design an algorithm for finding the three clusters in this case?

# A data set with clear cluster structure

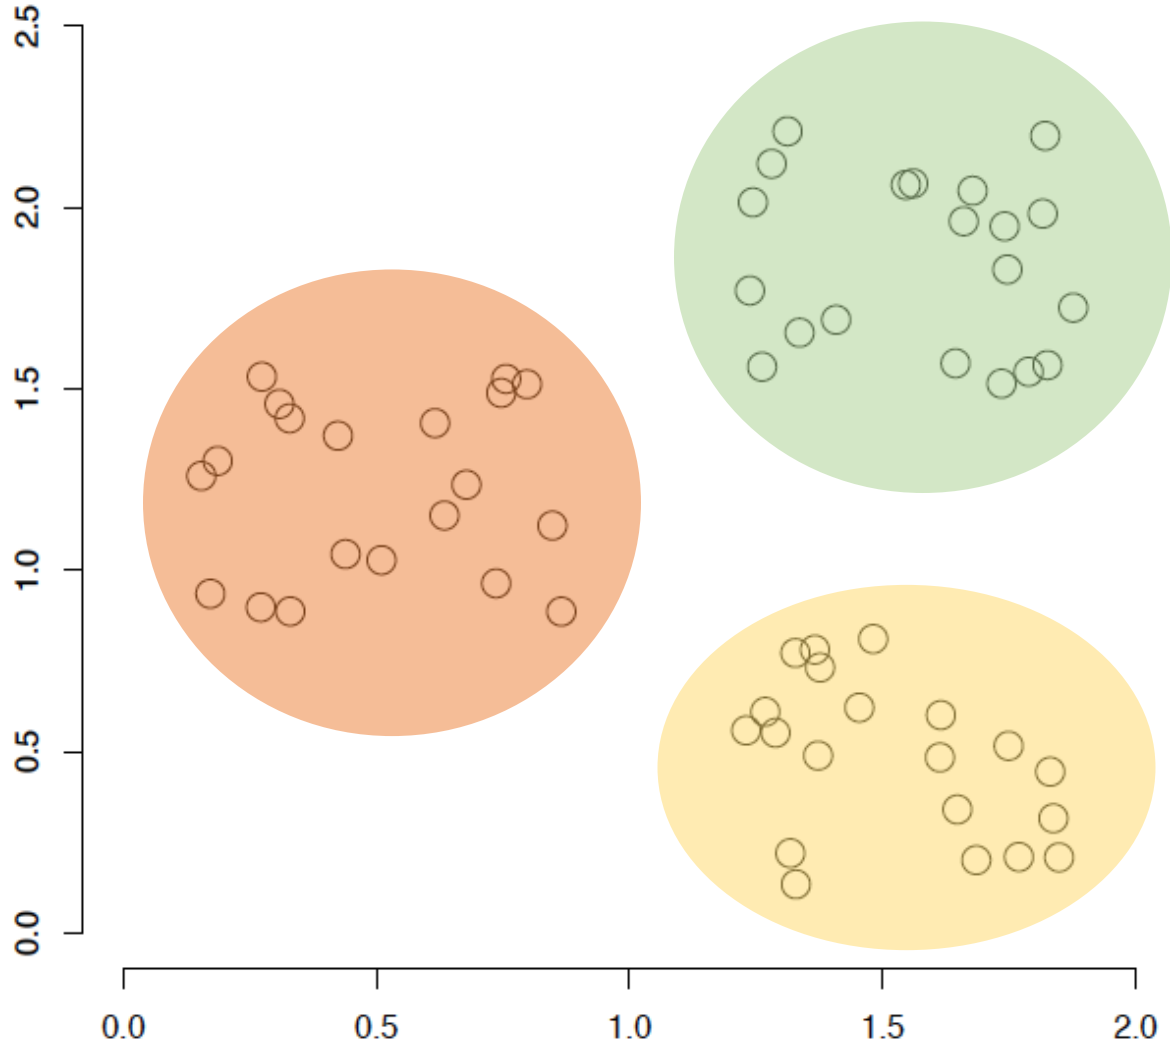


- 3 Clusters

# Applications of clustering in IR

- **Whole corpus analysis/navigation**
  - Better user interface: search without typing
- **For improving recall in search applications**
  - Better search results (like pseudo RF)
- **For better navigation of search results**
  - Effective “user recall” will be higher
- **For speeding up vector space retrieval**
  - Cluster-based retrieval gives faster search

# Hierarchy isn't clustering but is the kind of output you want from clustering

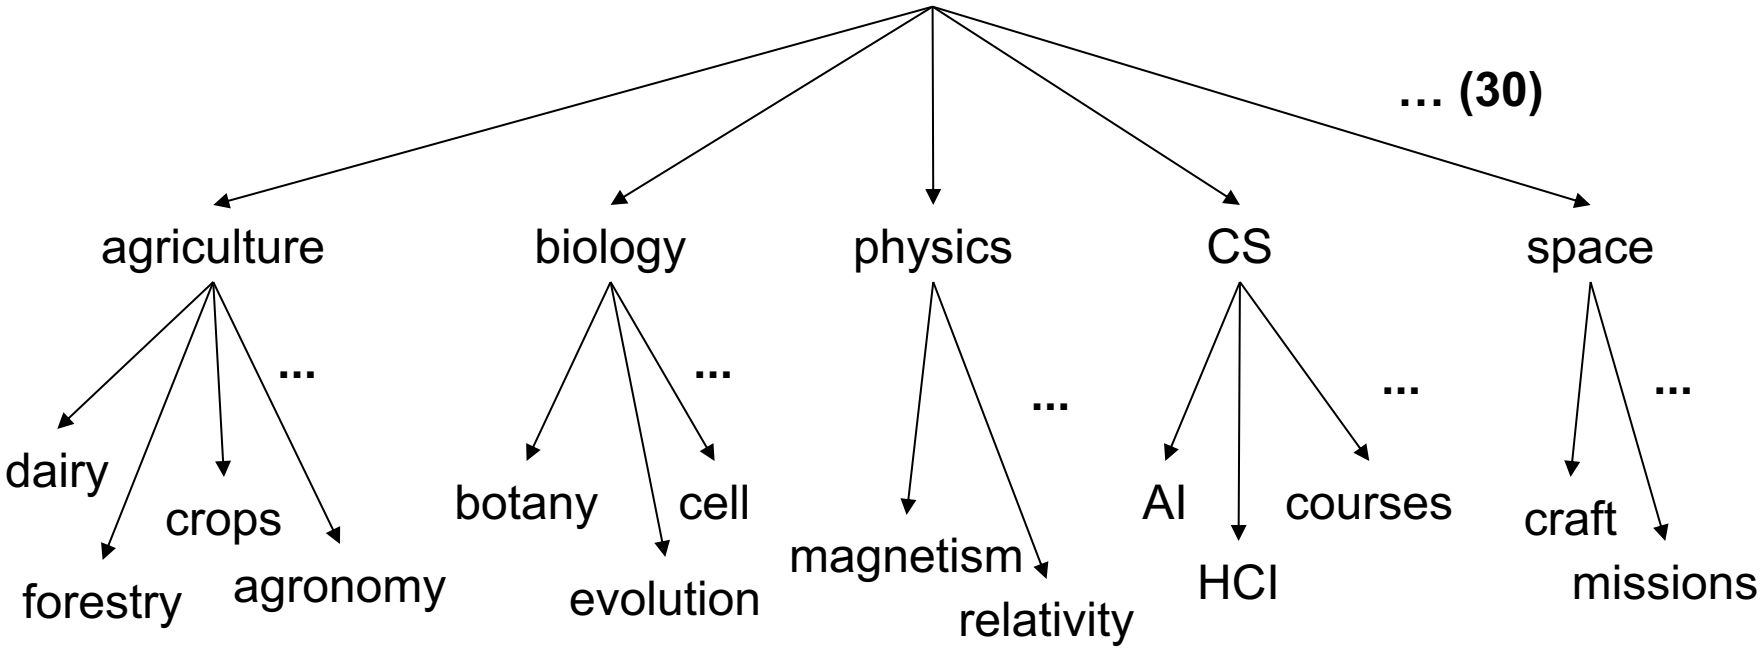


www.yahoo.com/Science

# Applications of clustering in IR

- **Whole corpus analysis/navigation**
  - Better user interface: search without typing
- **For improving recall in search applications**
  - Better search results (like pseudo Relevance Feedback)
- **For better navigation of search results**
  - Effective "user recall" will be higher
- **For speeding up vector space retrieval**
  - Cluster-based retrieval gives faster search

# Applications of clustering in IR

- **Whole corpus analysis/navigation**
  - Better user interface: search without typing
- **For improving recall in search applications**
  - Better search results (like pseudo Relevance Feedback)
- **For better navigation of search results**
  - Improved navigation of high recall
- **For speeding up vector space retrieval**
  - Cluster-based retrieval gives faster search

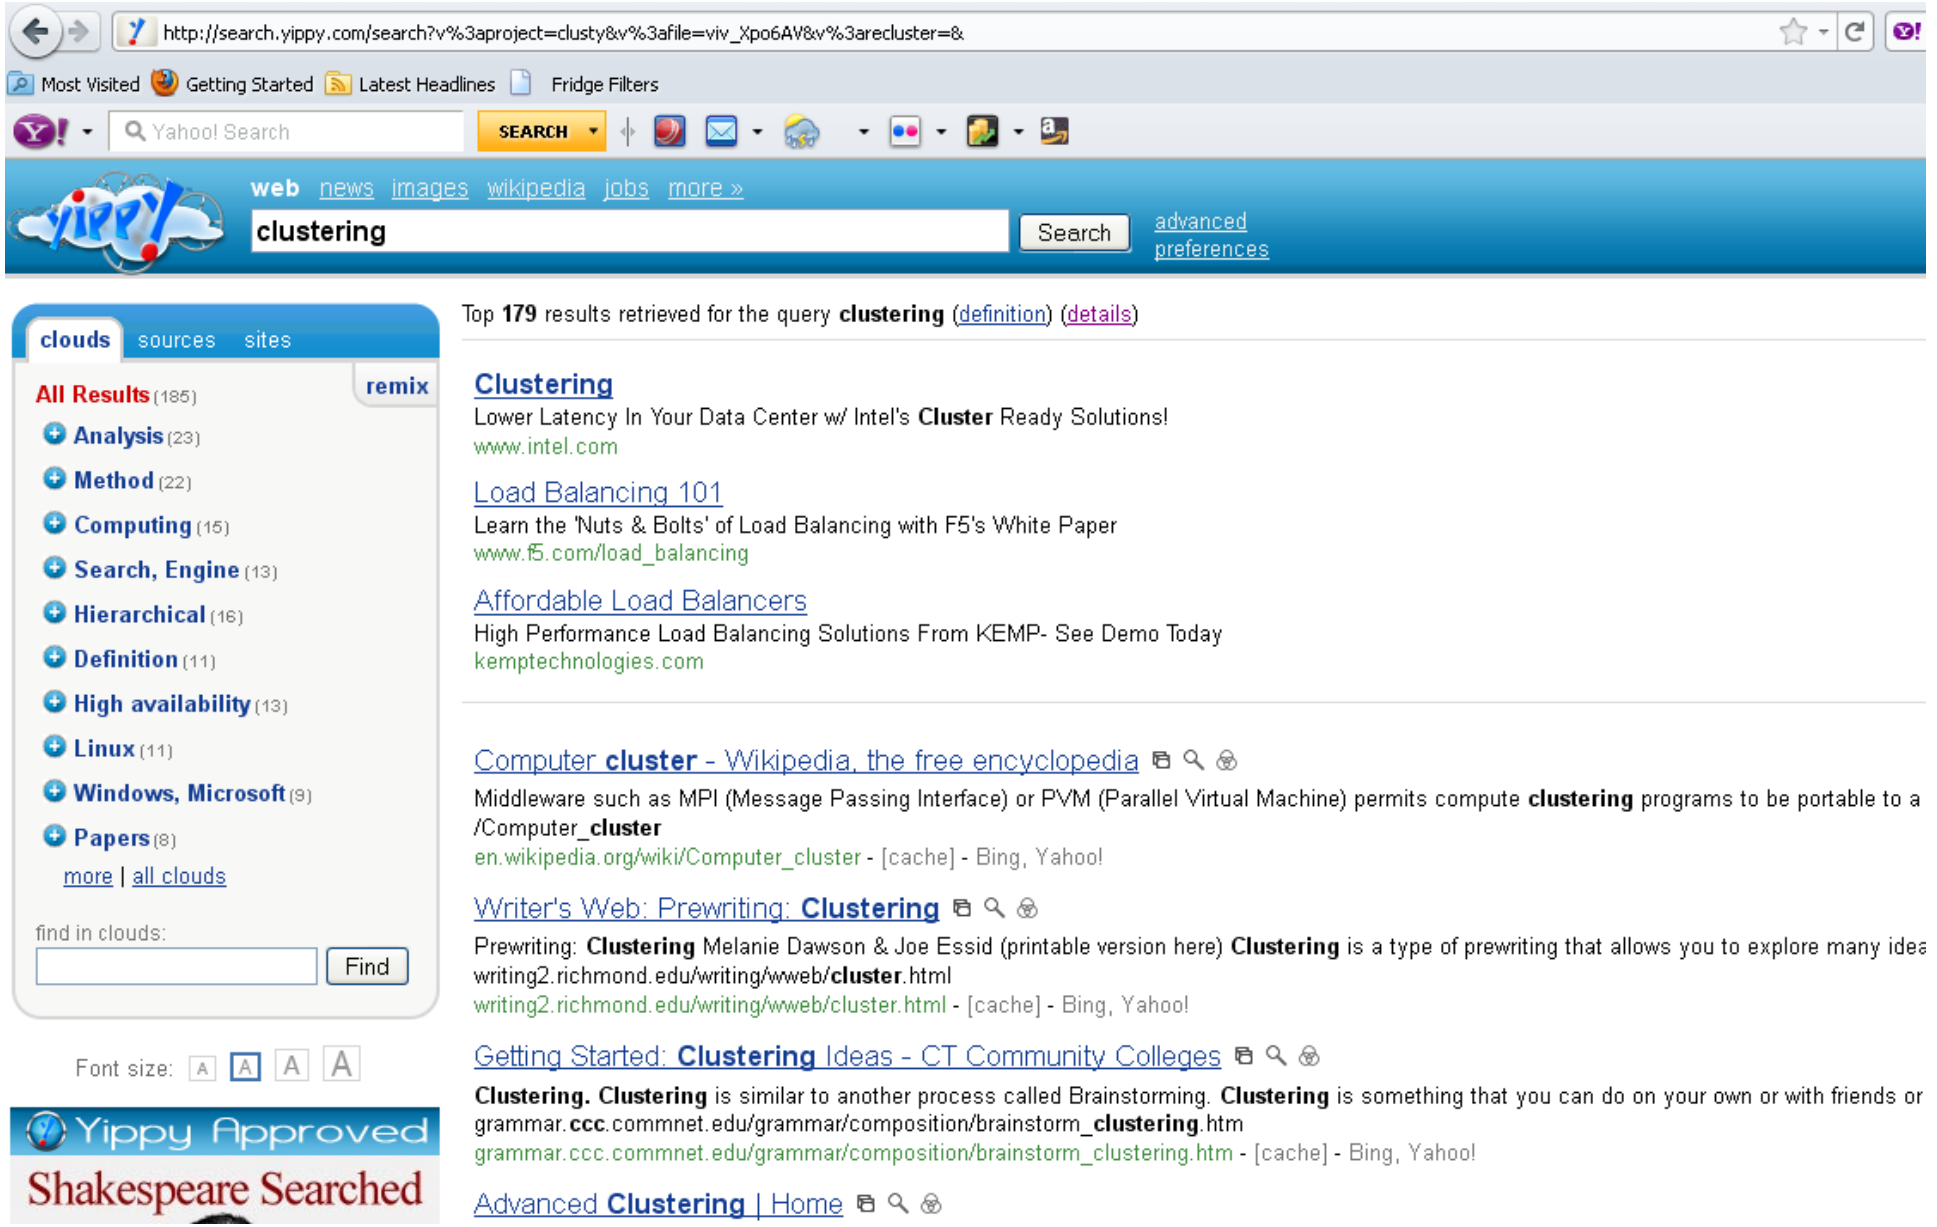
http://search.yippy.com/search?v%3aproject=clusty&v%3afile=viv_Xpo6AV&v%3arecluster=&
Most Visited
Getting Started
Latest Headlines
Fridge Filters
Yahoo! Search
SEARCH
web
news
images
wikipedia
jobs
more »
clustering
Search
advanced
preferences
clouds
sources
sites
remix
All Results (185)
Analysis (23)
Method (22)
Computing (15)
Search, Engine (13)
Hierarchical (16)
Definition (11)
High availability (13)
Linux (11)
Windows, Microsoft (9)
Papers (8)
more | all clouds
find in clouds:
Find
Font size:
Yippy Approved
Shakespeare Searched
Top 179 results retrieved for the query clustering (definition) (details)
Clustering
Lower Latency In Your Data Center w/ Intel's Cluster Ready Solutions!
www.intel.com
Load Balancing 101
Learn the 'Nuts & Bolts' of Load Balancing with F5's White Paper
www.f5.com/load_balancing
Affordable Load Balancers
High Performance Load Balancing Solutions From KEMP- See Demo Today
kemptechnologies.com
Computer cluster - Wikipedia, the free encyclopedia
Middleware such as MPI (Message Passing Interface) or PVM (Parallel Virtual Machine) permits compute clustering programs to be portable to a
/Computer_cluster
en.wikipedia.org/wiki/Computer_cluster - [cache] - Bing, Yahoo!
Writer's Web: Prewriting: Clustering
Prewriting: Clustering Melanie Dawson & Joe Essid (printable version here) Clustering is a type of prewriting that allows you to explore many idea
writing2.richmond.edu/writing/wweb/cluster.html
writing2.richmond.edu/writing/wweb/cluster.html - [cache] - Bing, Yahoo!
Getting Started: Clustering Ideas - CT Community Colleges
Clustering. Clustering is similar to another process called Brainstorming. Clustering is something that you can do on your own or with friends or
grammar.ccc.commnet.edu/grammar/composition/brainstorm_clustering.htm
grammar.ccc.commnet.edu/grammar/composition/brainstorm_clustering.htm - [cache] - Bing, Yahoo!
Advanced Clustering | Home

# Applications of clustering in IR

- **Whole corpus analysis/navigation**
  - Better user interface: search without typing
- **For improving recall in search applications**
  - Better search results (like pseudo Relevance Feedback)
- **For better navigation of search results**
  - Effective “user recall” will be higher
- **For speeding up vector space retrieval**
  - Cluster-based retrieval gives faster search
  - Match query with median of clusters

# Issues for clustering

- **Representation for clustering**
  - Document representation
  - Need a notion of similarity/distance
- **How many clusters?**
  - Number of clusters
  - Avoid too large or small number of clusters

# Hard vs. soft clustering

- **Hard clustering:** Each document belongs to exactly one cluster
  - More common and easier to do
- **Soft clustering:** A document can belong to more than one cluster.
  - Makes more sense for applications like creating browse-able hierarchies
  - You may want to put a pair of sneakers in two clusters: (i) sports apparel and (ii) shoes
  - You can only do that with a soft clustering approach.
- We will only do hard clustering today.

# Clustering Algorithms

- **Partitioning algorithms**
  - Usually start with a random (partial) partitioning
  - Refine it iteratively
    - *K* means clustering
- **Hierarchical algorithms**
  - Bottom-up
  - Top-down

# Partitioning Algorithms

- **Partitioning method**: Construct a partition of *n* documents into a set of *K* clusters
- **Given**: a set of documents and the number *K*
- **Find**: a partition of *K* clusters that optimizes the chosen partitioning criterion
  - Globally optimal
    - Intractable for many objective functions
  - Effective heuristic methods: *K*-means algorithm

# K-Means

- Assumes documents are real-valued vectors e.g. with TF-IDF.
- Clusters based on *centroids* (aka the *center of gravity* or mean) of points in a cluster, *c*:

$$\vec{\mu}(c) = \frac{1}{|c|} \sum_{\vec{x} \in c} \vec{x}$$

- Reassignment of instances to clusters is based on distance to the current cluster centroids.

# K-Means

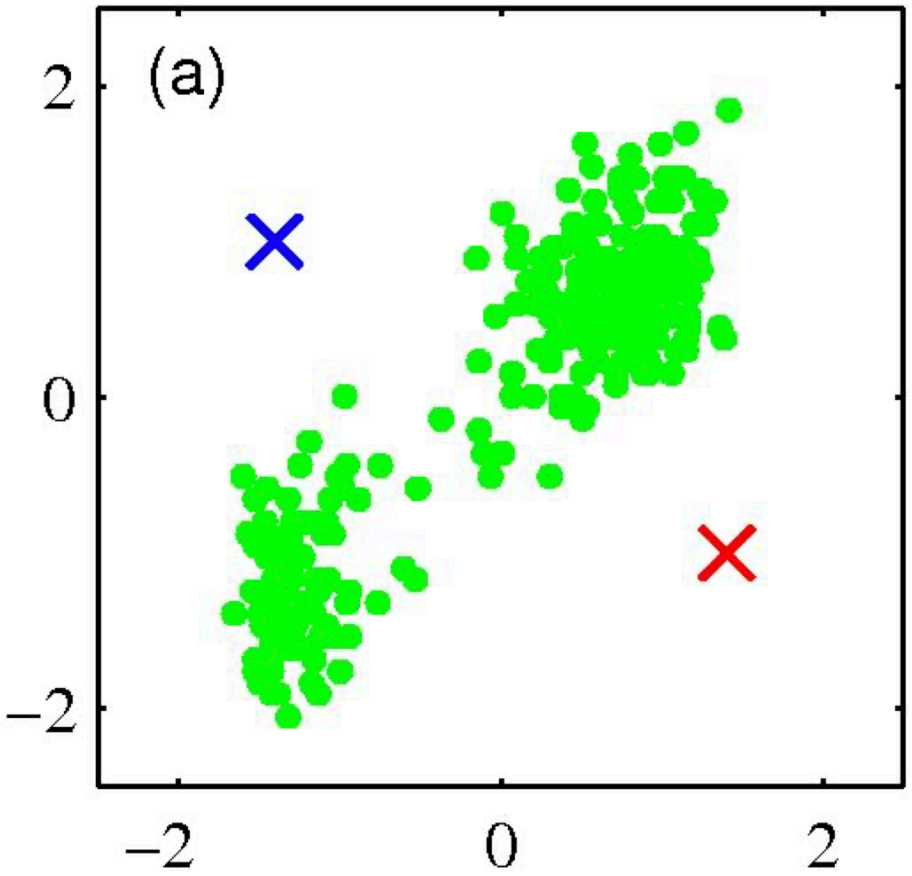


Pick *K* random points as cluster centres (means)

Shown here for *K*=2

# K-Means

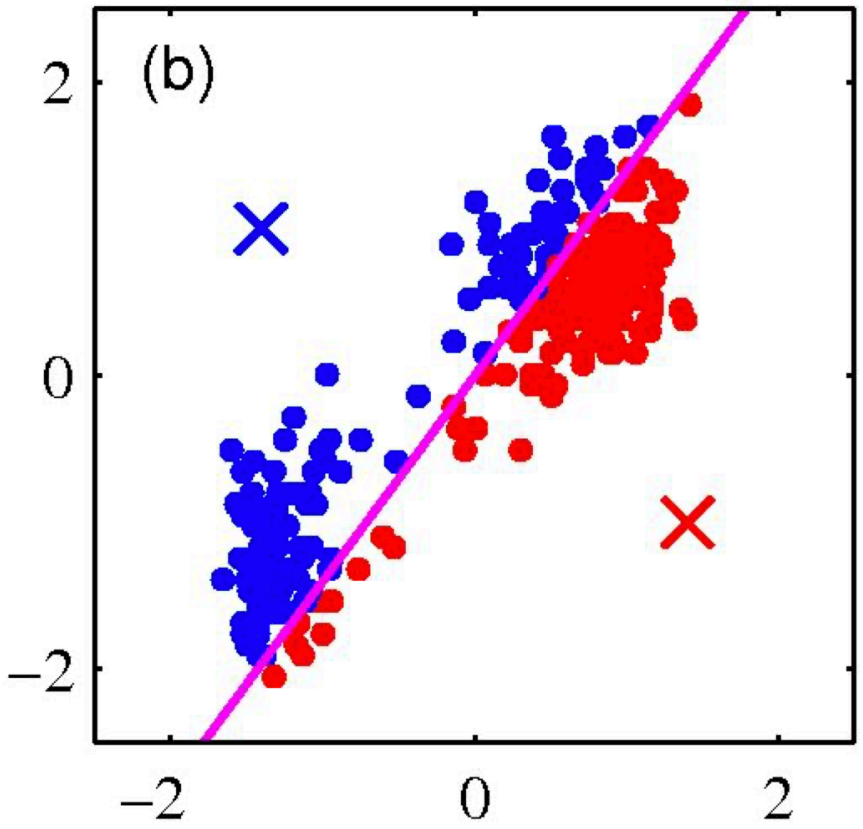


**Iterative Step 1**

• Assign data points to closest cluster canter

# K-Means

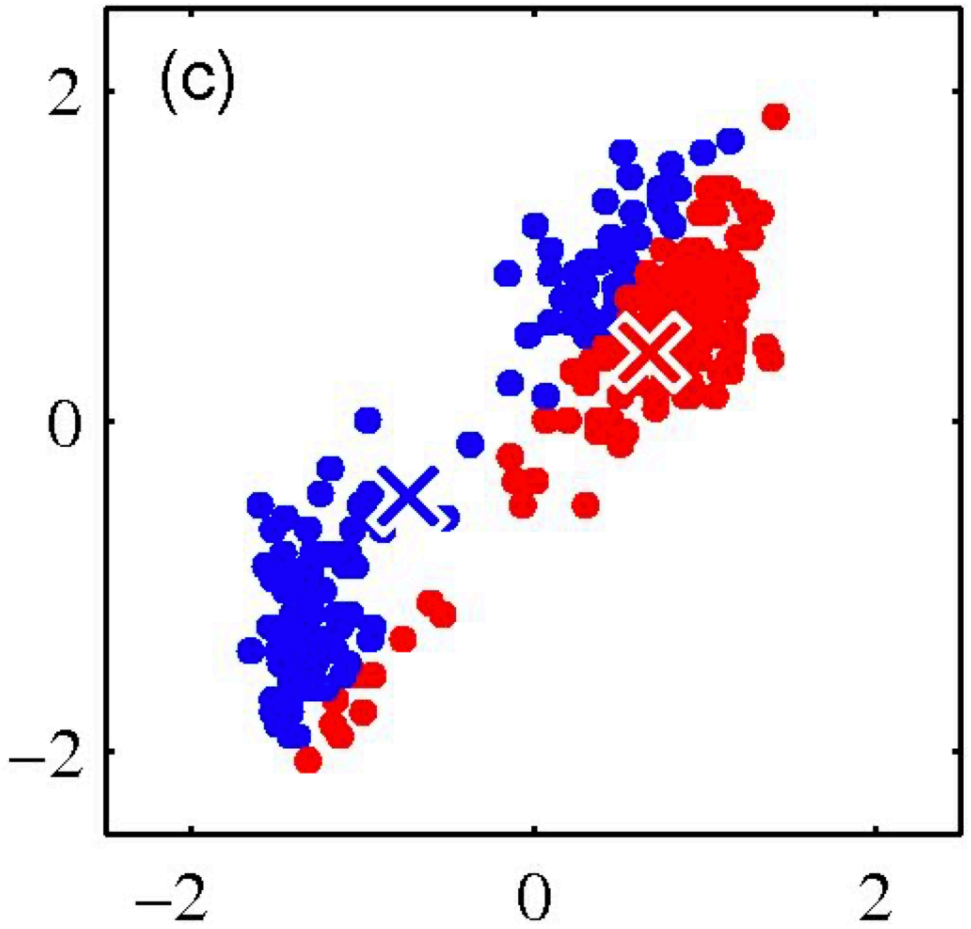


**Iterative Step 2**

• Change the cluster center to the average of the assigned points

# K-Means

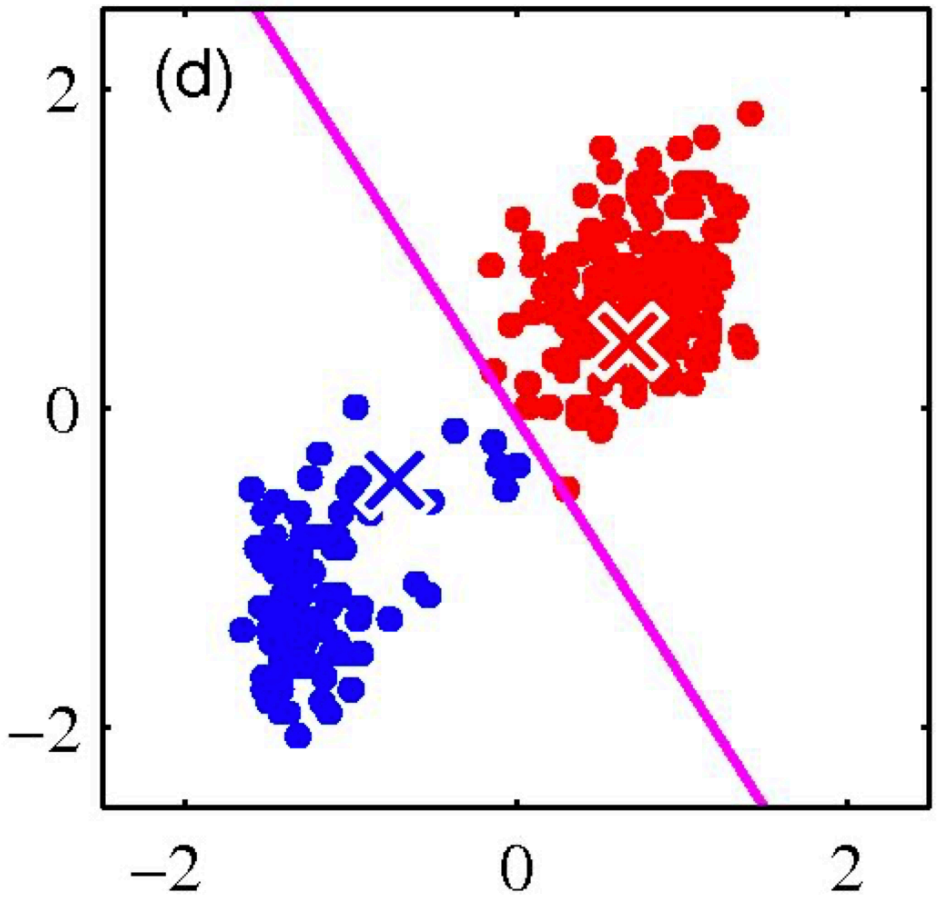


Repeat until convergence

# K-Means

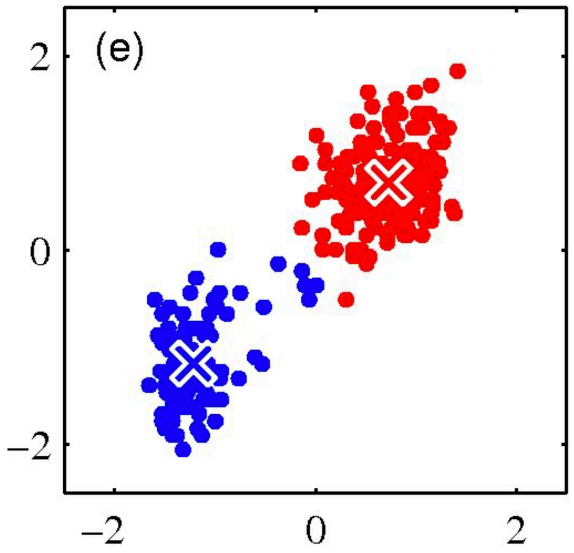


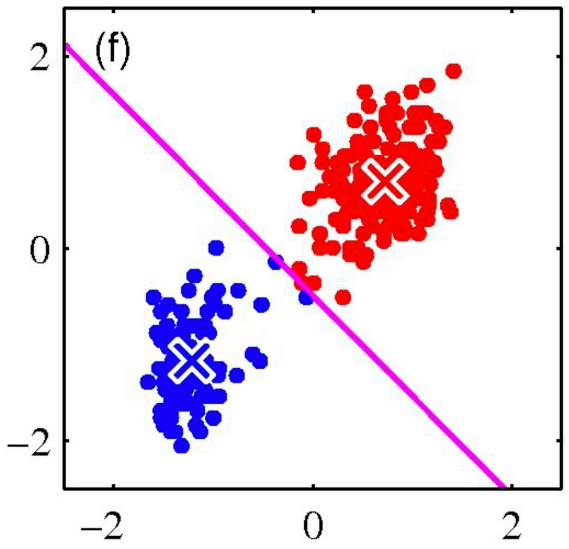


.....

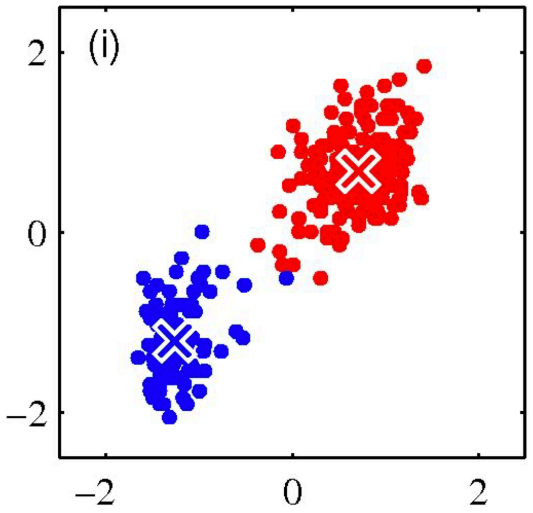

# Termination conditions

- Several possibilities, e.g.,
  - A fixed number of iterations.
  - Doc partition unchanged.

# How Many Clusters?

- Number of clusters *K* is given
  - Partition *n* docs into predetermined number of clusters
- Finding the “right” number of clusters is part of the problem
  - Given docs, partition into an “appropriate” number of subsets.
  - E.g., for query results - ideal value of *K* not known up
- Tradeoff between having more clusters (better focus within each cluster) and having too many clusters

# What is a Good Clustering?

- **Internal criterion:** A good clustering will produce high quality clusters in which:
  - the intra-class (that is, intra-cluster) similarity is high
  - the inter-class similarity is low
  - The measured quality of a clustering depends on both the document representation and the similarity measure used

# External criteria for clustering quality

- Quality measured by its ability to discover some or all of the hidden patterns or latent classes in gold standard data
- Assesses a clustering with respect to ground truth … requires *labeled data*
- Assume documents with $C$ gold standard classes, while our clustering algorithms produce $K$ clusters, $\omega_1, \omega_2, \ldots, \omega_K$ with $n_i$ members.

# External Evaluation of Cluster Quality

- Simple measure: purity, the ratio between the dominant class in the cluster $\omega_i$ and the size of cluster $\omega_i$

$$Purity(\omega_i) = \frac{1}{n_i} \max_j (n_{ij}) \quad j \in C$$

- Biased because having *n* clusters maximizes purity
- Others are entropy of classes in clusters (or mutual information between classes and clusters)

# Purity example

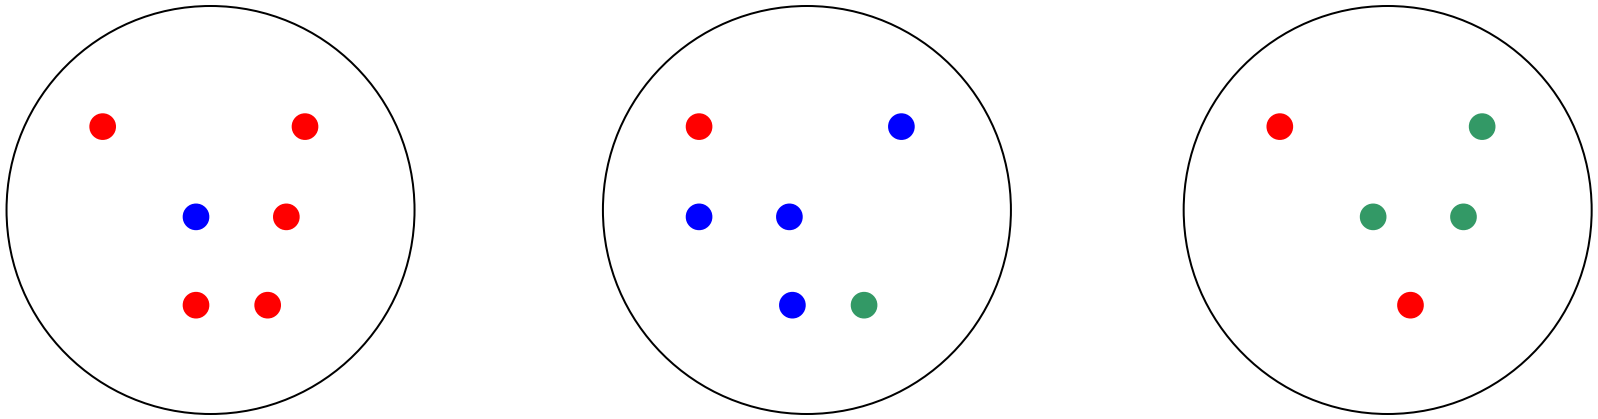

Cluster I  Cluster II  Cluster III

Cluster I: Purity = 1/6 (max(5, 1, 0)) = 5/6

Cluster II: Purity = 1/6 (max(1, 4, 1)) = 4/6

Cluster III: Purity = 1/5 (max(2, 0, 3)) = 3/5

# Final words

- In clustering, clusters are inferred from the data without human input (unsupervised learning)
- However, in practice, it's a bit less clear: there are many ways of influencing the outcome of clustering: number of clusters, similarity measure, representation of documents, ...