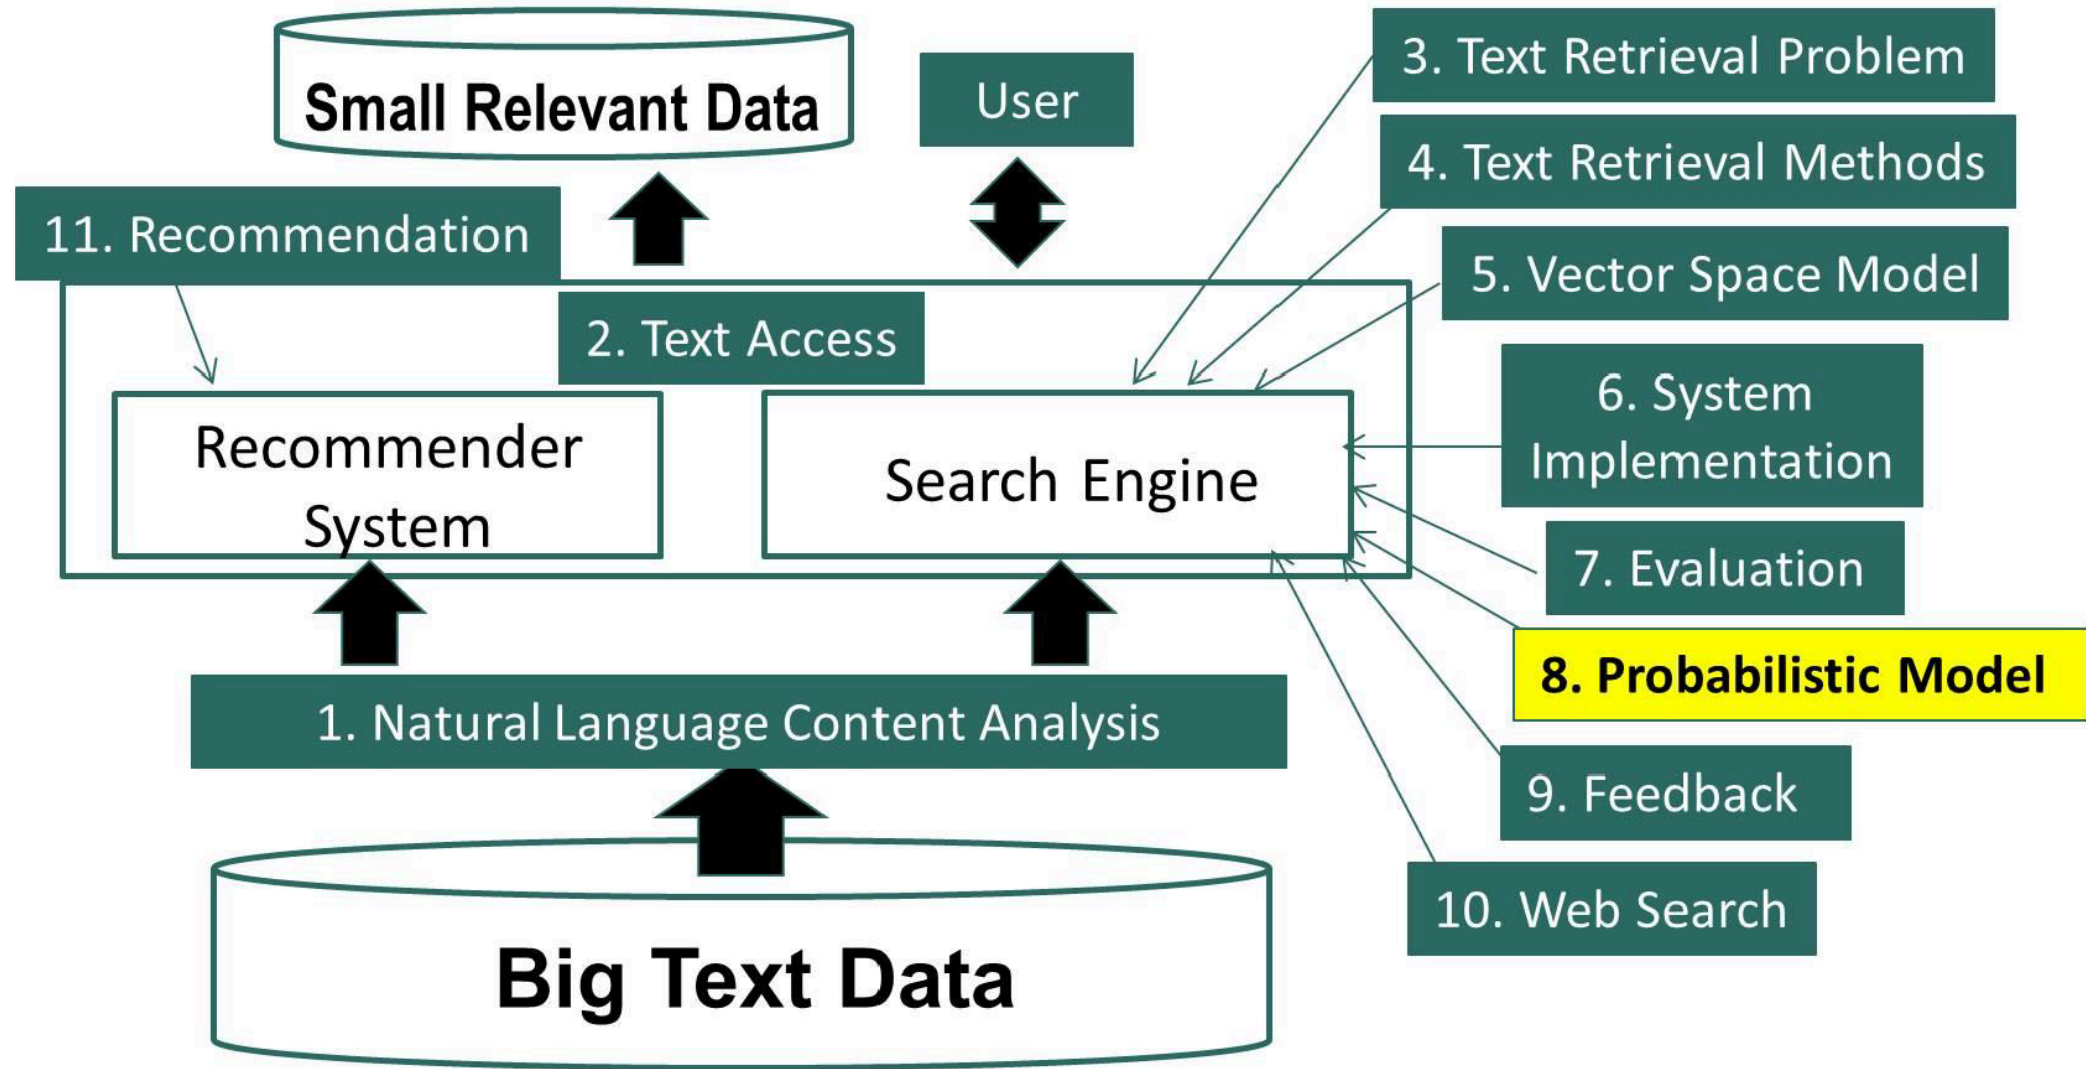


Information Retrieval & Text Mining

Probabilistic Retrieval Model: Smoothing

Dr. Saeed Ul Hassan

Probabilistic Retrieval Model: Smoothing




Ranking Function based on Query Likelihood

$$q = w_1 w_2 \dots w_n \quad p(q | d) = p(w_1 | d) \times \dots \times p(w_n | d)$$

$$f(q, d) = \log p(q | d) = \sum_{i=1}^n \log p(w_i | d) = \sum_{w \in V} c(w, q) \log p(w | d)$$

How should we estimate $p(w/d)$?

How to Estimate $p(w | d)$



How to smooth a LM

- Key Question: what probability should be assigned to an unseen word?
- Let the probability of an unseen word be proportional to its probability given by a reference LM
- One possibility: Reference LM = Collection LM

$$p(w | d) = \begin{cases} p_{Seen}(w | d) \\ \alpha_d p(w | C) \end{cases}$$

Discounted ML estimate
if w is seen in d
otherwise

Collection language model

Rewriting the Ranking Function with Smoothing

$$\log p(q | d) = \sum_{w \in V} c(w, q) \log p(w | d)$$

$$= \sum_{w \in V, c(w, d) > 0} c(w, q) \log p_{\text{Seen}}(w | d) + \sum_{w \in V, c(w, d) = 0} c(w, q) \log \alpha_d p(w | C)$$

Query words **matched** in d Query words **not matched** in d

$$\sum_{w \in V} c(w, q) \log \alpha_d p(w | C) + \sum_{w \in V, c(w, d) > 0} c(w, q) \log \alpha_d p(w | C)$$

All query words Query words **matched** in d

$$= \sum_{w \in V, c(w, d) > 0} c(w, q) \log \frac{p_{\text{Seen}}(w | d)}{\alpha_d p(w | C)} + |q| \log \alpha_d + \sum_{w \in V} c(w, q) \log p(w | C)$$

Benefit of Rewriting

- Better understanding of the ranking function
 - Smoothing with $p(w|C) \rightarrow$ TF-IDF weighting + length norm.

$$\log p(q | d) = \sum_{\substack{w_i \in d \\ w_i \in q}} \left[\log \frac{p_{\text{Seen}}(w_i | d)}{\alpha_d p(w_i | C)} \right] + n \log \alpha_d + \boxed{\sum_{i=1}^n \log p(w_i | C)}$$

- Enable efficient computation

Benefit of Rewriting

- Better understanding of the ranking function
 - Smoothing with $p(w|C) \rightarrow$ TF-IDF weighting + length norm.

$$\log p(q | d) = \sum_{\substack{w_i \in d \\ w_i \in q}} \left[\log \frac{\text{TF weighting} \left(\frac{p_{\text{Seen}}(w_i | d)}{\alpha_d p(w_i | C)} \right) \right] + n \log \alpha_d + \boxed{\sum_{i=1}^n \log p(w_i | C)}$$

matched query terms (points to the summation index)

IDF weighting (points to $p(w_i | C)$)

Doc length normalization (points to $n \log \alpha_d$)

Ignore for ranking (points to the boxed summation term)

- Enable efficient computation

Summary

- Smoothing of $p(w | d)$ is necessary for query likelihood
- General idea: smoothing with $p(w | C)$
 - The probability of an unseen word in d is assumed to be proportional to $p(w | C)$
 - Leads to a general ranking formula for query likelihood with TF-IDF weighting and document length normalization
 - Scoring is primarily based on sum of weights on matched query terms
- However, how exactly should we smooth?