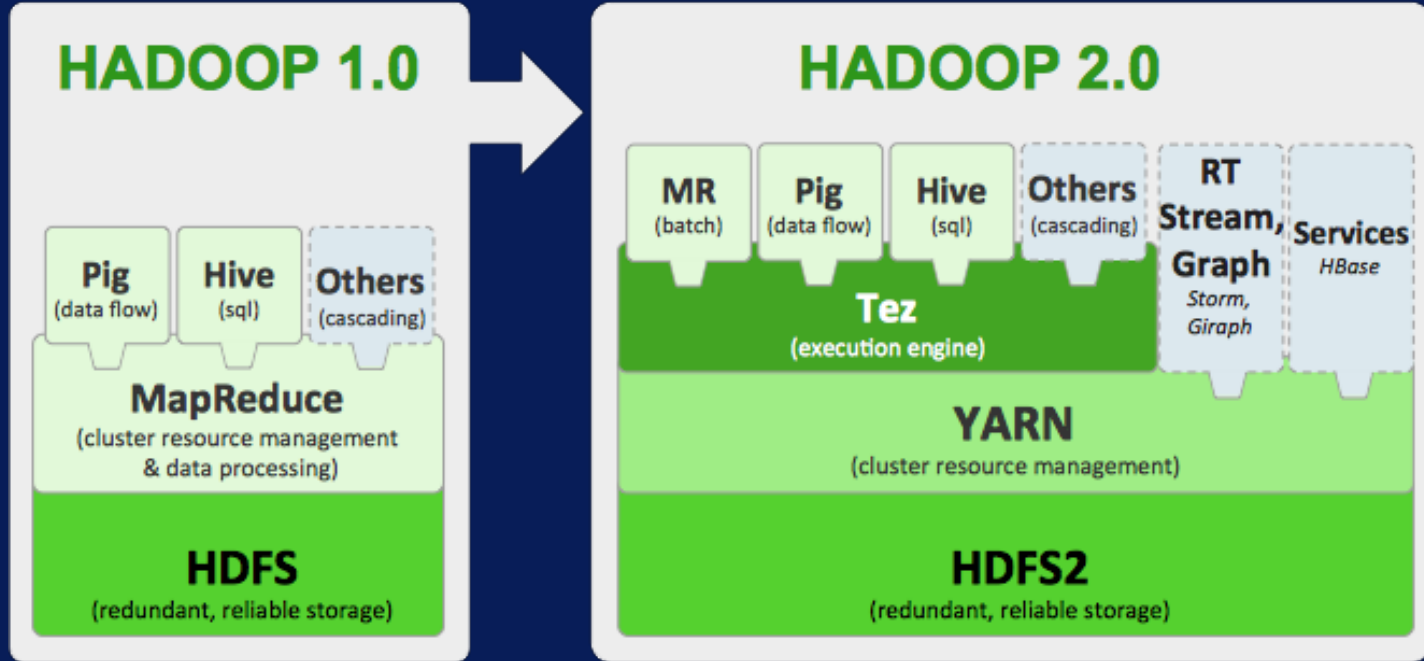


# The Basic Hadoop Components

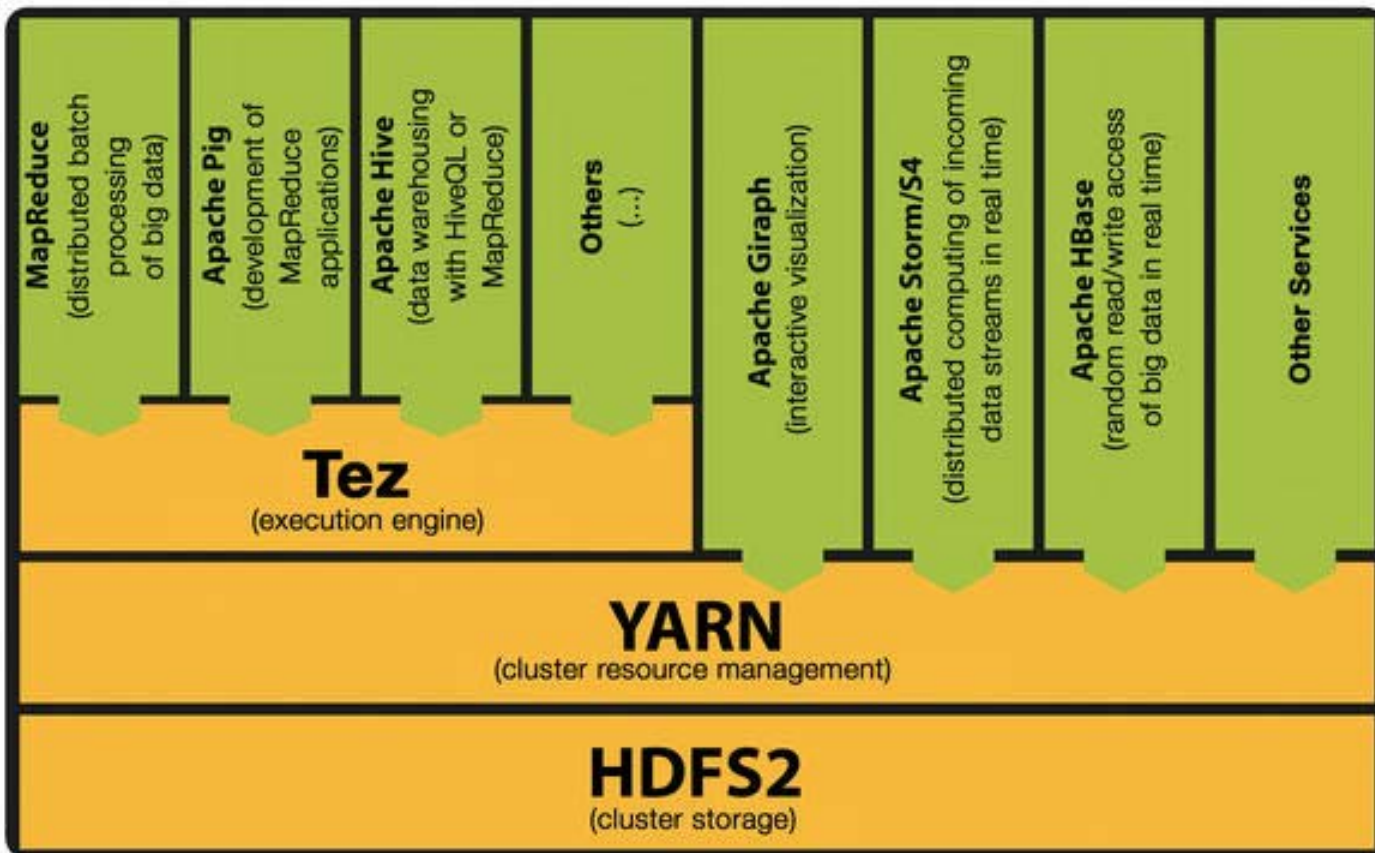
- *Hadoop Common* - libraries and utilities
- *Hadoop Distributed File System (HDFS)* – a distributed file-system
- *Hadoop YARN* – a resource-management platform, scheduling
- *Hadoop MapReduce* – a programming model for large scale data processing

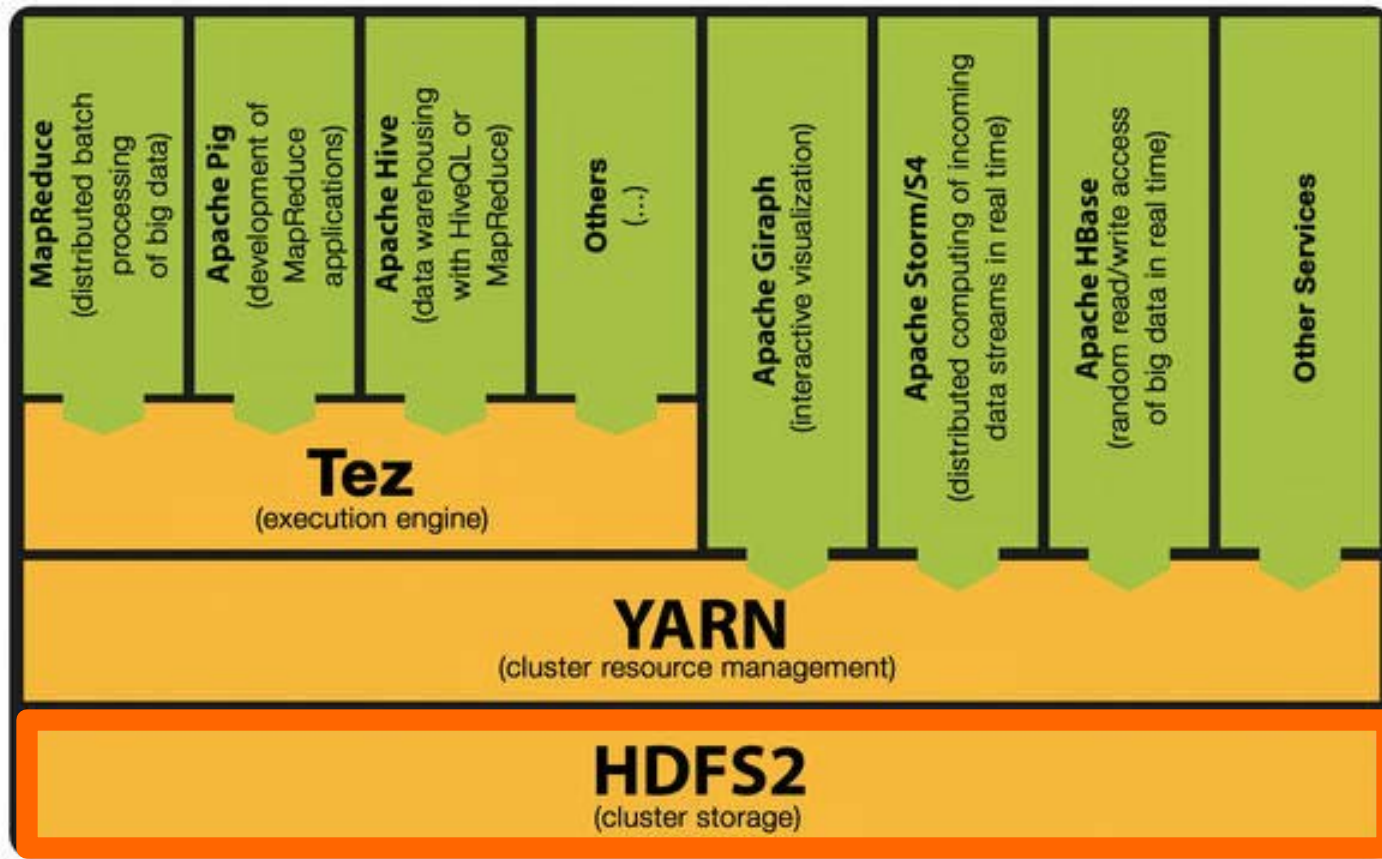
# Hadoop Stack Transition



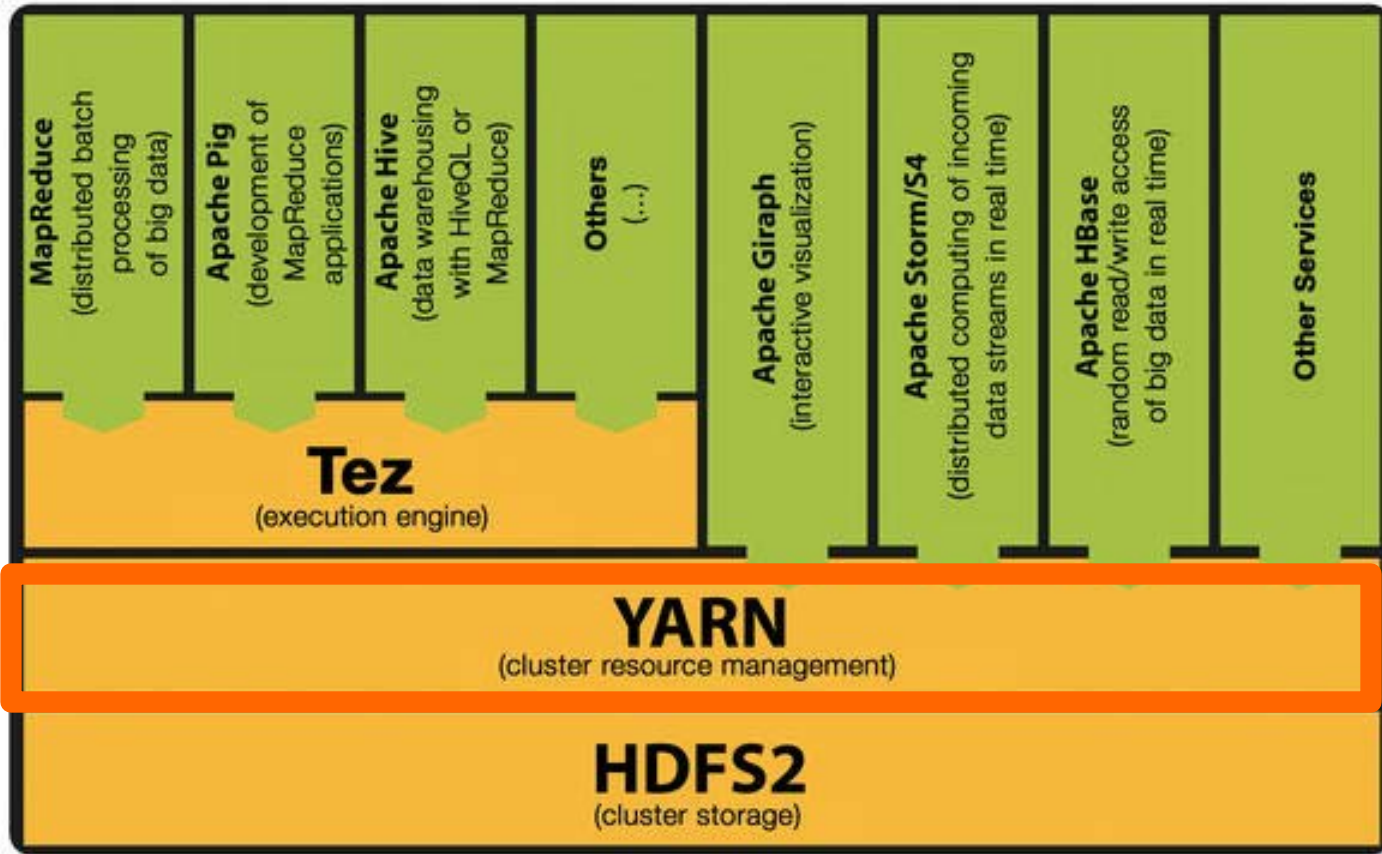
# Applications and Frameworks

- **HBase** – a scalable data warehouse with support for large tables.
- **Hive** – a data warehouse infrastructure that provides data summarization and ad hoc querying
- **Pig** – A high-level data-flow language and execution framework for parallel computation
- **Spark** – a fast and general compute engine for Hadoop data. Wide range of applications – ETL, Machine Learning, stream processing, and graph analytics.

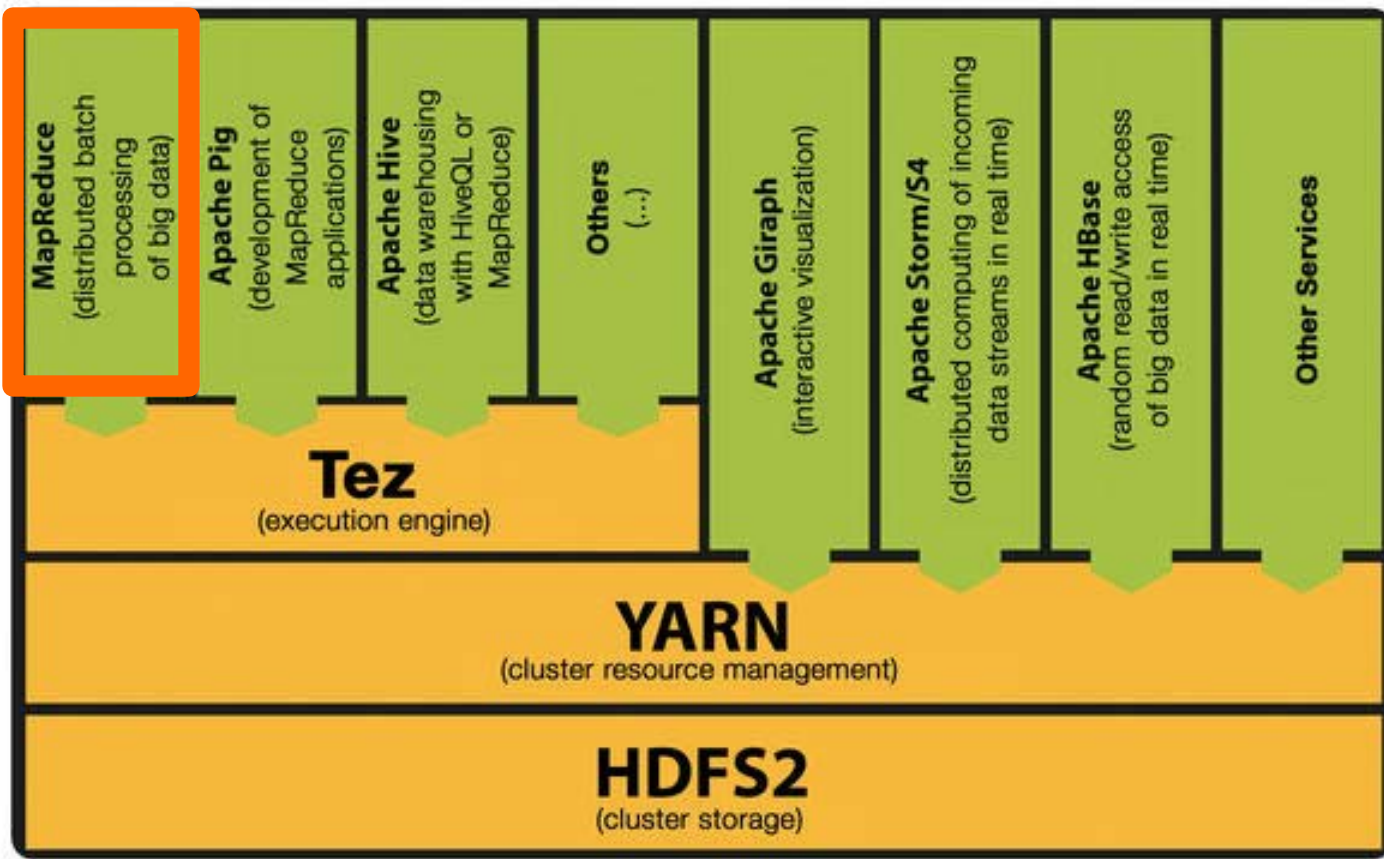




Distributed Filesystem

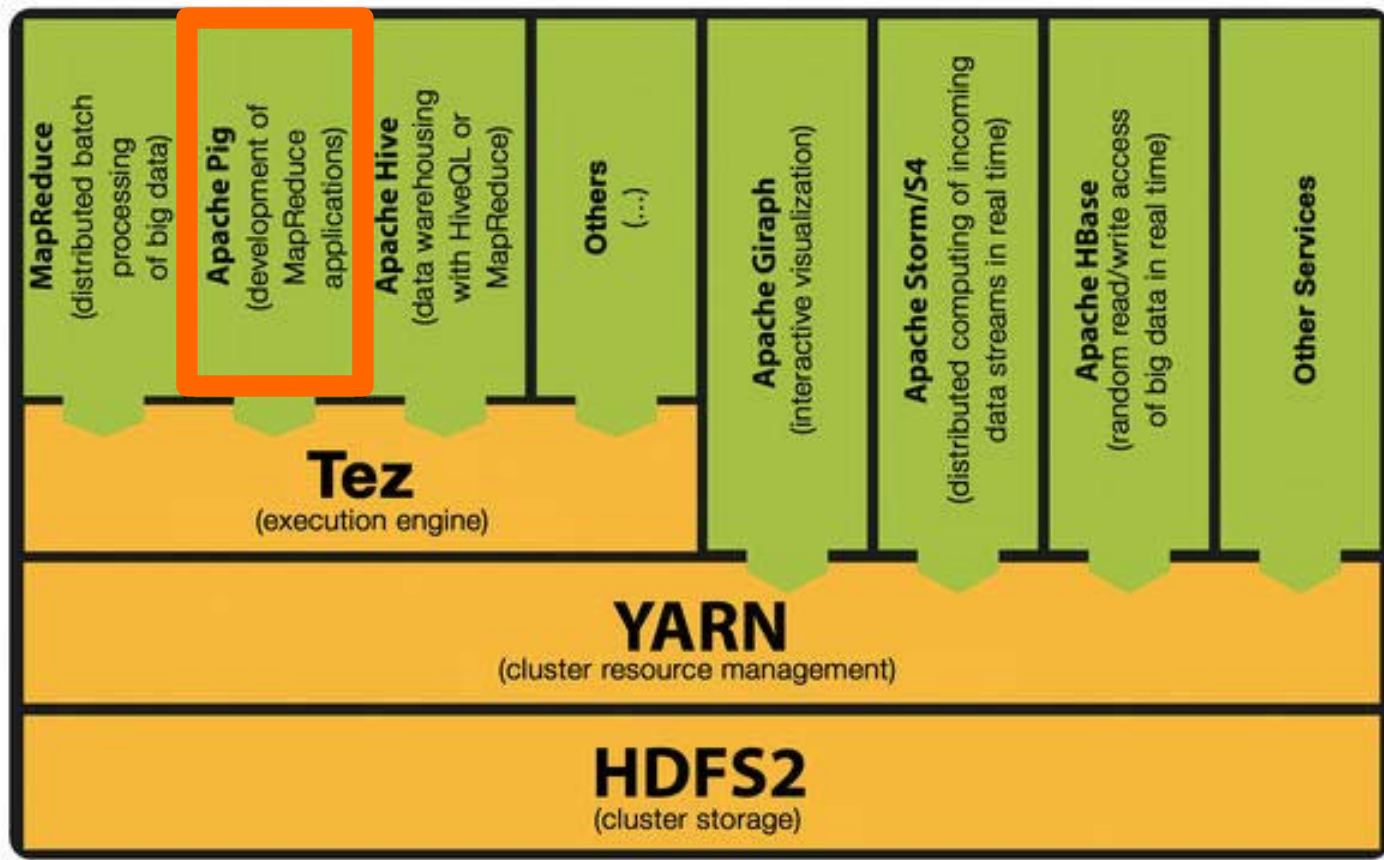


Resource Management, Scheduling



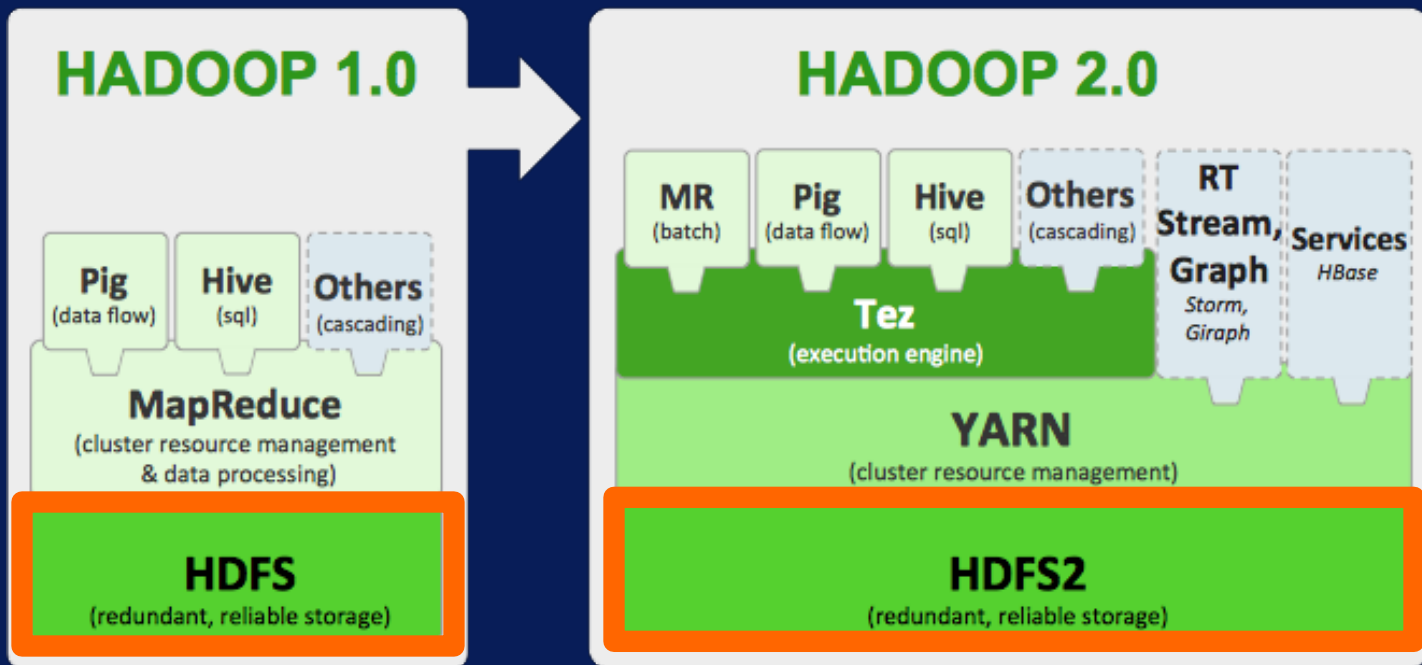
YARN-based system for  
parallel processing of data







# HDFS and HDFS2

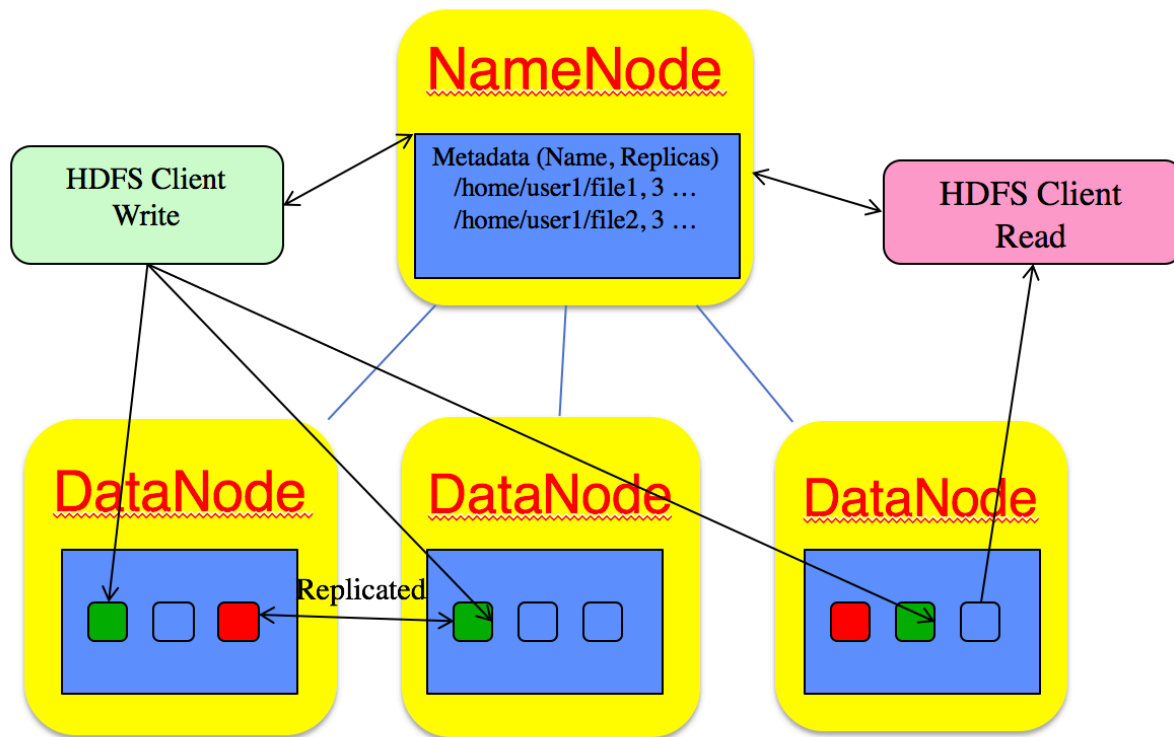


# Original HDFS Design Goals

- *Resilience to hardware failure*
- *Streaming data access*
- *Support for large dataset, scalability to hundreds/thousands of nodes with high aggregate bandwidth*
- *Application locality to data*
- *Portability across heterogeneous hardware and software platforms*

# Original HDFS Design

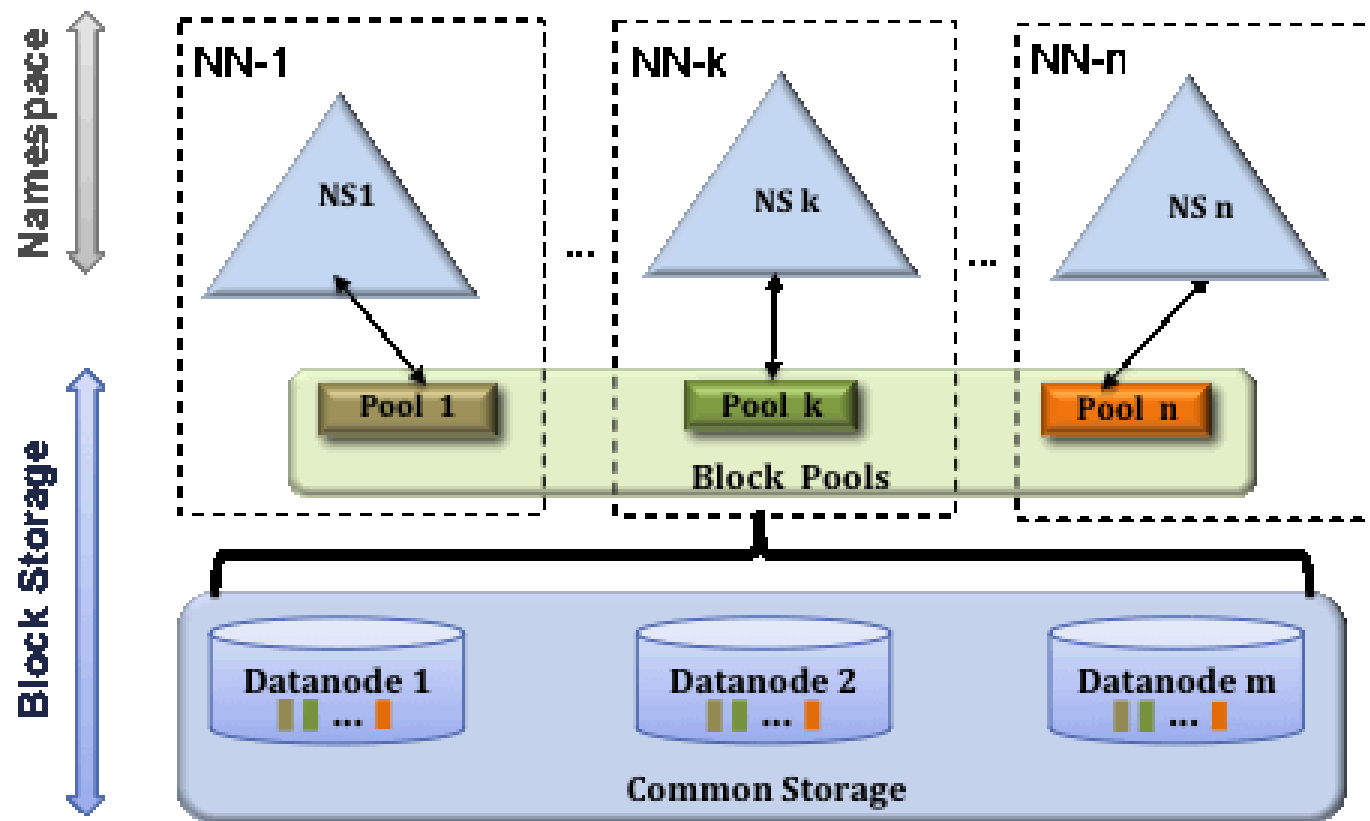
- **Single NameNode** - a master server that manages the file system namespace and regulates access to files by clients.
- **Multiple DataNodes** – typically one per node in the cluster. Functions:
  - Manage storage
  - Serving read/write requests from clients
  - Block creation, deletion, replication based on instructions from NameNode



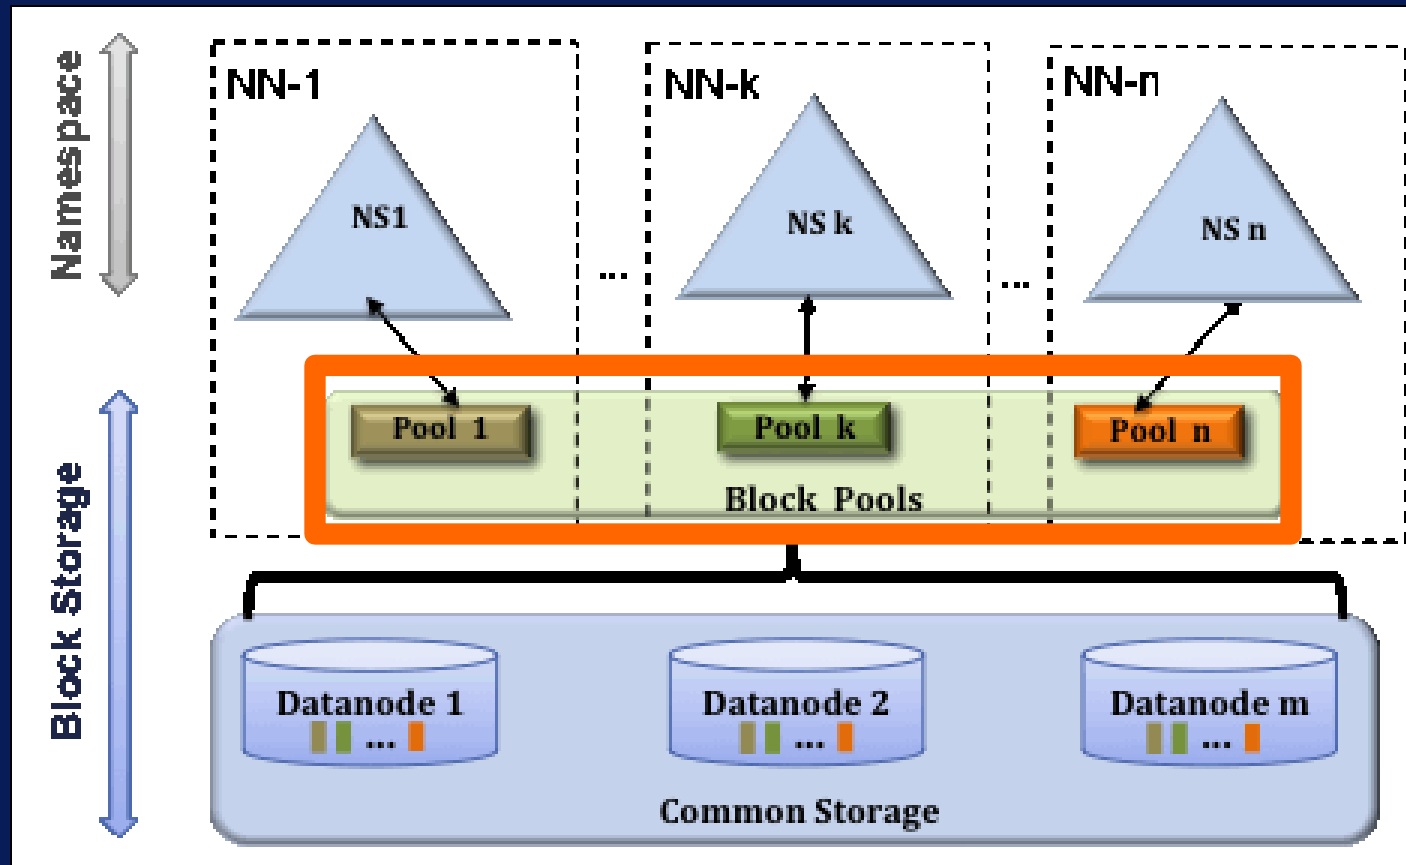
# HDFS in Hadoop 2

- *HDFS Federation*
- *Multiple Namenode servers*
- *Multiple namespaces*
- *High Availability – redundant NameNodes*
- *Heterogeneous Storage and Archival Storage*
  - *ARCHIVE, DISK, SSD, RAM\_DISK*

# Federation



# Federation: Block Pools





# Federation: Benefits

- *Allows namespace scaling*
- *Scales up filesystem read/write throughput*
- *Isolation*

# MapReduce Framework

- *Software framework – for writing parallel data processing applications*
- *MapReduce job splits data into chunks*
- *Map tasks process data chunks*
- *Framework sorts map output*
- *Reduce tasks use sorted map data as input*

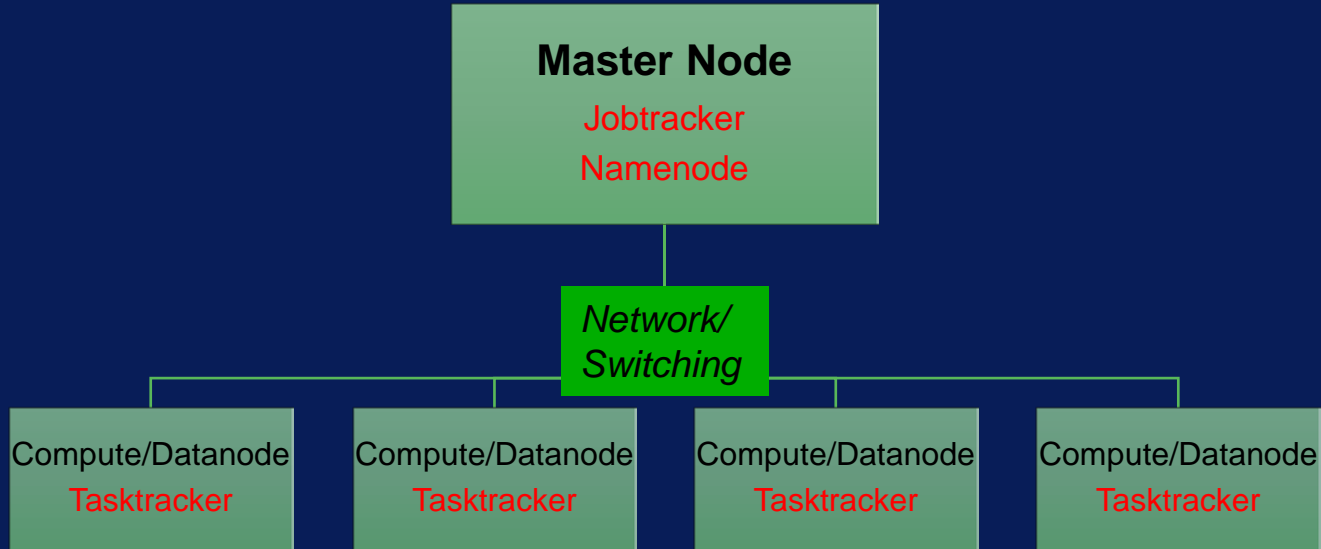
# MapReduce Framework

- *Typically compute and storage nodes are the same.*
- *MapReduce tasks and HDFS running on the same nodes*
- *Can schedule tasks on nodes with data already present.*

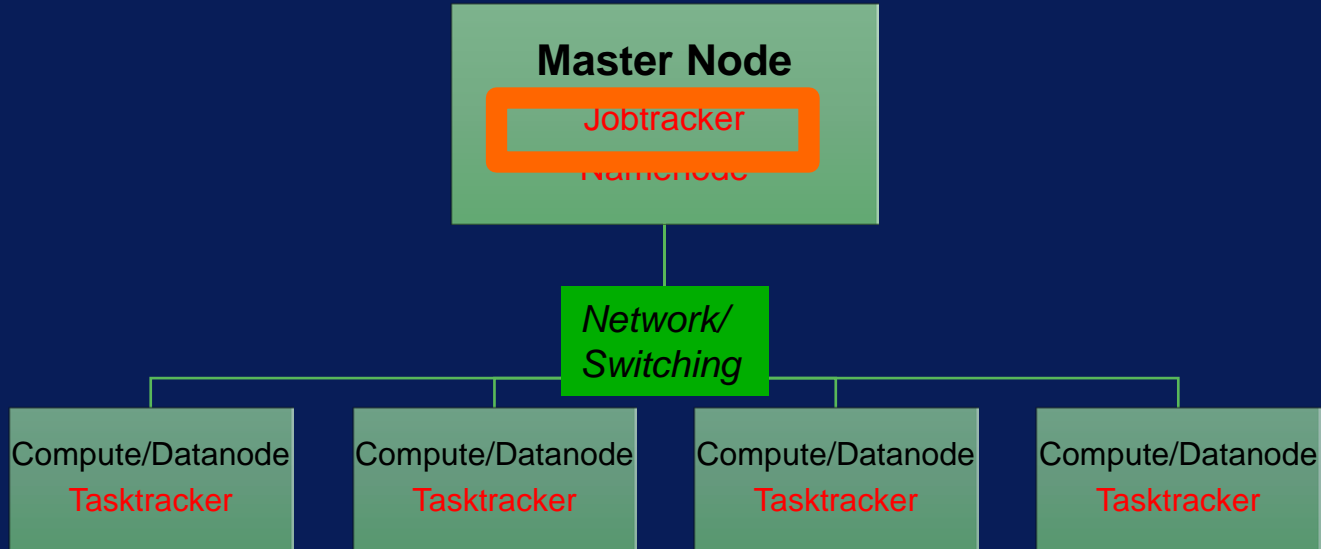
# Original MapReduce Framework

- *Single master JobTracker*
- *JobTracker schedules, monitors, and re-executes failed tasks.*
- *One slave TaskTracker per cluster node*
- *TaskTracker executes tasks per JobTracker requests.*

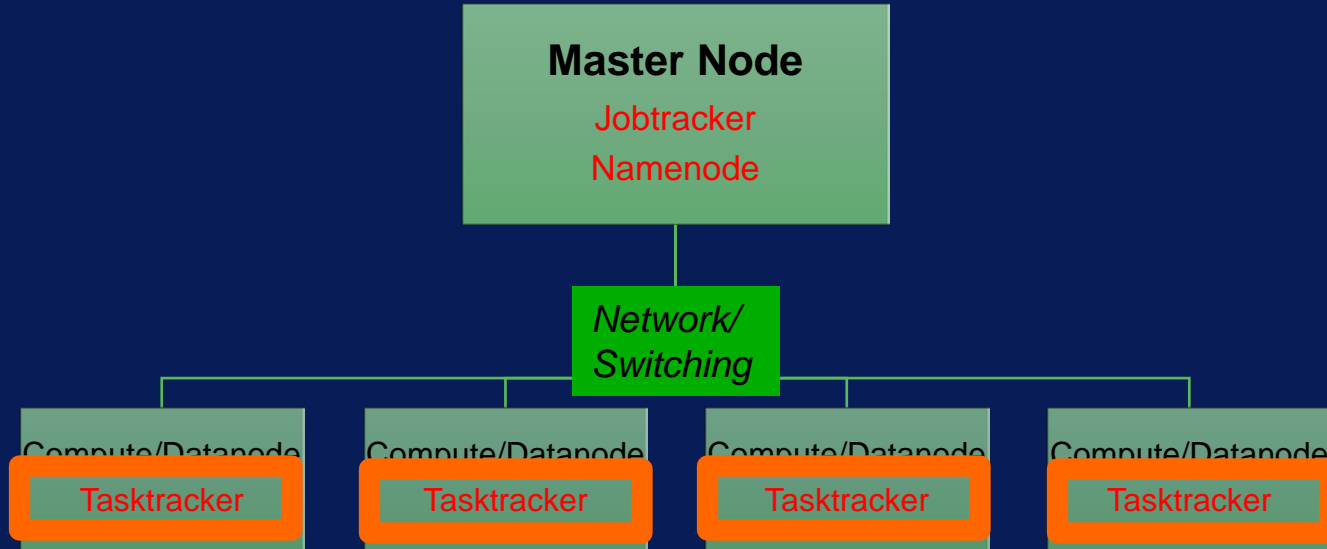
# Original Hadoop Architecture



# Original Hadoop Architecture



# Original Hadoop Architecture

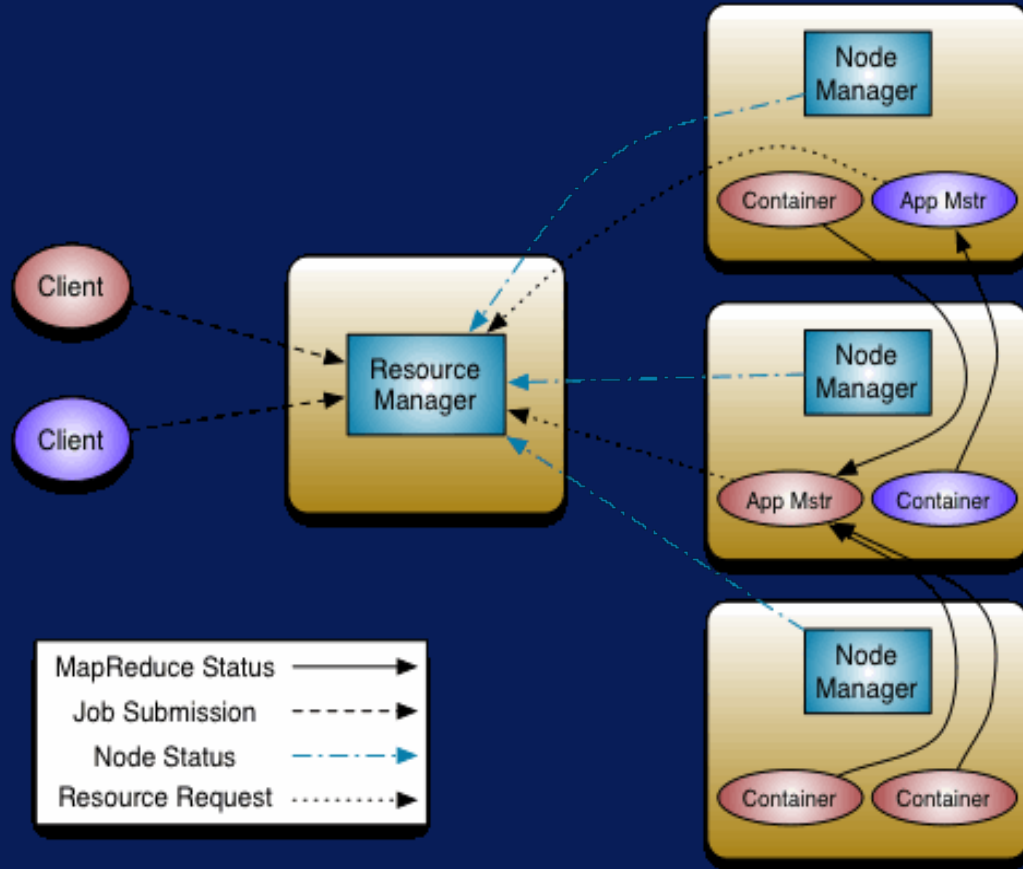




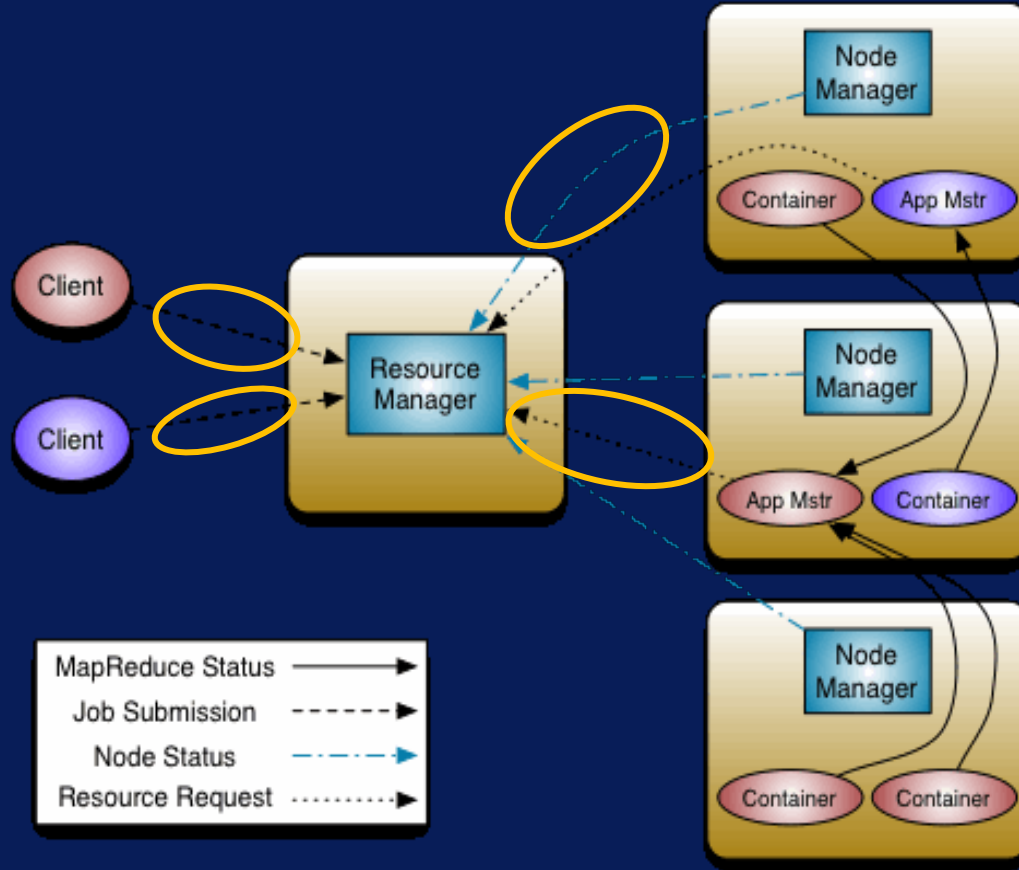
# YARN: NexGen MapReduce

- *Main idea – Separate resource management and job scheduling/monitoring.*
- *Global ResourceManager (RM)*
- *NodeManager on each node*
- *ApplicationMaster – one for each application*

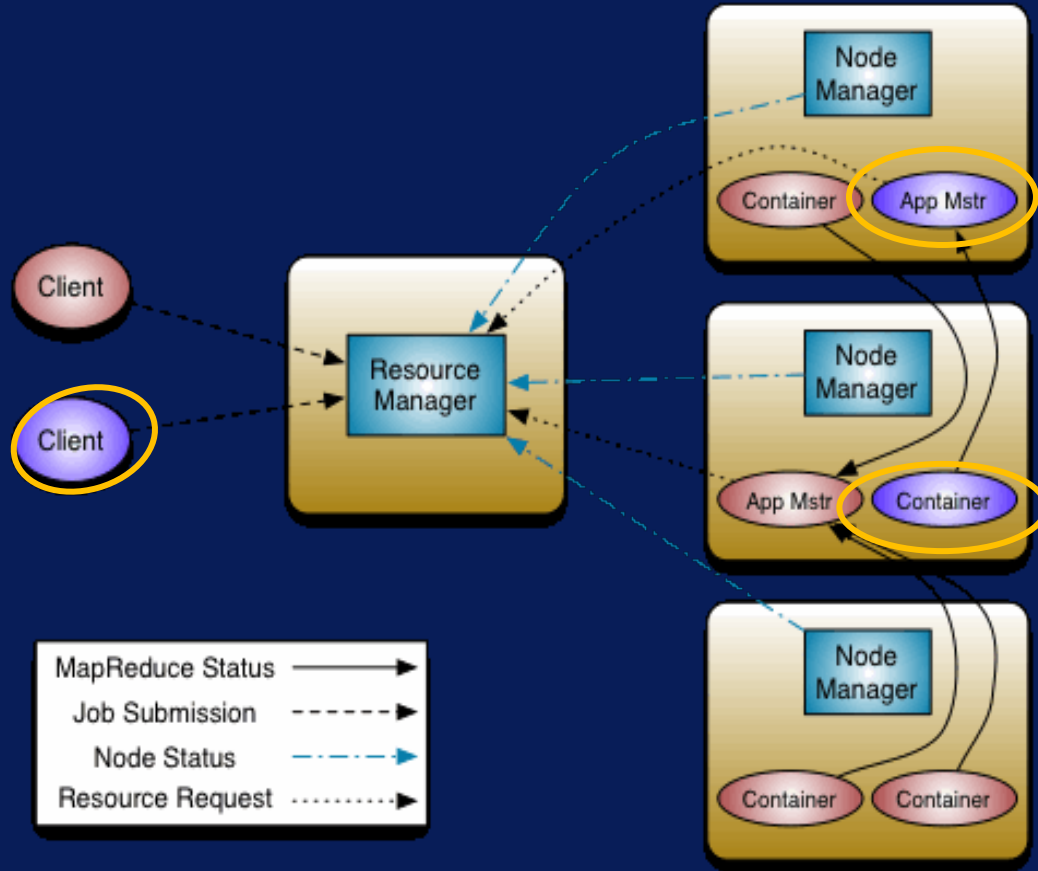
# YARN Architecture



# YARN Architecture



# YARN Architecture



# Additional YARN Features

- *High Availability ResourceManager*
- *Timeline Server*
- *Use of Cgroups*
- *Secure Containers*
- *YARN – web services  
REST APIs*