

Big Data Analytics Introduction to Hive

What is Hive?

**Data warehouse
infrastructure build on top
of Hadoop for querying
and managing large data
sets**

Why Hive?

Hadoop is great!

MapReduce is very low level


Lack of expressiveness

**Higher level data processing
languages are needed**

Hive Features

Designed for OLAP

*On Line
Analytical
Processing*




SQL type language for querying

*HiveQL or
HQL*



It is familiar, fast, scalable, and
extensible

*Can plug in
map/reduce scripts in
language of choice*



Hive is NOT

Relational database

Designed for Online Transaction Processing (OLTP)

*Online transaction processing, or **OLTP**, is a class of information systems that facilitate and manage transaction-oriented applications, typically for data entry and retrieval transaction processing (Wikipedia)*

History

**Early Hive development work
started at Facebook in 2007**

**Hive is an Apache project under
Hadoop**

<http://hive.apache.org>

*ETL =
Extract, Transform
and Load*

*Data Warehouse
Infrastructure for
Hadoop*


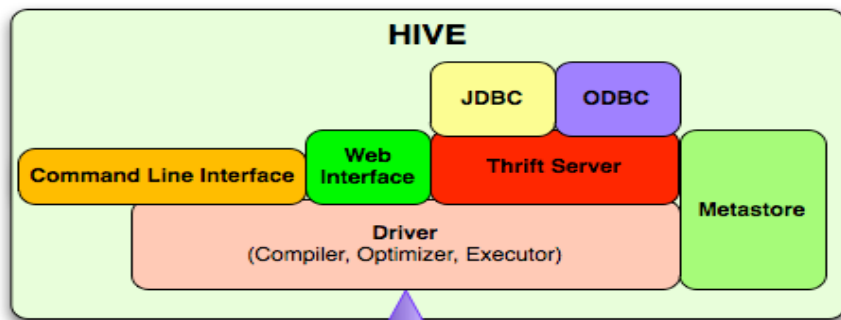
*SQL-like query
language (QL)*




*Enables developers
to utilize custom
mappers and reducers*

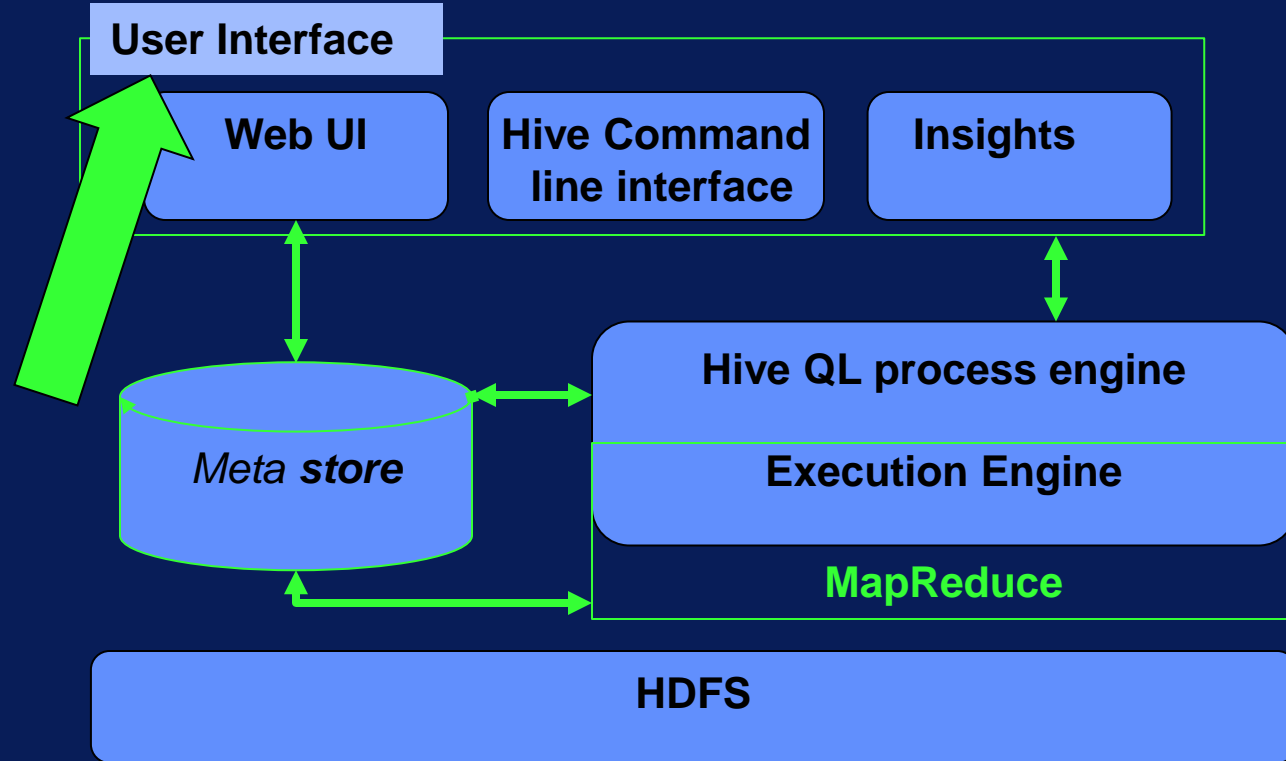
*Provides tools to
enable ETL on
large data*

Hive Architecture and Components

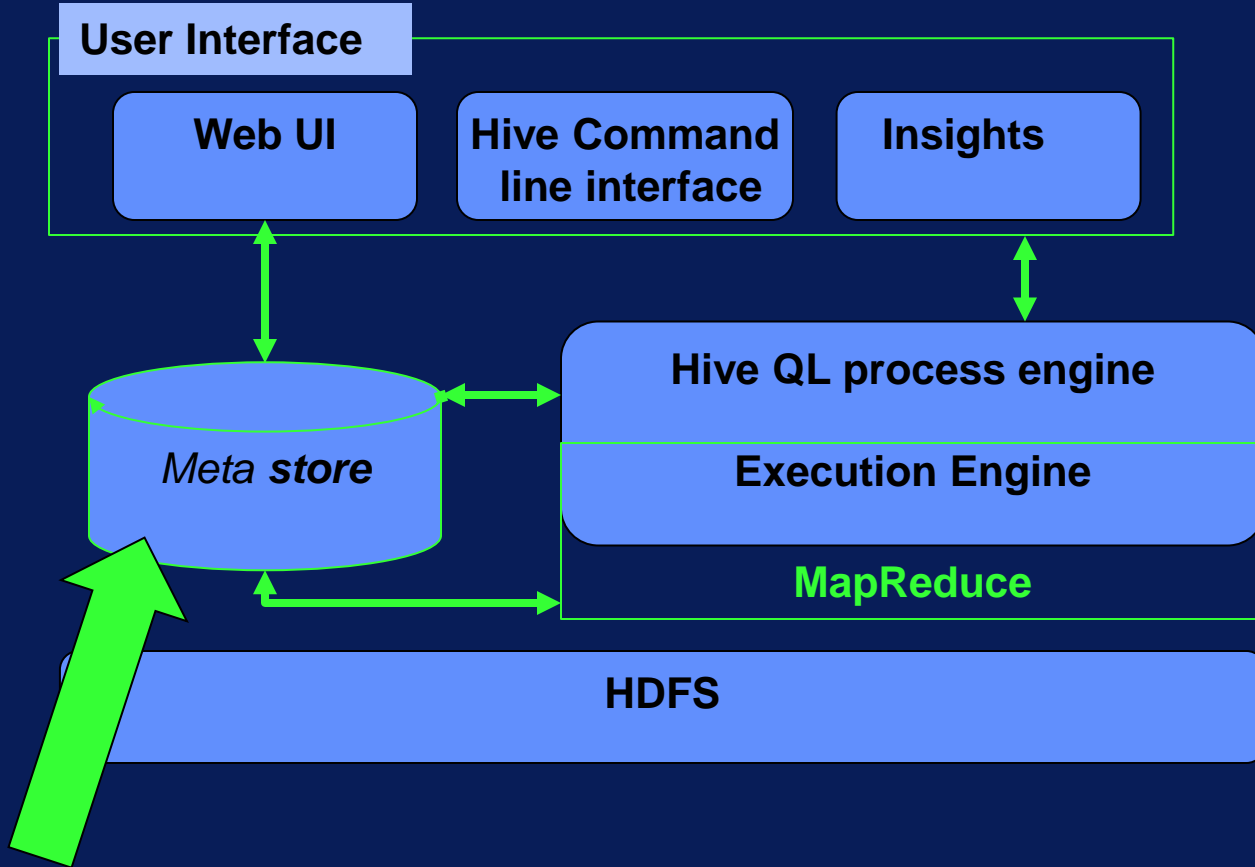




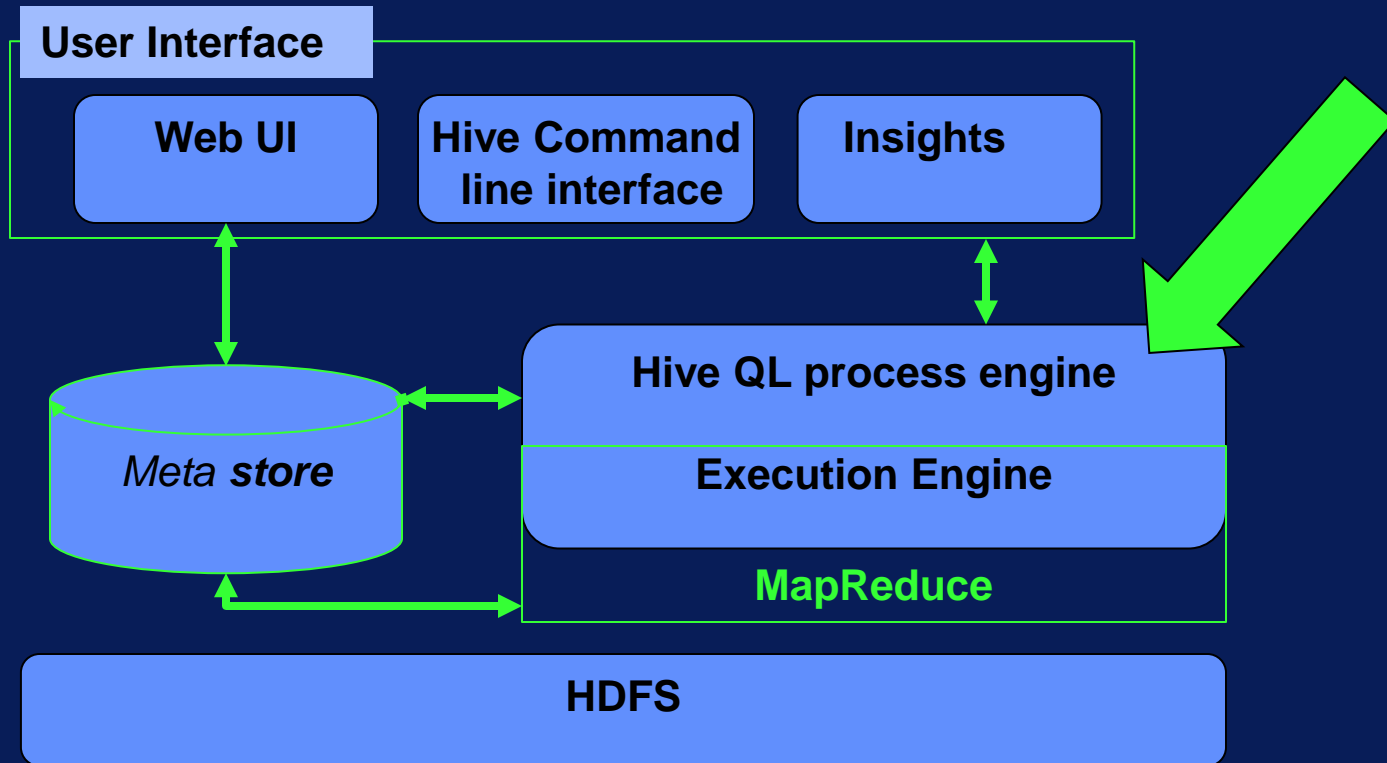
Hive Architecture



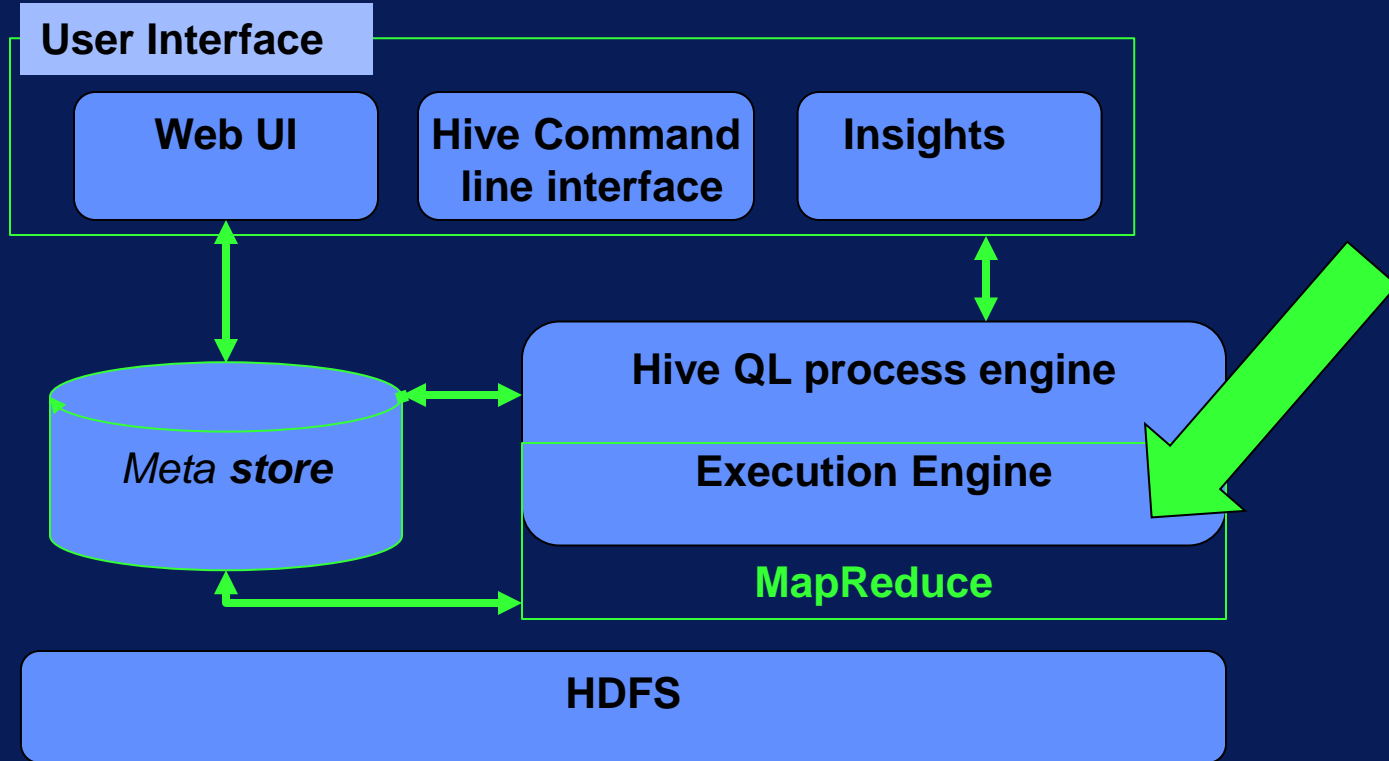
Hive Architecture

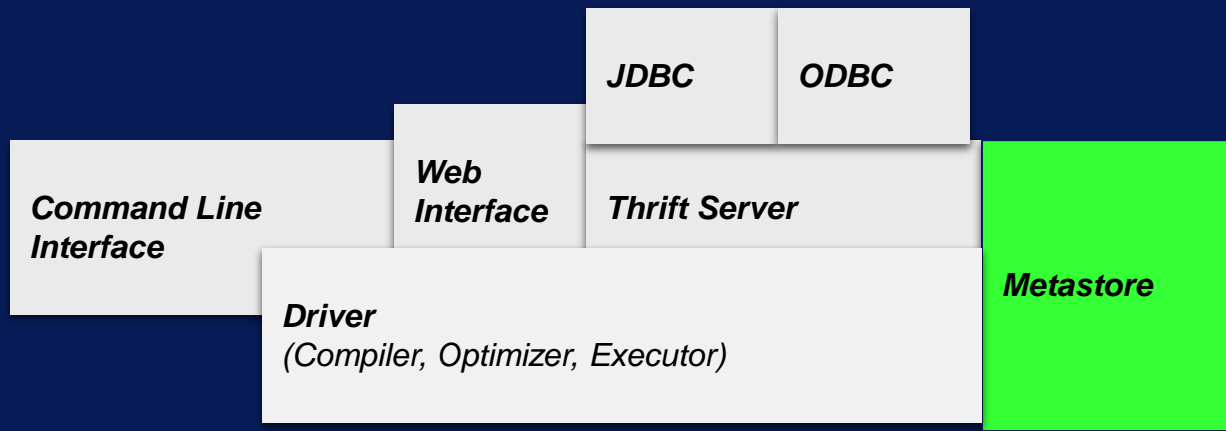


Hive Architecture



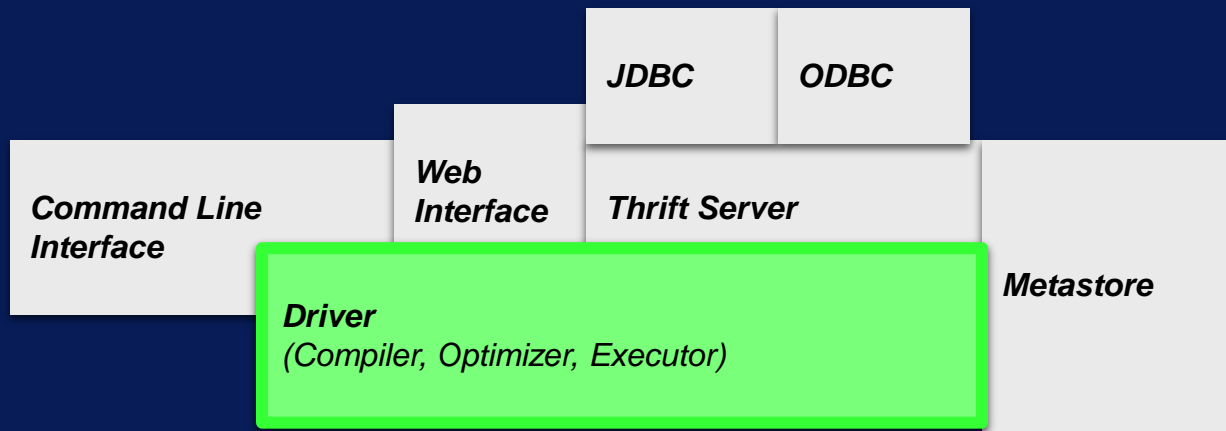
Hive Architecture





Metastore

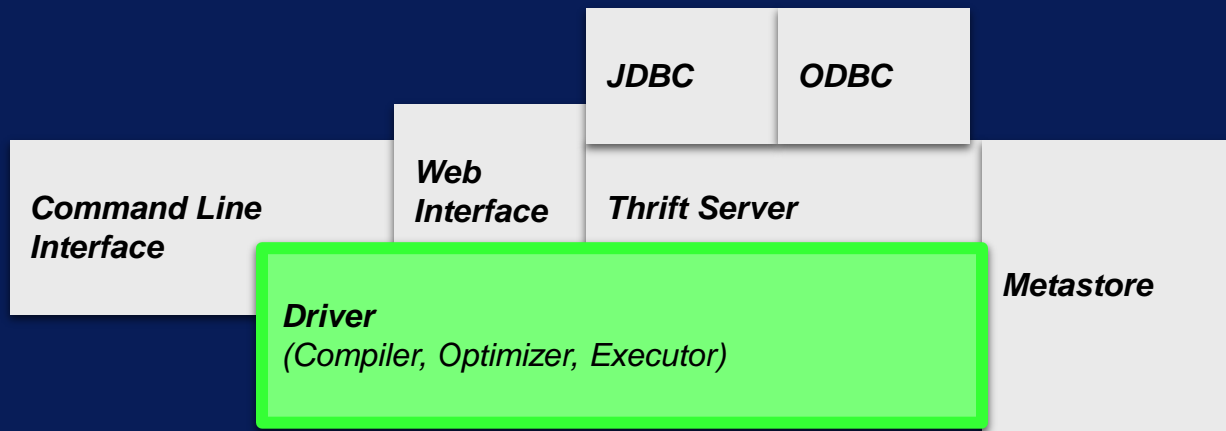
Stores the system catalog and meta data about tables, columns, partitions etc.



Driver

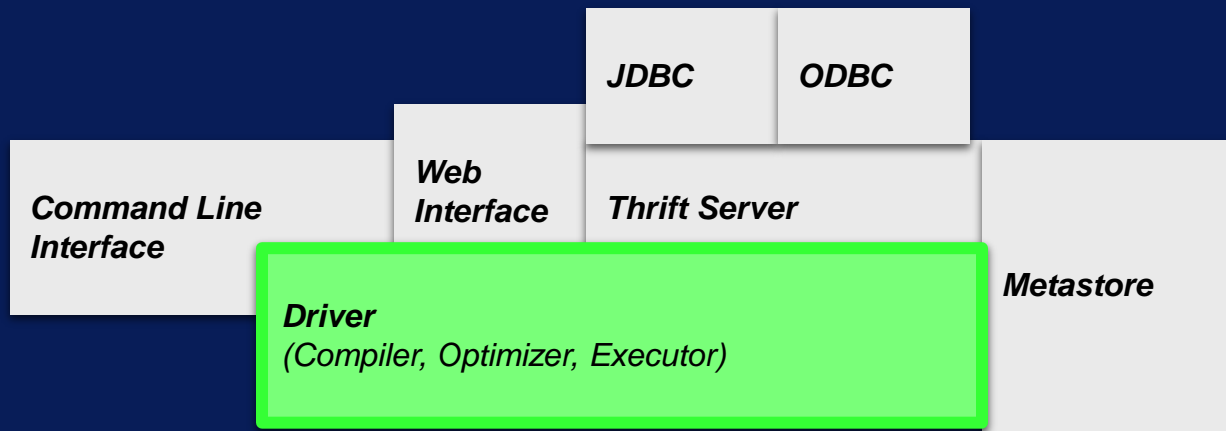
Manages the lifecycle of a HiveQL statement

Maintains a session handle and any session statistics



Query Compiler

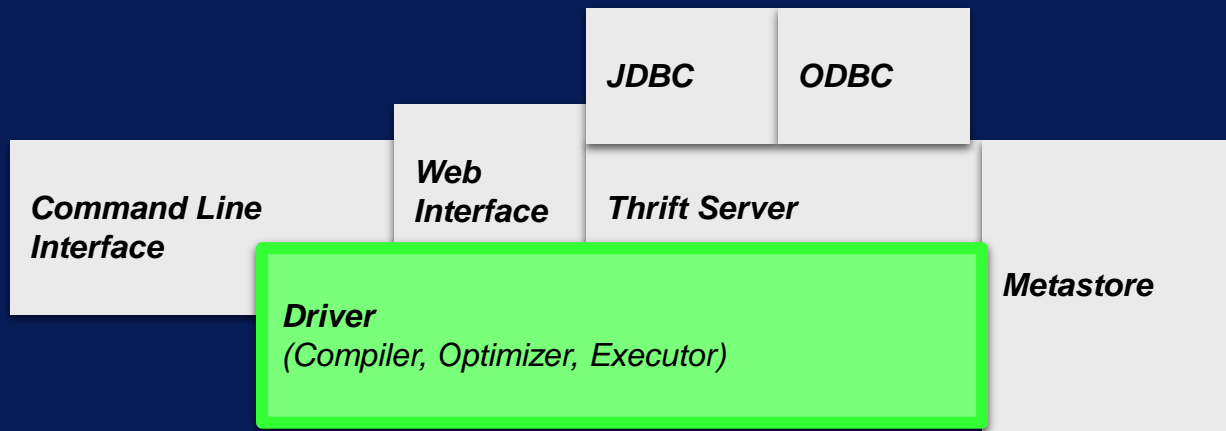
The component that compiles HiveQL into a directed acyclic graph of map/reduce tasks



Optimizer


Consists of a chain of transformations

Performs Column Pruning , Partition Pruning, Repartitioning of Data



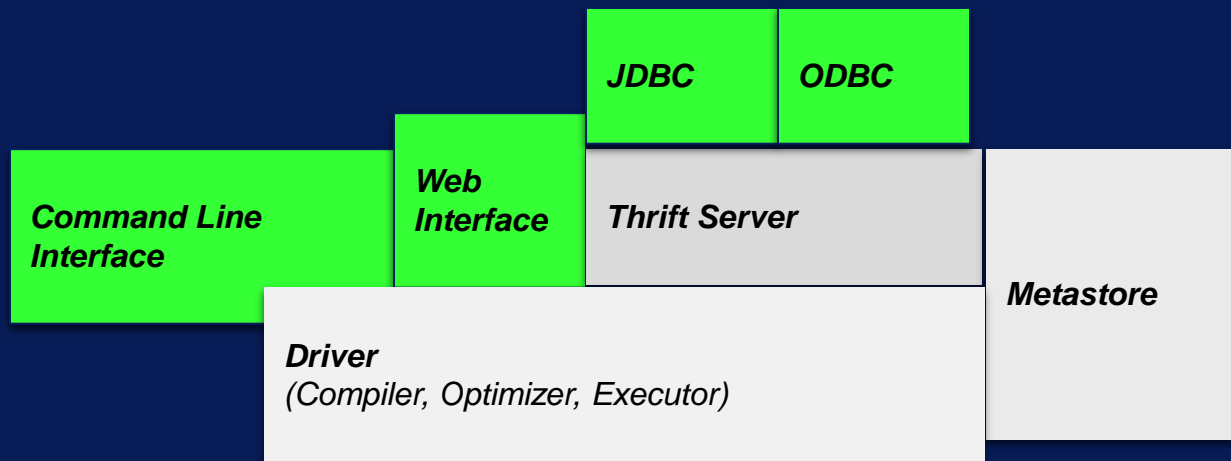
Execution Engine

Executes the tasks produced by the compiler in proper dependency order
Interacts with the underlying Hadoop instance



HiveServer

Provides a Thrift interface and a JDBC/ODBC server Enables Hive integration with other applications

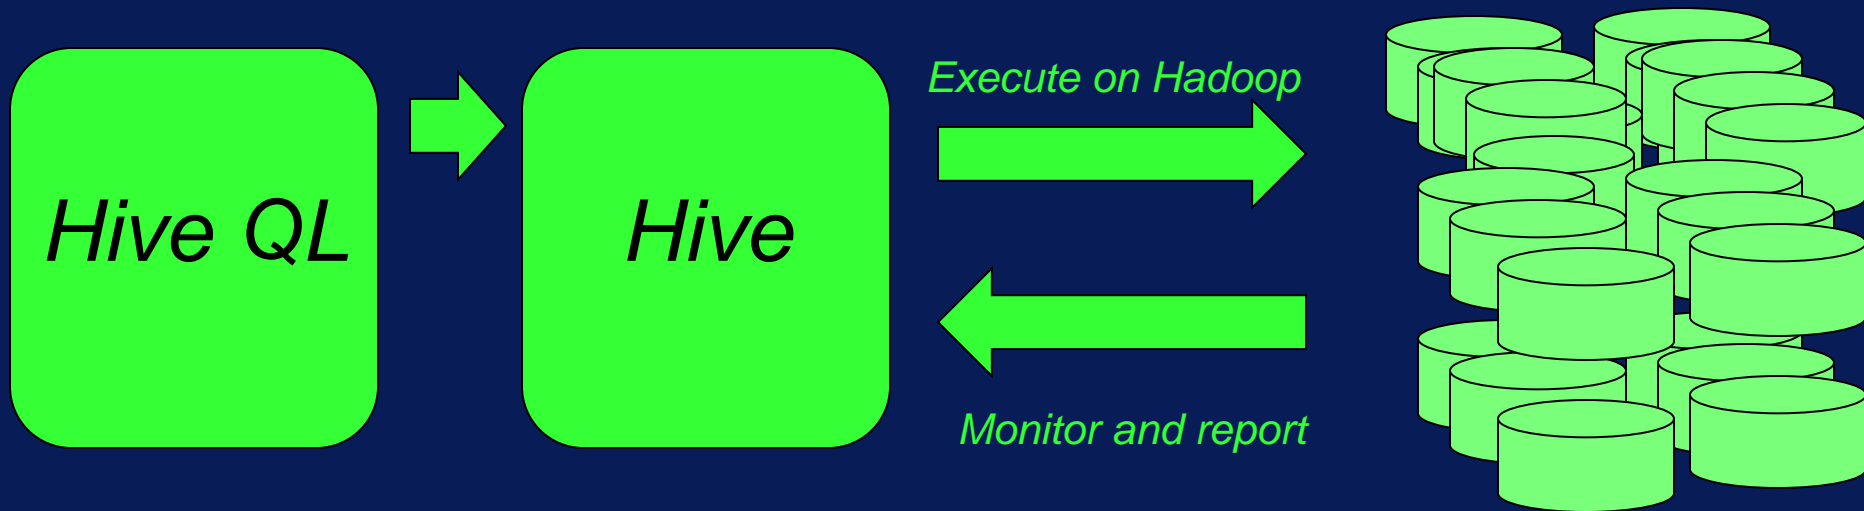


Client Components

Command Line Interface(CLI)

Web UI

JDBC/ODBC driver



Hive

Hadoop

Execution Engine

Job Tracker

MapReduce

Task Tracker

HDFS

Interface

1

Driver

Compiler

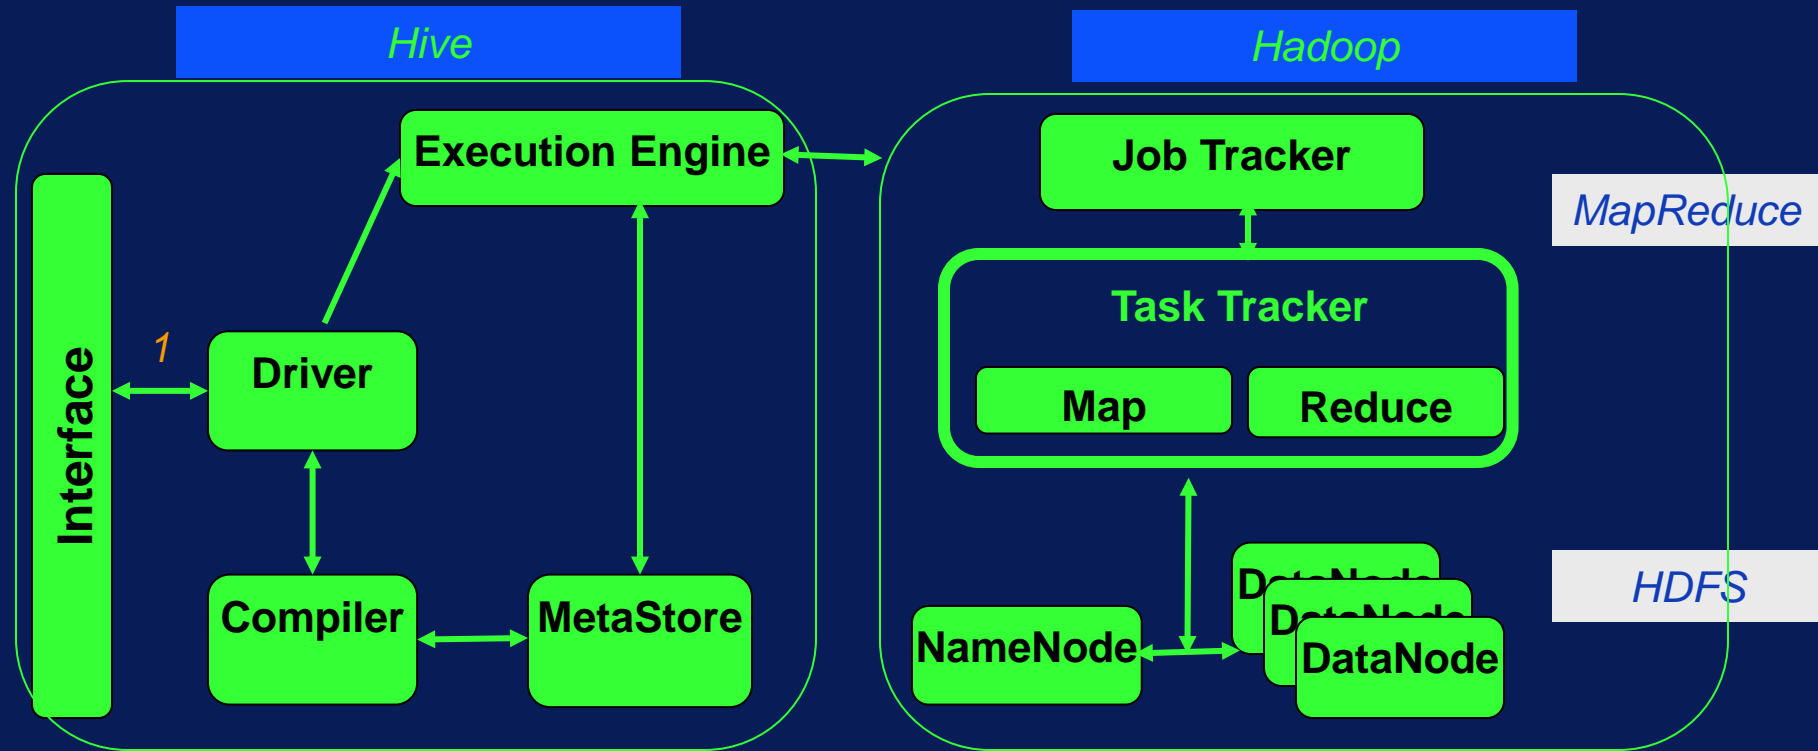
MetaStore

Map

Reduce

NameNode

DataNode
DataNode
DataNode



Hive

Hadoop

Execution Engine

Job Tracker

MapReduce

Task Tracker

Map

Reduce

HDFS

NameNode

DataNode
DataNode
DataNode

Interface

Driver

Compiler

MetaStore

Hive

Hadoop

Execution Engine

Job Tracker

MapReduce

Task Tracker

Map

Reduce

HDFS

NameNode

DataNode

DataNode

DataNode

DataNode

Interface

Driver

Compiler

MetaStore

1

2

3

Hive

Hadoop

Execution Engine

Job Tracker

MapReduce

Task Tracker

Map

Reduce

HDFS

NameNode

DataNode

DataNode

DataNode

Interface

Driver

Compiler

MetaStore

1

2

3

4

Hive

Hadoop

Execution Engine

Job Tracker

MapReduce

Task Tracker

Map

Reduce

HDFS

NameNode

DataNode

DataNode

DataNode

DataNode

Interface

Driver

Compiler

MetaStore

1

2

5

3

4

Hive

Hadoop

Execution Engine

Job Tracker

MapReduce

Task Tracker

HDFS

Interface

Driver

Compiler

MetaStore

Map

Reduce

NameNode

DataNode

1

6

2

5

3

4

Hive

Hadoop

Execution Engine

Job Tracker

MapReduce

Task Tracker

HDFS

Interface

Driver

Compiler

MetaStore

Map

Reduce

NameNode

DataNode

1

6

2

5

3

4

7b

7a
8

Hive

Hadoop

Execution Engine

Job Tracker

MapReduce

Task Tracker

HDFS

Interface

Driver

Compiler

MetaStore

Map

Reduce

NameNode

DataNode
DataNode
DataNode

1

6

9

5

2

3

4

7a

8

7b

Hive

Hadoop

Execution Engine

Job Tracker

MapReduce

Task Tracker

HDFS

Interface

Driver

Compiler

MetaStore

Map

Reduce

NameNode

DataNode

1
10

6
5
2

3
4

9

7b

7a
8

Hive's Data Units

Databases

Tables

Partitions

Buckets (or clusters)

Very similar to SQL and Relational DBs

3-Levels: Tables → Partitions → Buckets

Data Model

Table maps to a HDFS directory

Partition maps to sub-directories under the table

Bucket maps to files under each partition

Tables

Similar to tables in relational DBs

Each table has corresponding
directory in HDFS

Partitions

Analogous to dense indexes on partition columns

Nested sub-directories in HDFS for each combination of partition column values

Allows users to efficiently retrieve rows

Hive Data Structures

Traditional Database concepts

Supports primitive types

Additional types and structures

Hive Data Structures

Traditional database concepts

Tables

Rows

Columns

Partitions

Hive Data Structures

Basic types

Integers

Floats

Doubles

Strings

Hive File Formats

Hive enables users store different file formats

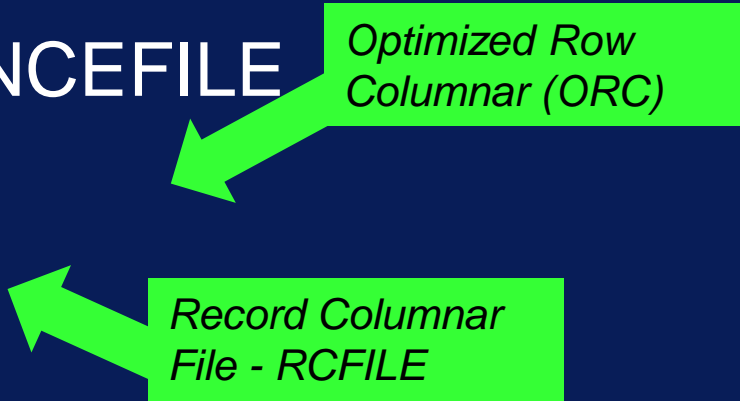
Performance improvements

TEXTFILE

SEQUENCEFILE

ORC

RCFILE



*Optimized Row
Columnar (ORC)*

*Record Columnar
File - RCFILE*

Hive Commands

Hive Interface

Command Line interface

Web interface or Hue

Java Database connectivity

Hive Commands

Database

Set of Tables - name conflicts resolution

Table

Set of Rows - have the same columns

Row

A single record - a set of columns

Column

Value and type for a single value

Tables Commands

- **SHOW TABLES**
- **CREATE TABLE**
- **ALTER TABLE**
- **DROP TABLE**

Hive Commands

```
CREATE TABLE mytable (myint INT, bar STRING)  
PARTITIONED BY (ds STRING);
```

```
SHOW TABLES '.*my';
```

*A table in Hive is an HDFS
directory in Hadoop*

```
ALTER TABLE mytable ADD COLUMNS (new_col  
INT);
```


```
DROP TABLE mytable;
```

Hive Commands

Schema is known at creation time (like DB schema)

Partitioned tables have “sub-directories”, one for each partition

```
CREATE TABLE mypeople (  
  id          int,  
  name        string  
)  
partitioned by (date string)
```



Hive Query Language

JOIN

```
SELECT t1.a1 as c1, t2.b1 as c2  
FROM t1 JOIN t2 ON (t1.a2 = t2.b2);
```

INSERTION

```
INSERT OVERWRITE TABLE t1  
SELECT * FROM t2;
```


Format Rows

CREATE TABLE mypeople (id INT, name STRING)

ROW FORMAT

**DELIMITED FIELDS TERMINATED BY <output
format>**

LINES TERMINATED BY '\n';



Loading Data into HIVE

HDFS

```
LOAD DATA INPATH 'mybigdata'  
[OVERWRITE] INTO TABLE mypeople;
```

Local file system

```
LOAD DATA LOCAL INPATH 'mybigdata'  
INTO TABLE mypeople;
```

Partitions

```
LOAD DATA INPATH 'myweblogs' INTO TABLE  
mypeople PARTITION (dt=12-12-2020);
```


BUCKETS

Set `hive.enforce.bucketing` property to true

```
CREATE TABLE mycustomers(id INT, purchases DOUBLE,  
name STRING)  
CLUSTERED BY id into 32 BUCKETS;
```

BUCKETS

```
SELECT min(cost) FROM mysales  
TABLESAMPLE (BUCKET 10 OUT OF 32 ON  
rand());
```

VIEWS

Similar to SQL Views

Virtual table in Metastore

SHOW TABLES

JOINS

LEFT OUTER JOIN

RIGHT OUTER JOIN

FULL OUTER JOIN

```
hive> SELECT c.ID, c.NAME, c.AGE, o.AMOUNT FROM  
CUSTOMERS c JOIN ORDERS o ON  
(c.ID = o.CUSTOMER_ID);
```

DROP TABLE

- **DROP TABLE MyCustomers;**

DELETE PARTITION

```
ALTER TABLE MyCustomers DROP  
PARTITION (col2=100);
```