

Lecture 3

Mechanics of Bitcoin

Recap: Bitcoin consensus

Bitcoin consensus gives us:

- Append-only ledger
- Decentralized consensus
- Miners to validate transactions

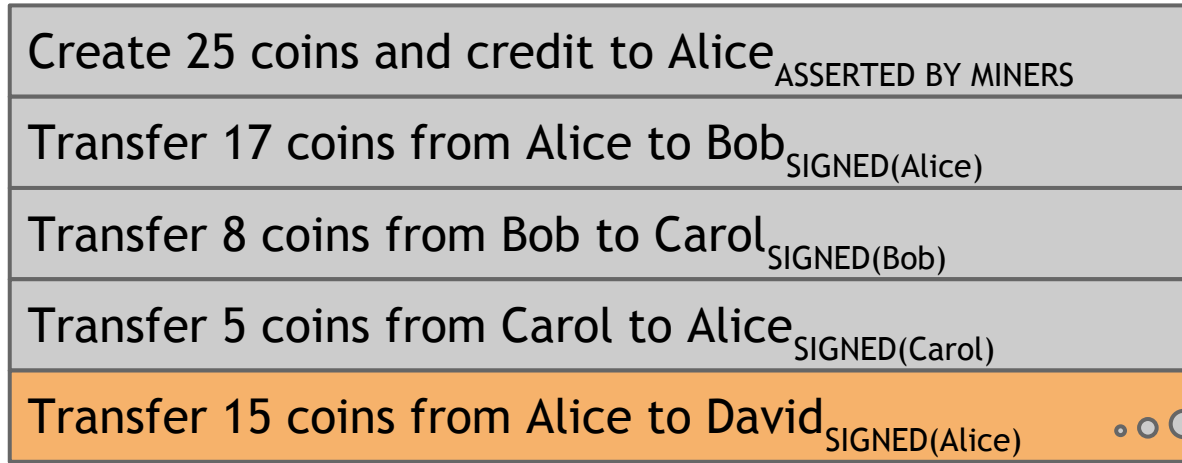
assuming a currency exists to motivate miners!

Lecture 3.1:

Bitcoin transactions

An account-based ledger (*not* Bitcoin)

time

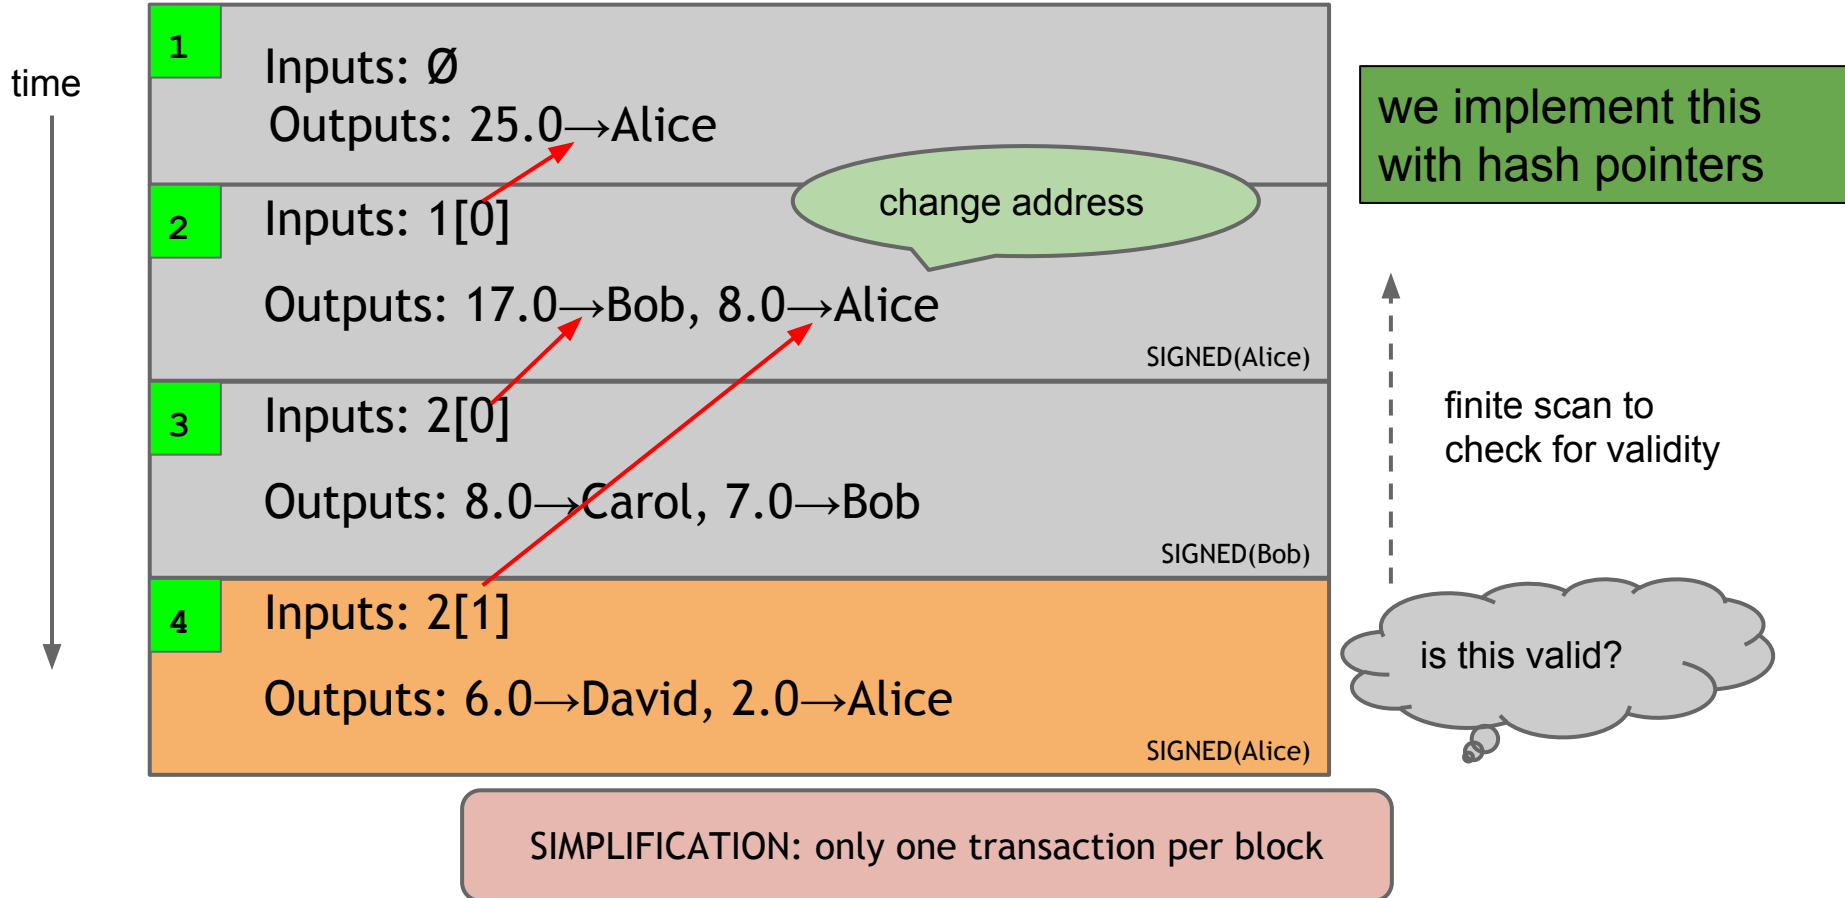


might need to
scan backwards
until genesis!

is this valid?

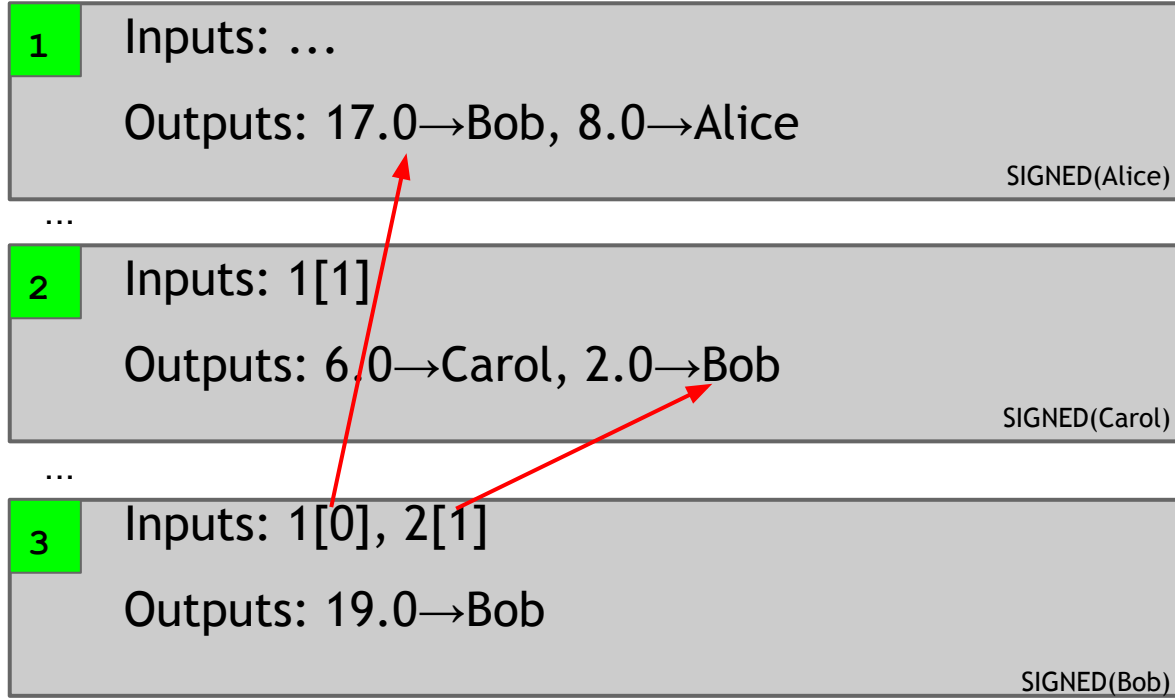
SIMPLIFICATION: only one transaction per block

A transaction-based ledger (Bitcoin)



Merging value

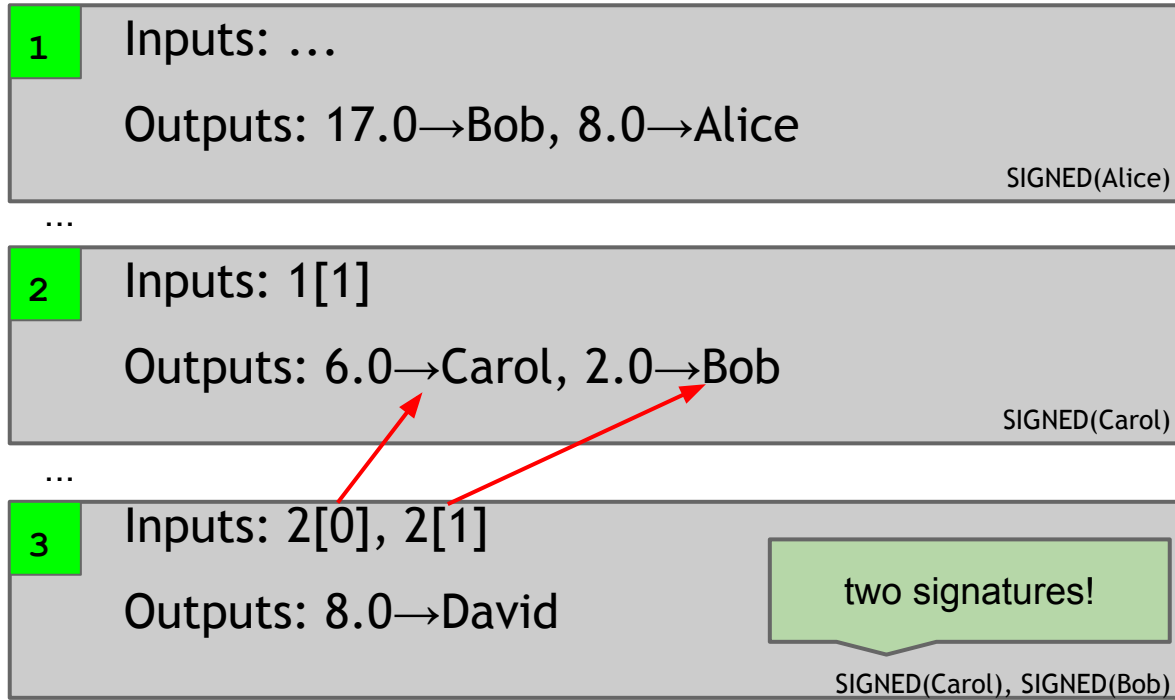
time



SIMPLIFICATION: only one transaction per block

Joint payments

time

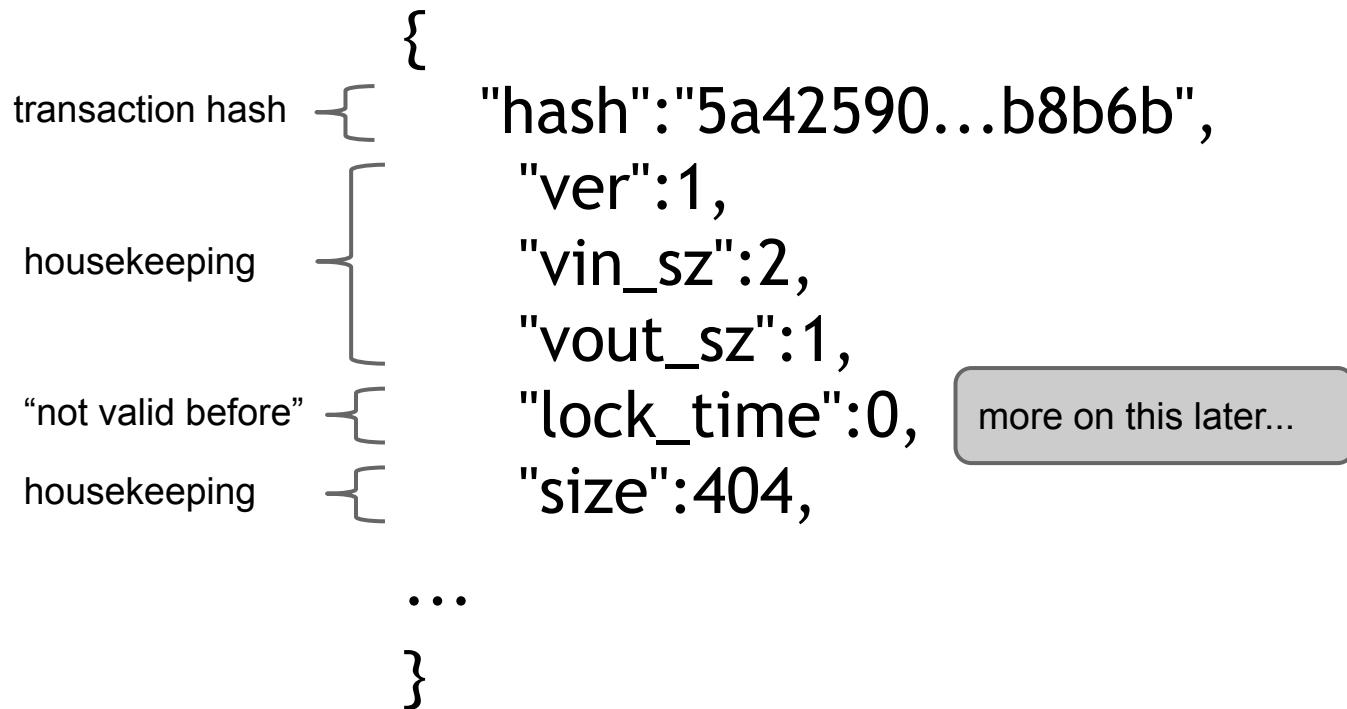


SIMPLIFICATION: only one transaction per block

The real deal: a Bitcoin transaction



The real deal: transaction metadata



The real deal: transaction inputs

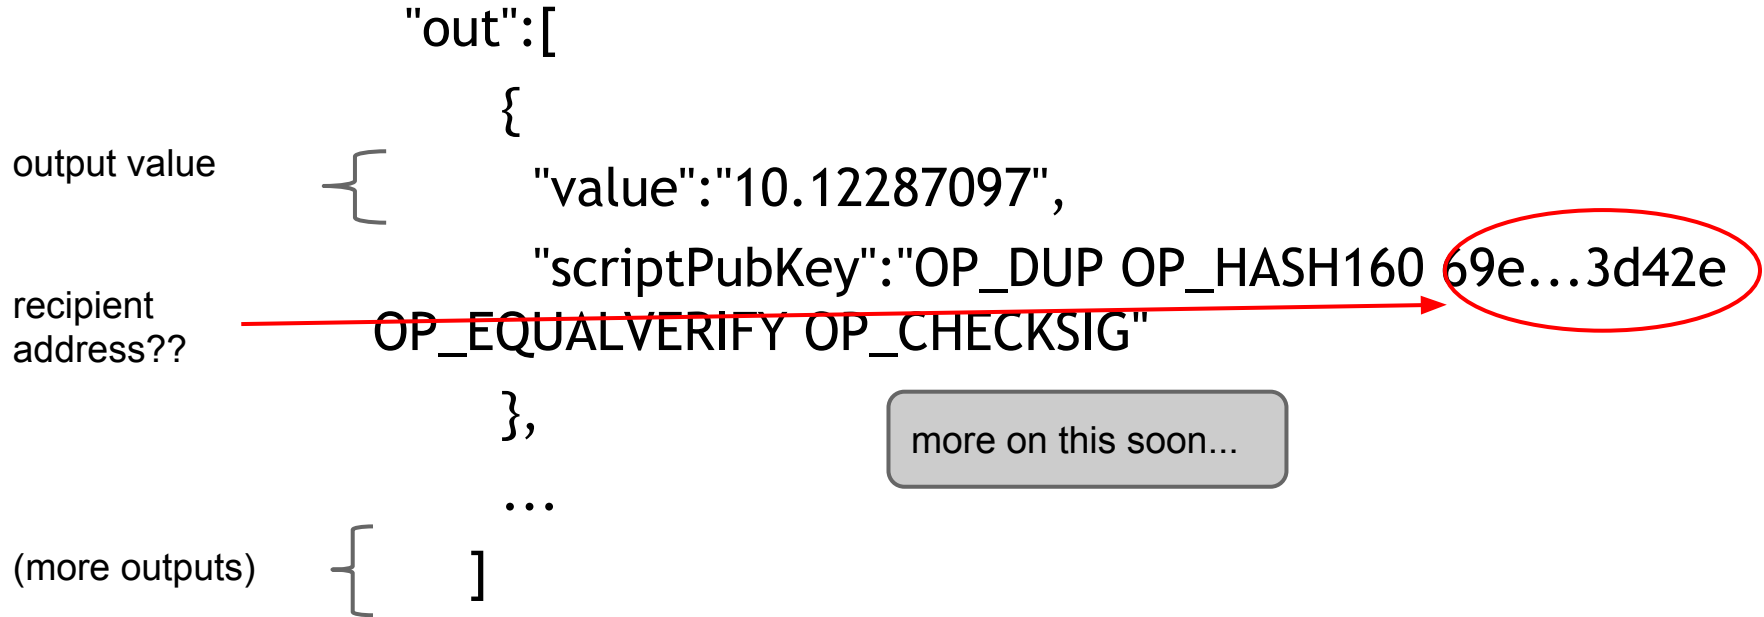
```
    "in":[  
      {  
        "prev_out":{  
          "hash":"3be4...80260",  
          "n":0  
        },  
        "scriptSig":"30440....3f3a4ce81"  
      },  
      ...  
    ],
```

previous transaction {

signature {

(more inputs) {

The real deal: transaction outputs



Lecture 3.2:

Bitcoin scripts

Output “addresses” are really *scripts*

OP_DUP

OP_HASH160

69e02e18...

OP_EQUALVERIFY OP_CHECKSIG

Input “addresses” are *also* scripts

scriptSig

30440220...
0467d2c9...

scriptPubKey

OP_DUP
OP_HASH160
69e02e18...
OP_EQUALVERIFY OP_CHECKSIG

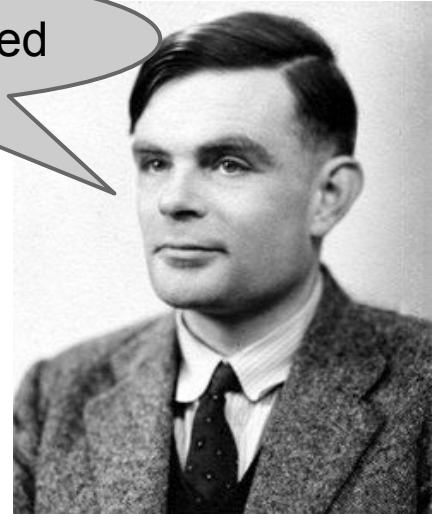
TO VERIFY: Concatenated script must execute completely with no errors

Bitcoin scripting language (“Script”)


Design goals

- Built for Bitcoin (inspired by Forth)
- Simple, compact
- Support for cryptography
- Stack-based
- Limits on time/memory
- No looping

I am not impressed



Bitcoin script execution example



`<sig> <pubKey> OP_DUP OP_HASH160 <pubKeyHash?> OP_EQUALVERIFY OP_CHECKSIG`

Bitcoin script instructions

256 opcodes total (15 disabled, 75 reserved)

- Arithmetic
- If/then
- Logic/data handling
- Crypto!
 - Hashes
 - Signature verification
 - Multi-signature verification


OP_CHECKMULTISIG

- Built-in support for joint signatures
- Specify n public keys
- Specify t
- Verification requires t signatures



BUG ALERT: Extra data value
popped from the stack and ignored

Bitcoin scripts in practice (as of 2014)

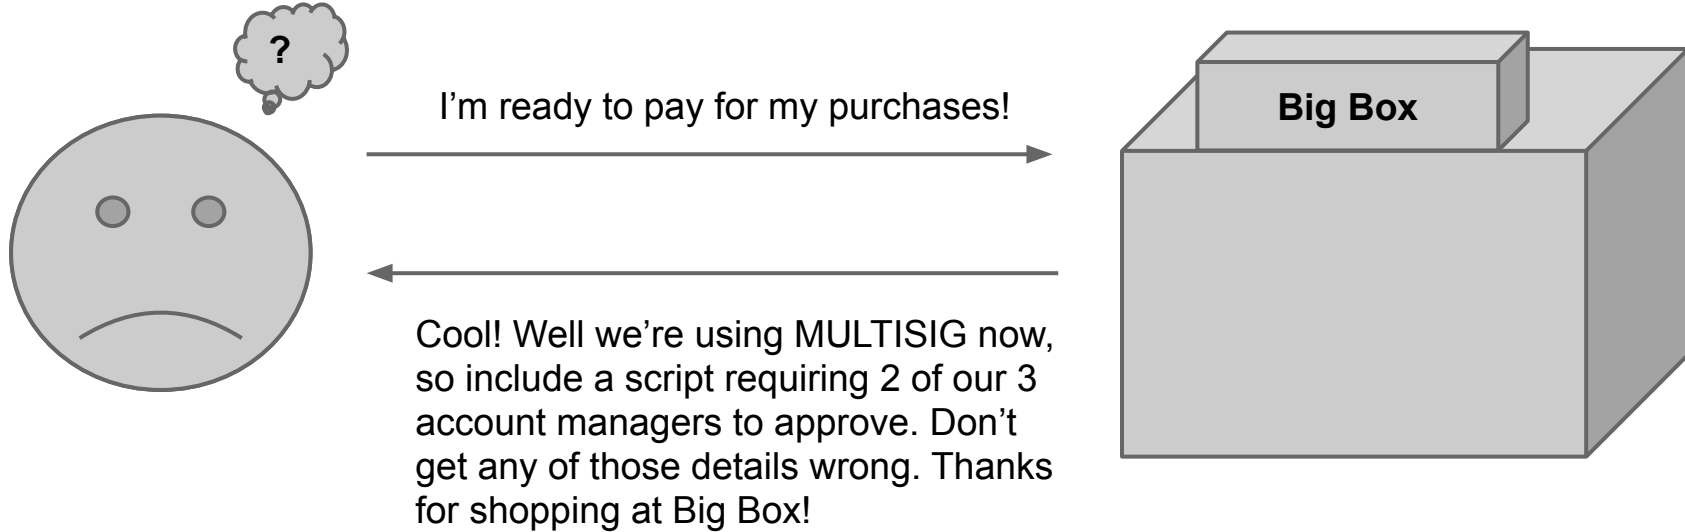
- Most nodes whitelist known scripts
- 99.9% are simple signature checks
- ~0.01% are MULTISIG 
- ~0.01% are **Pay-to-Script-Hash**
- Remainder are errors, proof-of-burn

Proof-of-burn

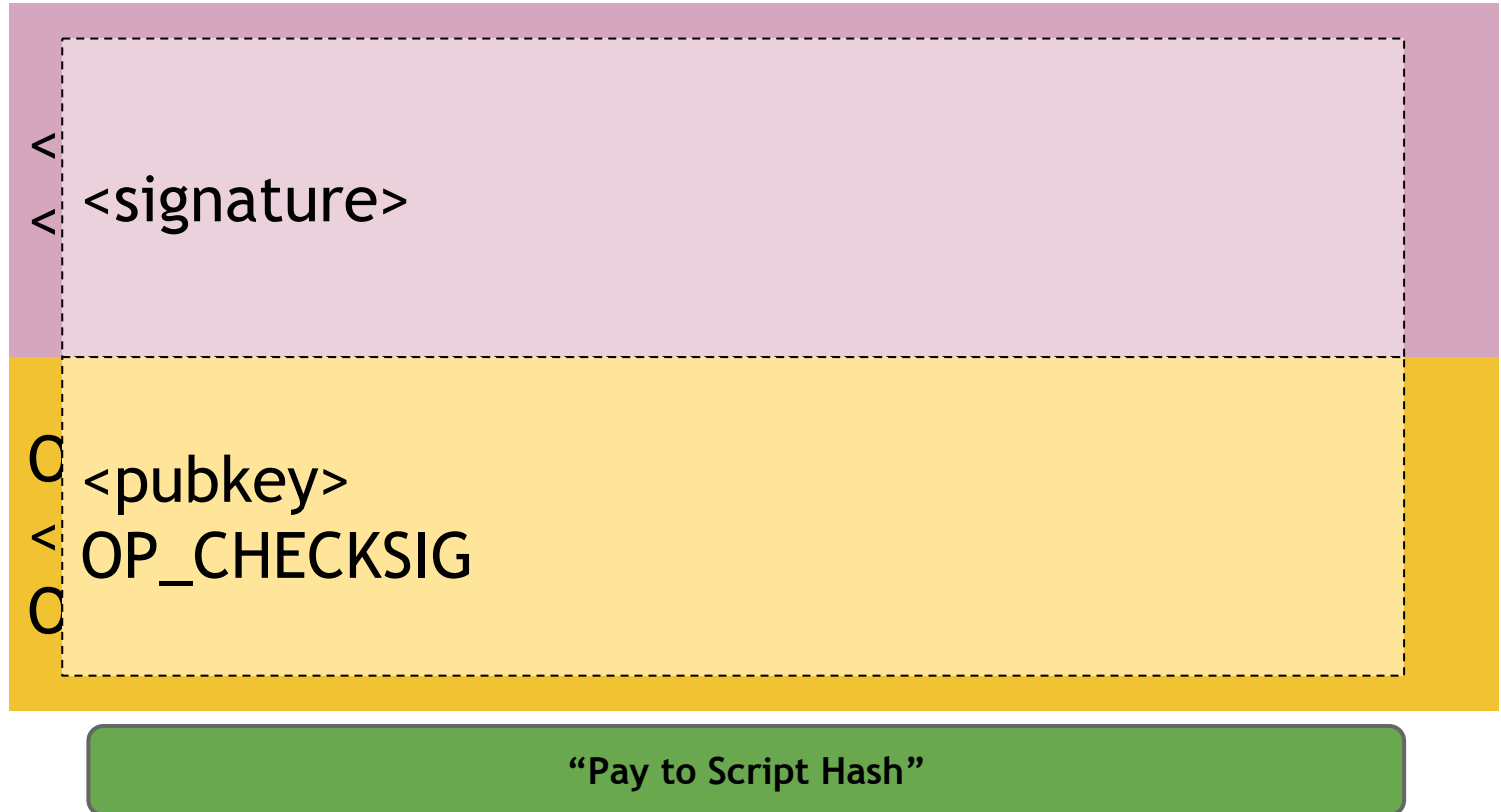
nothing's going to redeem that 😞

OP_RETURN
<arbitrary data>

Should senders specify scripts?



Idea: use the hash of redemption script




Pay to script hash



I'm ready to pay for my purchases!



Great! Here's our address: 0x3454

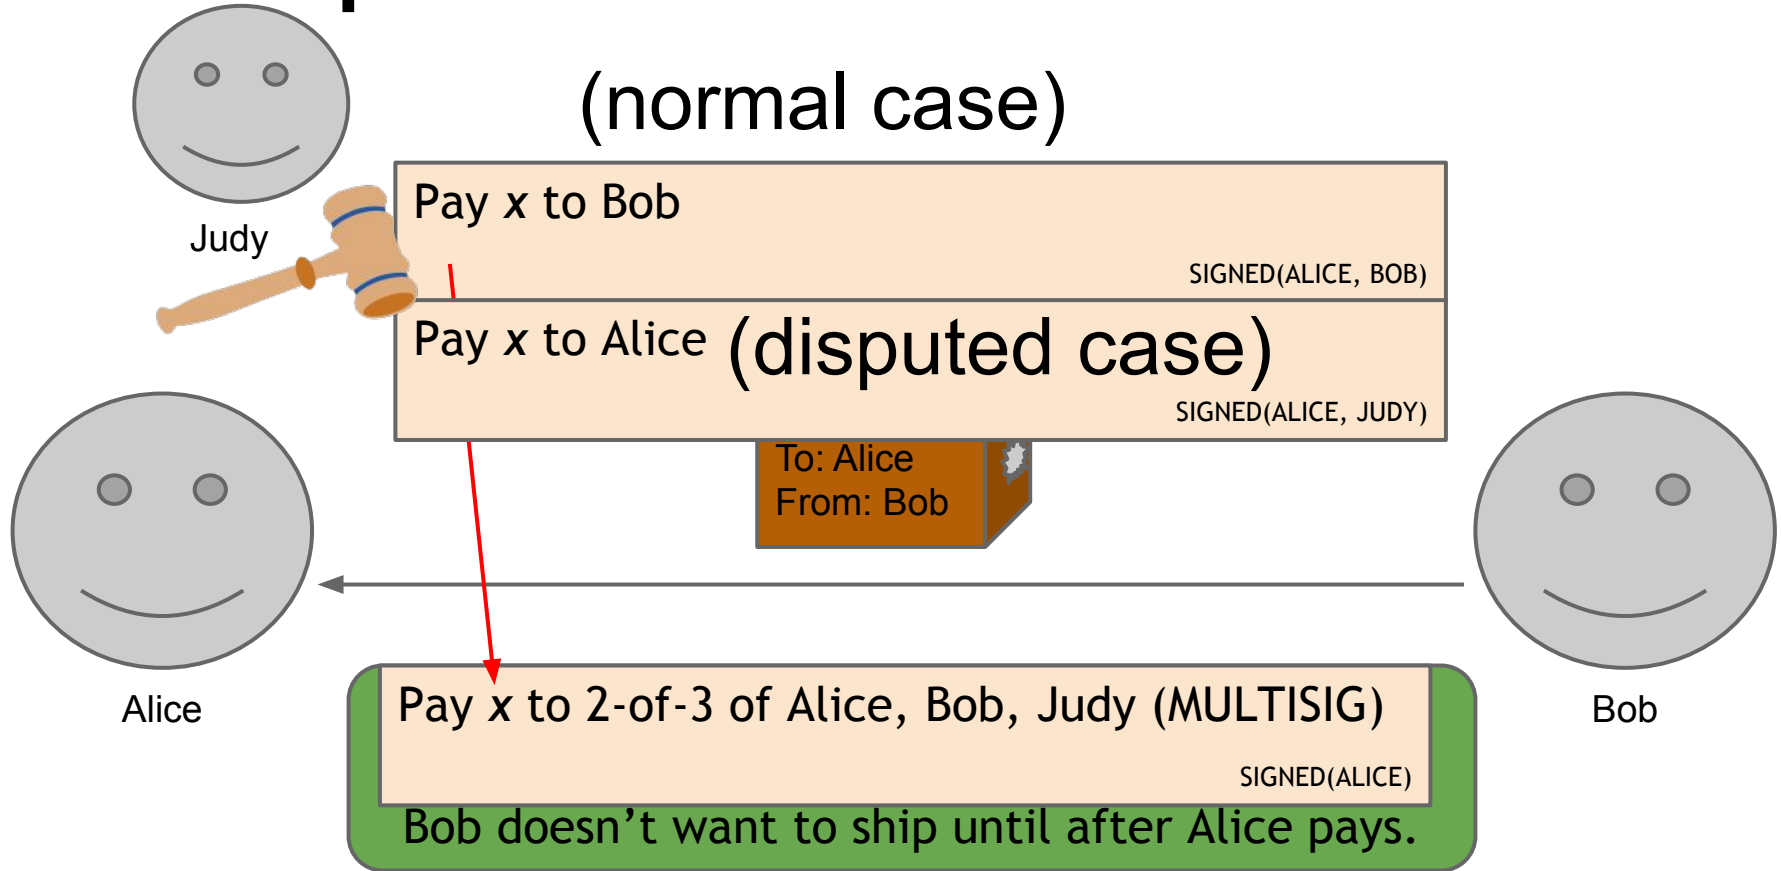


Lecture 3.3:


Applications of Bitcoin scripts

Example 1: Escrow transactions

(normal case)



Example 2: Green addresses



Bank

004 days since last double spend!

Faraday cage

Pay x to Bob, y to Bank

No double spend

SIGNED(BANK)

Alice

PROBLEM: Alice wants to pay Bob.
Bob can't wait 6 verifications to guard against double-spends, or is offline completely.

Bob

Example 3: Efficient micro-payments

What if Bob never signs??

all of these could
be
double-spends!

Input: x ; Pay 42 to Bob, 58 to Alice

SIGNED(ALICE) SIGNED(BOB)

...

Alice demands a timed refund transaction before starting

Input: x ; Pay 100 to Alice, LOCK until time t

SIGNED(ALICE) SIGNED(BOB)

I'm done!

Input: x ; Pay 03 to Bob, 97 to Alice

SIGNED(ALICE) _____

I'll publish!

Input: x ; Pay 02 to Bob, 98 to Alice

SIGNED(ALICE) _____

Input: x ; Pay 01 to Bob, 99 to Alice

SIGNED(ALICE) _____

PROBLEM: Alice wants to pay Bob for each

Input: y ; Pay 100 to Bob/Alice (MULTISIG)

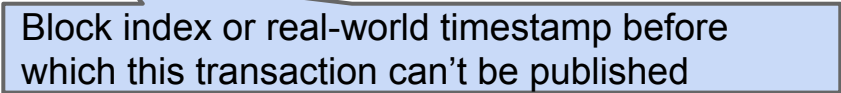
SIGNED(ALICE)

Alice

Bob

lock_time

```
{  
  "hash":"5a42590...b8b6b",  
  "ver":1,  
  "vin_sz":2,  
  "vout_sz":1,  
  "lock_time":315415,  
  "size":404,  
  ...  
}
```



Block index or real-world timestamp before
which this transaction can't be published

More advanced scripts

- Multiplayer lotteries
- Hash pre-image challenges
- Coin-swapping protocols
 - Don't miss the lecture on anonymity!

“Smart contracts”

Lecture 3.4:

Bitcoin blocks

Bitcoin blocks

Why bundle transactions together?

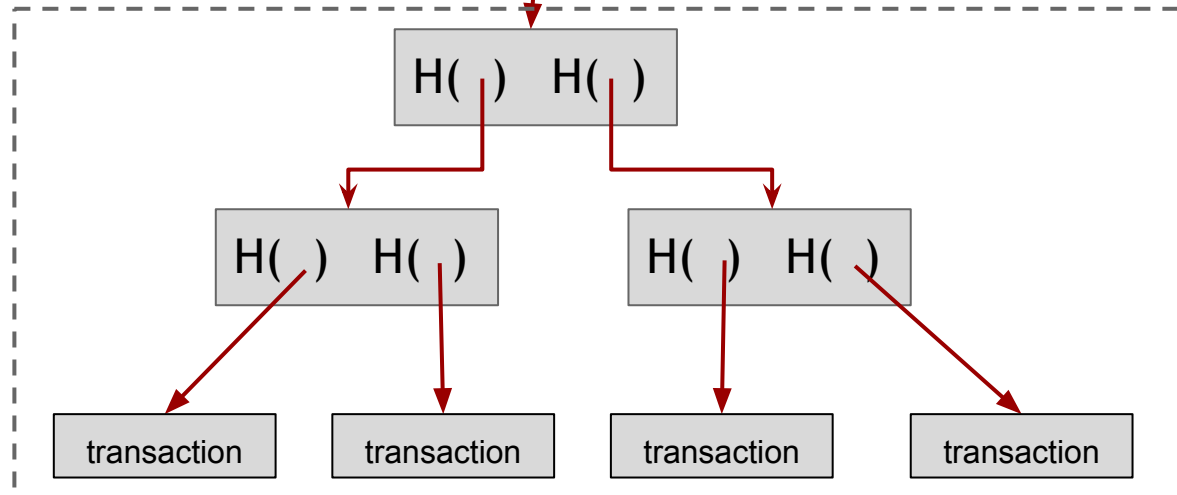
- Single unit of work for miners
- Limit length of hash-chain of blocks
 - Faster to verify history

Bitcoin block structure

Hash chain of blocks



Hash tree (Merkle tree) of transactions in each block



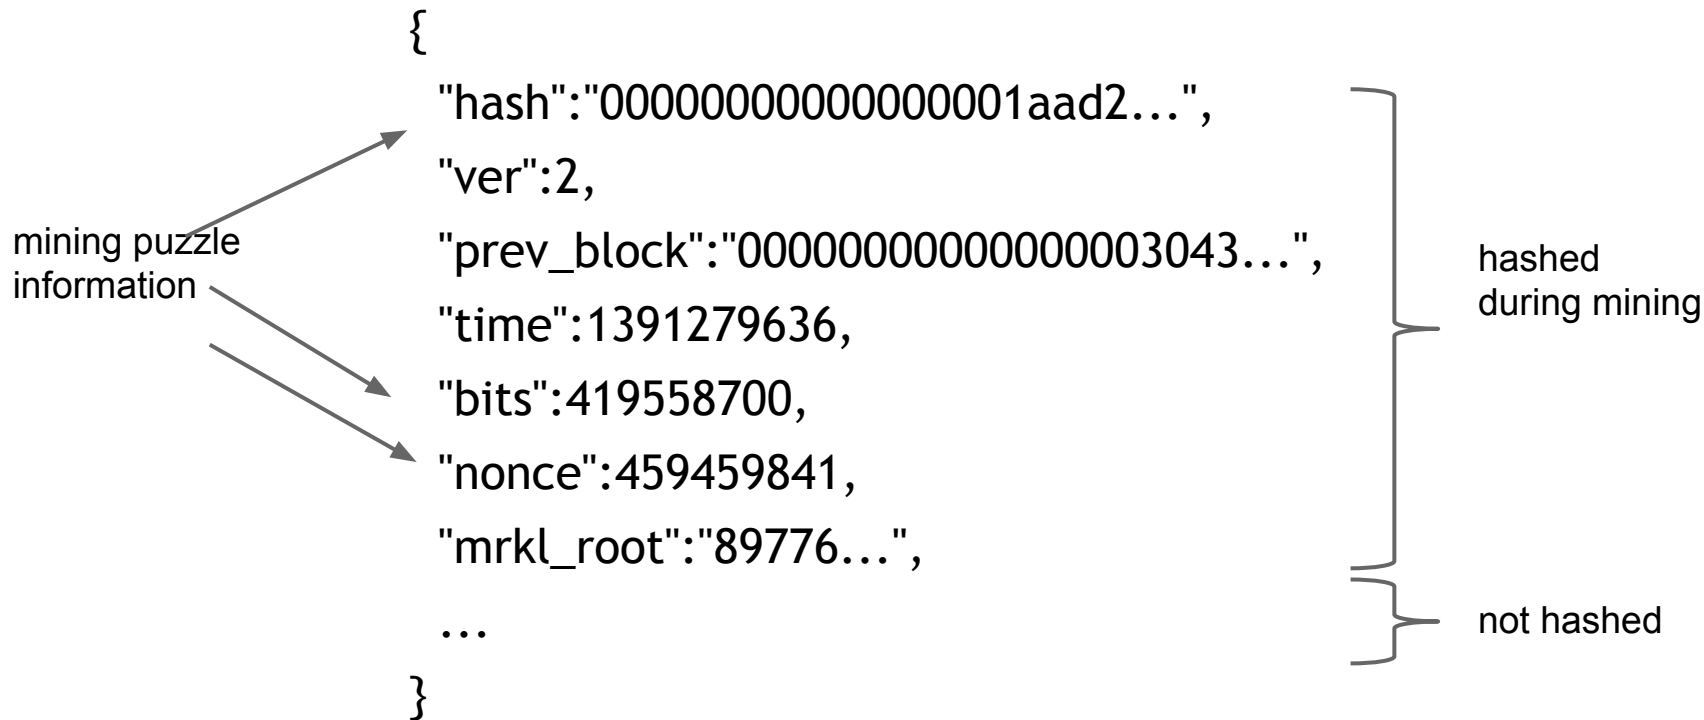
The real deal: a Bitcoin block

block header

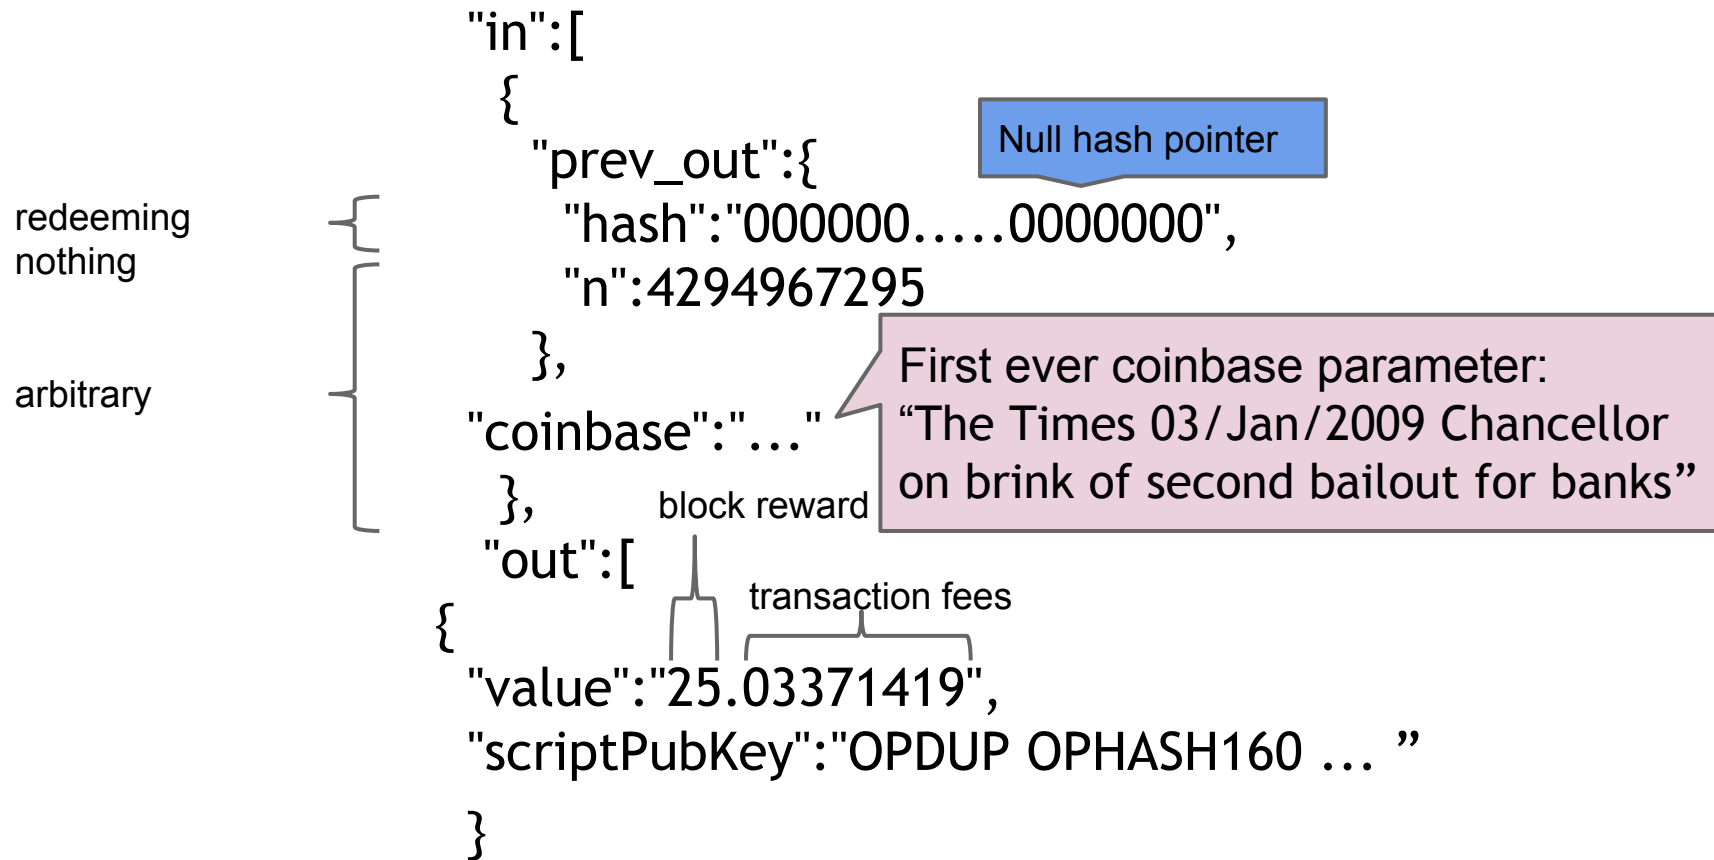
transaction data

```
{
  "hash": "0000000000000001aad2...",
  "ver": 2,
  "prev_block": "0000000000000003043...",
  "time": 1391279636,
  "bits": 419558700,
  "nonce": 459459841,
  "mrkl_root": "89776...",
  "n_tx": 354,
  "size": 181520,
  "tx": [
    ...
  ],
  "mrkl_tree": [
    "6bd5eb25...",
    ...
    "89776cdb..."
  ]
}
```

The real deal: a Bitcoin block header



The real deal: coinbase transaction



See for yourself!

Transaction View information about a bitcoin transaction

151b750d1f13e76d84e82b34b12688811b23a8e3119a1cba4b4810f9b0ef408d

1KryFUt9tXHvaoCYTNPbqpWPJKQ717YmL5

➔

1KvrdrQ3oGqMAiDTMEYCcdDSnVaGNW2YZh

1KryFUt9tXHvaoCYTNPbqpWPJKQ717YmL5

1.0194 BTC

3.458 BTC

9 Confirmations

4.4774 BTC

Summary	
Size	257 (bytes)
Received Time	2014-08-05 01:55:25
Included In Blocks	314018 (2014-08-05 02:00:40 +5 minutes)
Confirmations	9 Confirmations
Relayed by IP ⓘ	Blockchain.info
Visualize	View Tree Chart

Inputs and Outputs	
Total Input	4.4775 BTC
Total Output	4.4774 BTC
Fees	0.0001 BTC
Estimated BTC Transacted	1.0194 BTC
Scripts	Show scripts & coinbase

blockchain.info (and many other sites)

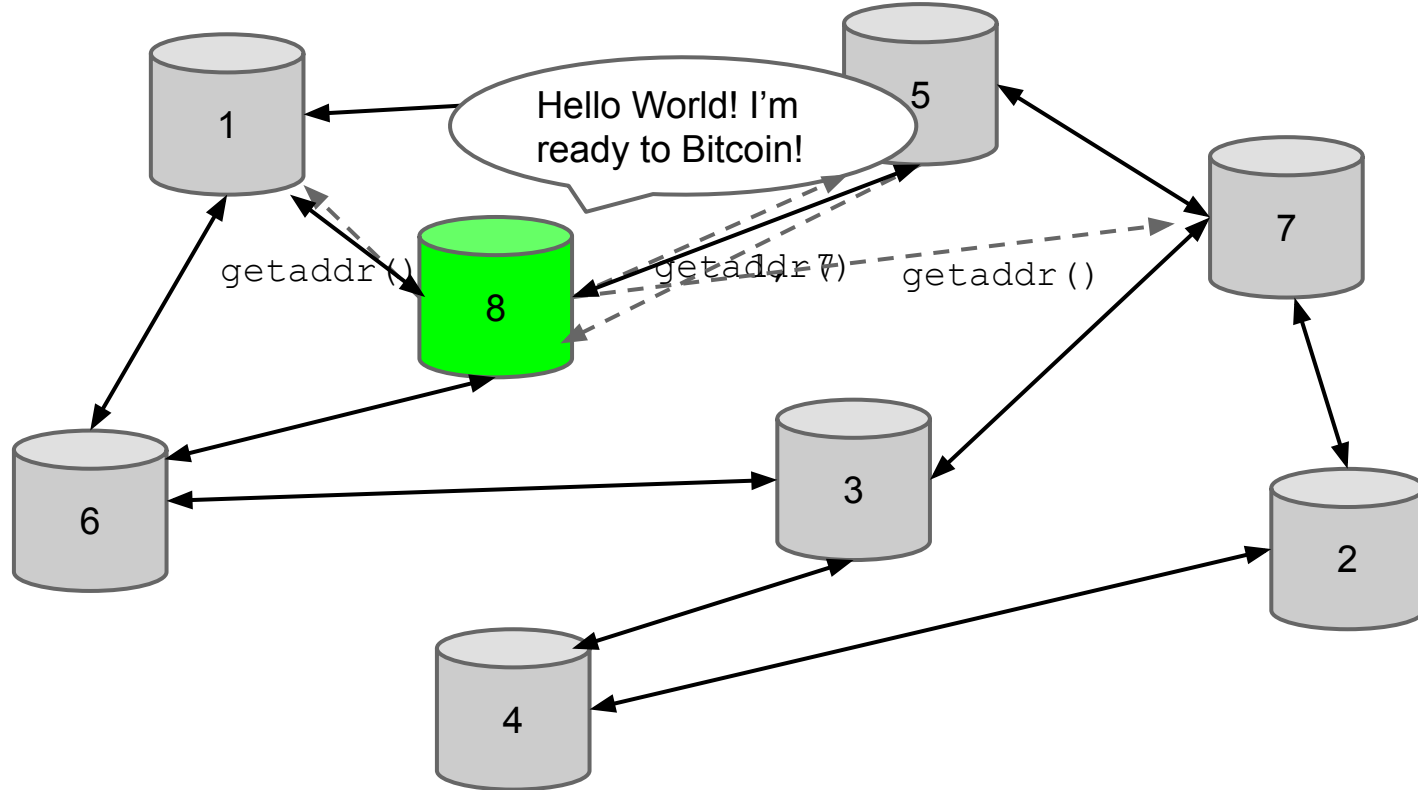
Lecture 3.5:

The Bitcoin network


Bitcoin P2P network

- Ad-hoc protocol (runs on TCP port 8333)
- Ad-hoc network with random topology
- All nodes are equal
- New nodes can join at any time
- Forget non-responding nodes after 3 hr

Joining the Bitcoin P2P network



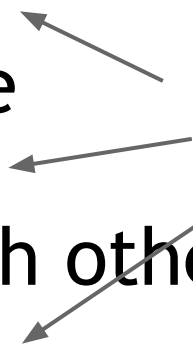
Transaction propagation (flooding)



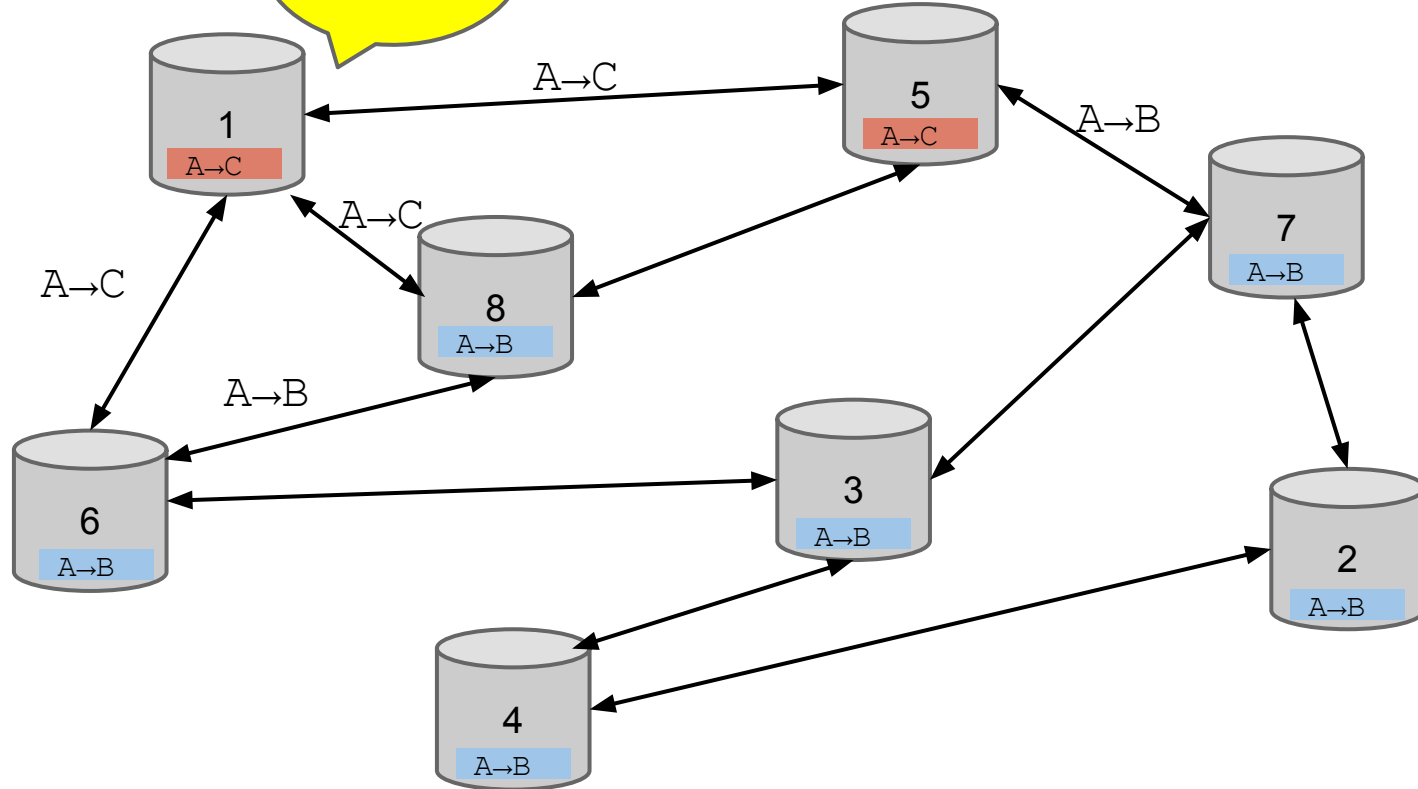
Should I relay a proposed transaction?

- Transaction valid with current block chain
- (default) script matches a whitelist
 - Avoid unusual scripts
- Haven't seen before
 - Avoid infinite loops
- Doesn't conflict with others I've relayed
 - Avoid double-spends

Sanity checks only...
Some nodes may ignore them!

Three arrows originate from the green box and point to specific items in the list: one points to 'Avoid unusual scripts', one points to 'Avoid infinite loops', and one points to 'Avoid double-spends'.

Nodes differ on transaction pool



Race conditions

Transactions or blocks may *conflict*


- Default behavior: accept what you hear first
- Network position matters
- Miners may implement other logic!

Stay tune for our lecture on mining!

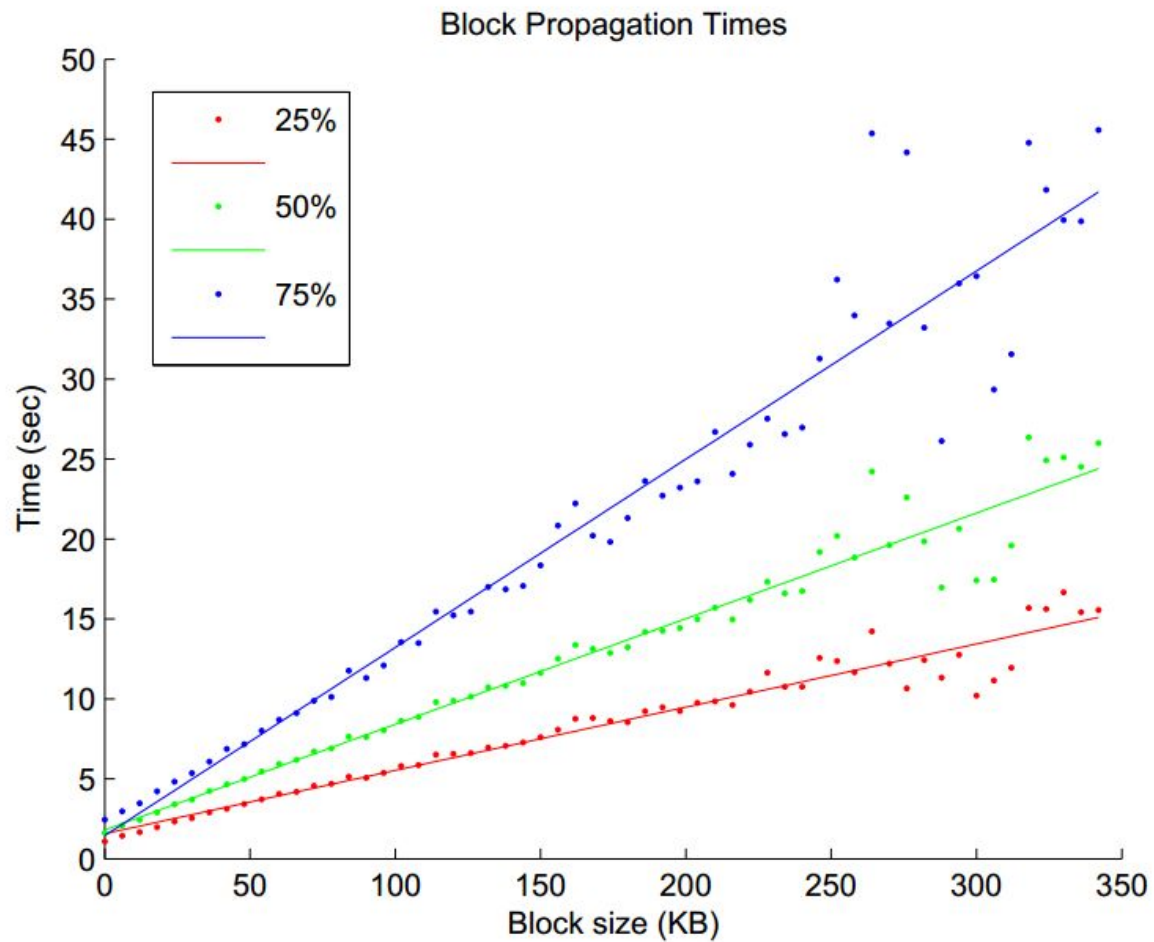
Block propagation nearly identical

Relay a new block when you hear it if:

- Block meets the hash target
- Block has all valid transactions
 - Run *all* scripts, even if you wouldn't relay
- Block builds on current longest chain
 - Avoid forks



Sanity check
Also may be ignored...



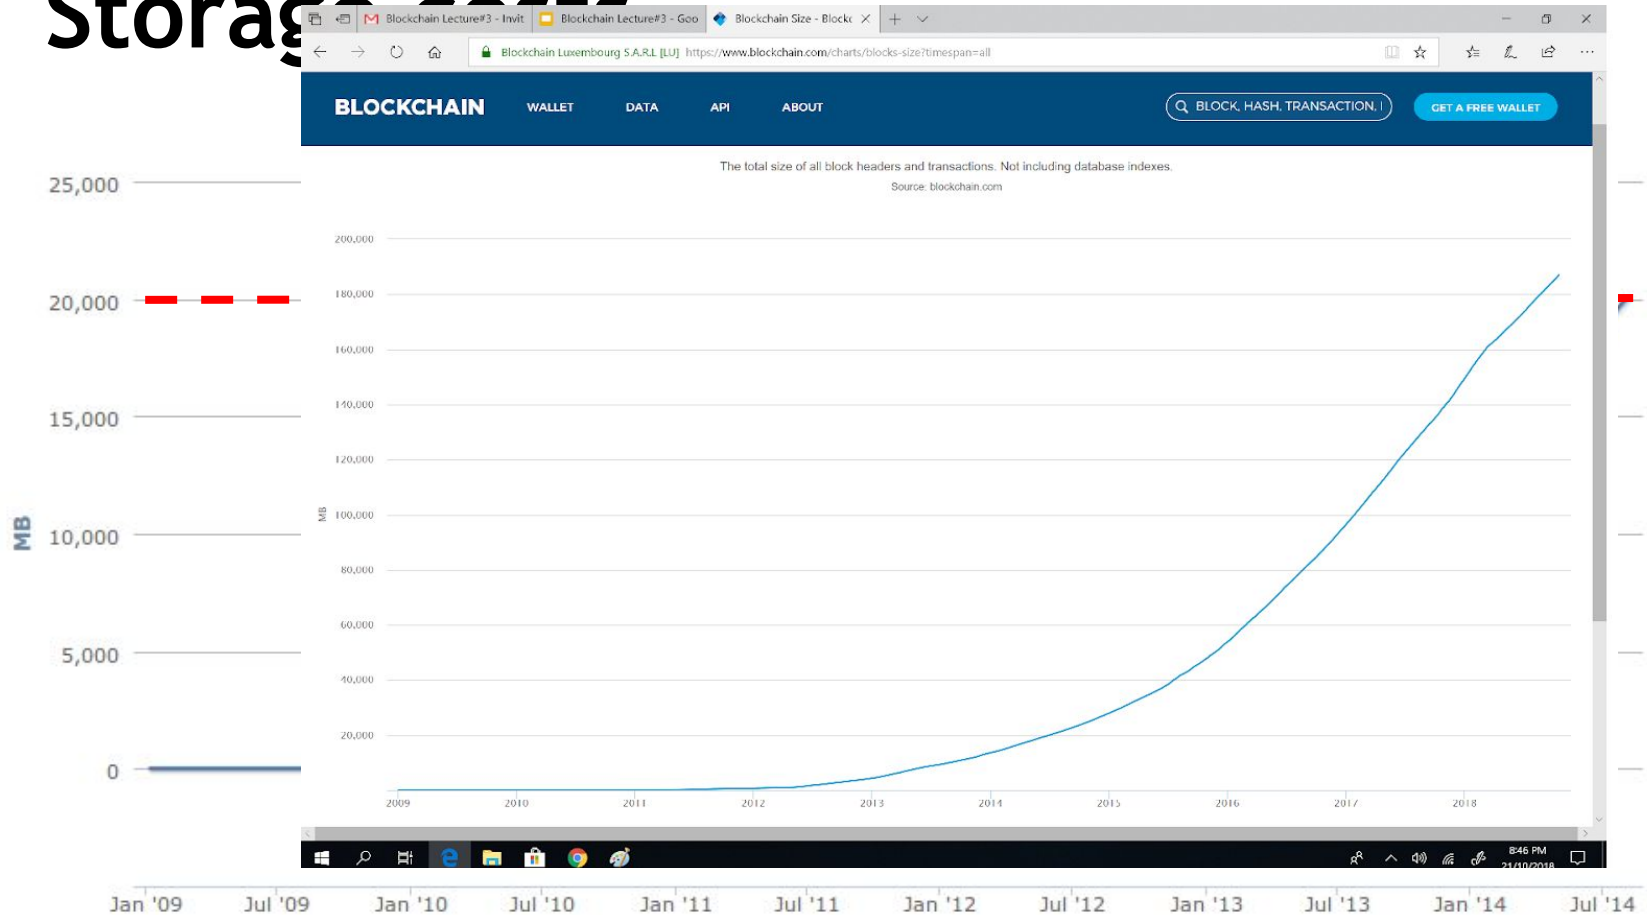
How big is the network?

- Impossible to measure exactly
- Estimates-up to 1M IP addresses/month
- Only about 5-10k “full nodes”
 - Permanently connected
 - Fully-validate
- This number may be dropping!

Fully-validating nodes

- Permanently connected
- Store entire block chain
- Hear and forward every node/transaction

Storage costs



Tracking the UTXO set

- **U**nspent **T**ransaction **O**utput
 - Everything else can be stored on disk
- Currently ~12 M UTXOs
 - Out of 44 M transactions
- Can easily fit into RAM

Thin/SPV clients (not fully-validating)

Idea: don't store everything

- Store block headers only
- Request transactions as needed
 - To verify incoming payment
- Trust fully-validating nodes

1000x cost savings! (20 GB-23MB)

Software diversity

- About 90% of nodes run “Core Bitcoin” (C++)
 - Some are out of date versions
- Other implementations running successfully
 - BitcoinJ (Java)
 - Libbitcoin (C++)
 - btcd (Go)
- “Original Satoshi client”

Lecture 3.6:

Limitations & improvements

Hard-coded limits in Bitcoin

- 10 min. average creation time per block
- 1 M bytes in a block
- 20,000 signature operations per block
- 100 M *satoshis* per bitcoin
- 23M total bitcoins maximum
- 50,25,12.5... bitcoin mining reward

} These affect
economic
balance of power
too much to
change now

Throughput limits in Bitcoin

- 1 M bytes/block (10 min)
- >250 bytes/transaction
- 7 transactions/sec 😞

Compare to:


- VISA: 2,000-10,000 transactions/sec
- PayPal: 50-100 transaction/sec

Cryptographic limits in Bitcoin

- Only 1 signature algorithm (ECDSA/P256)
- Hard-coded hash functions

Crypto primitives might break by 2040...

“Hard-forking” changes to Bitcoin



Soft forks

Observation: we can add new features which only *limit* the set of valid transactions

Need majority of nodes to enforce new rules

Old nodes will approve

RISK: Old nodes might mine now-invalid blocks

Soft fork example: pay to script hash

<signature>
<<pubkey> OP_CHECKSIG>

OP_HASH160
<hash of redemption script>
OP_EQUAL

Old nodes will just approve the hash, not run the embedded script

Soft fork possibilities

- New signature schemes
- Extra per-block metadata
 - Shove in the coinbase parameter
 - Commit to UTXO tree in each block

Hard forks

- New op codes
- Changes to size limits
- Changes to mining rate
- Many small bug fixes

Currently seem very unlikely to happen

Stay tuned for our lecture on altcoins!

In the next lecture...

Human beings aren't Bitcoin nodes

- How do people interact with the network?
- How do people exchange bitcoins for cash?
- How do people securely store bitcoins?

Currency needs to work for people, not nodes